



THE UNIVERSITY  
*of* EDINBURGH

# Advanced Databases

Spring 2020

---

Lecture #01:

## Course Intro & Relational Databases

Milos Nikolic

# COURSE OVERVIEW

This course is on the **design** and **implementation** of disk-oriented database management systems

This is **not** a course on how to use a database to build applications or how to administer a database

Recommended prerequisite: Database Systems ([INFR10070](#))

# COURSE OUTLINE

Relational Databases

Storage and File Structure

Indexing

Query Evaluation

Query Optimisation

Transaction Management

Distributed Databases

# LEARNING OUTCOMES

Gain insights into how DBMSs function internally

Implement major components of a database system

Learn data management techniques that can help **YOU, the future scientist**, to transform data into knowledge

Useful concepts for CS and other sciences in general

# COURSE LOGISTICS

Course policies + schedule

Refer to course web page: <https://course.inf.ed.ac.uk/adbs>

Two in-class tutorials, no practical labs

All discussion + announcements will be on Piazza

<https://piazza.com/ed.ac.uk/spring2020/infr11011>

**Sign up now** with your student email address

# TEXTBOOK


## Database Management Systems

3<sup>rd</sup> Edition

R. Ramakrishnan and J. Gehrke

Most lectures will closely follow this book

Lectures will also refer to research papers



# ASSESSMENT

## Coursework (30%)

No requirement to pass the coursework

## Written Exam (70%)

School of Informatics uses a Common Marking Scheme

1<sup>st</sup> class or MSc distinction: 70% and above

# COURSEWORK

Implement features in an educational database system in C++

1 formative assignment (0%)

Purpose: get feedback & accommodate to DBMS and C++

1 summative assignment (30%)

Project: implement features + run experiments + write report

School of Informatics has a policy on coursework deadlines



# PLAGIARISM POLICY

All assignments must be your own work

They are **not** group assignments

You may **not** copy source code from other people or the web

You may **not** use public repositories to host your code

See [UoE Academic misconduct](#) for more information



# ACKNOWLEDGEMENTS

The lecture notes draw on notes by several people to which I am grateful, in particular:

D. Olteanu and T. Furche from Oxford

P. Guagliardo and S. Viglas from Edinburgh

A. Pavlo from CMU

T. Grust from Tübingen

J. Gehrke from Microsoft

# DATABASES

# DATABASE

Organised collection of inter-related data that models some aspect of the real-world

Databases are the core component of most computer applications

- Banking
- Web and mobile apps
- Sales
- Online retailers
- Human resources



# DATABASE EXAMPLE

Create a database that models a **university organisation** to keep track of students, instructors, and courses

Application program examples:

- Add new students, instructors, and courses

- Register students for courses and generate class rosters

- Assign grades to students, compute GPA, and generate transcripts

# FLAT FILE STRAWMAN

Store our database as comma-separated value (CSV) files

**Instructor(name, dept, salary)**

```
"Jones", "CS", 95000  
"Smith", "Physics", 75000  
"Gold", "CS", 62000
```

instructor.csv

**Course(name, instructor, year)**

```
"Databases", "Jones", 2018  
"Quantum M.", "Smith", 2017  
"Compilers", "Jones", 2017
```

course.csv

Apps have to parse the files each time they want to read/update records

# FLAT FILE STRAWMAN

Example: Get the names of all computer science instructors

**Instructor(name, dept, salary)**

```
"Jones", "CS", 95000  
"Smith", "Physics", 75000  
"Gold", "CS", 62000
```

instructor.csv



```
for line in file:  
    record = parse(line)  
    if "CS" == record[1]:  
        print record[0]
```

Tight coupling between application logic and physical storage

# FLAT FILE: DRAWBACKS

## Data redundancy and inconsistency

Multiple file formats, duplication of information in different files

## Difficulty in accessing data

Need to write a new program to carry out each new task

## Data isolation

Due to multiple files and formats writing new programs is difficult

## Integrity problems

Integrity constraints (e.g.,  $\text{balance} > 0$ ) become “buried” in program code

Hard to add new constraints or change existing ones



# FLAT FILE: DRAWBACKS (CONT.)

## Atomicity of updates

Failures may leave database in an inconsistent state with partial updates carried out

Ex: Moving money between two accounts should either complete or not happen at all

## Concurrent access by multiple users

Concurrent access needed for performance

Uncontrolled concurrent accesses can lead to inconsistencies

## Security problems

Hard to provide user access to some, but not all, data

**Database systems offer solutions to all the above problems**

# DATABASE MANAGEMENT SYSTEM (DBMS)

Software that stores, manages, and facilitates access to databases

Mediates interactions between users and databases

Provide users with an **abstract view** of the data

# LEVELS OF ABSTRACTIONS

View Level

Simplifies interaction with the database, hides info (e.g., salary) for security purposes

Logical Level

Describes data stored in the DB

```
type instructor = record
  name: string;
  dept: string;
  salary: integer;
end
```

Physical Level

Describes how a record is stored



**Data independence:** Insulate users from changes in lower levels

# DATA MODELS

## Data model

collection of concepts for describing the data in a database

## Schema

description of a particular collection of data, using a given model

## Models in practice

relational, key-value, graph, document, array, hierarchical, network

**Most DBMSs implement the relational data model**

# RELATIONAL MODEL

Proposed in 1970 by Edgar T. Codd

## Structure


The definition of relations and their contents

## Integrity

Ensure the database's contents satisfy constraints

## Manipulation

How to access and modify a database's contents



# RELATIONAL MODEL

Data organised in relations (tables)

## Relation schema

relation name + distinct (typed) attributes + constraints

## Relation is a set of records

Order of records is irrelevant

## Record is a set of attribute values

Values are (normally) atomic / scalar

**Instructor(name, dept, salary)**

name	dept	salary
Jones	CS	95000
Smith	Physics	75000
Gold	CS	62000

# RELATIONAL MODEL: PRIMARY KEYS

A primary key uniquely identifies a single tuple

DBMSs may automatically  
create an internal primary key  
if you don't define one

**Instructor(id, name, dept, salary)**

id	name	dept	salary
123	Jones	CS	95000
456	Smith	Physics	75000
789	Gold	CS	62000

Auto-generation of unique integer primary keys

**SEQUENCE** (SQL:2003)

**AUTO\_INCREMENT** (MySQL)

# RELATIONAL MODEL: FOREIGN KEYS

A foreign key specifies that an attribute from one relation has to map to a tuple in another relation



**Instructor(id, name, dept, salary)**

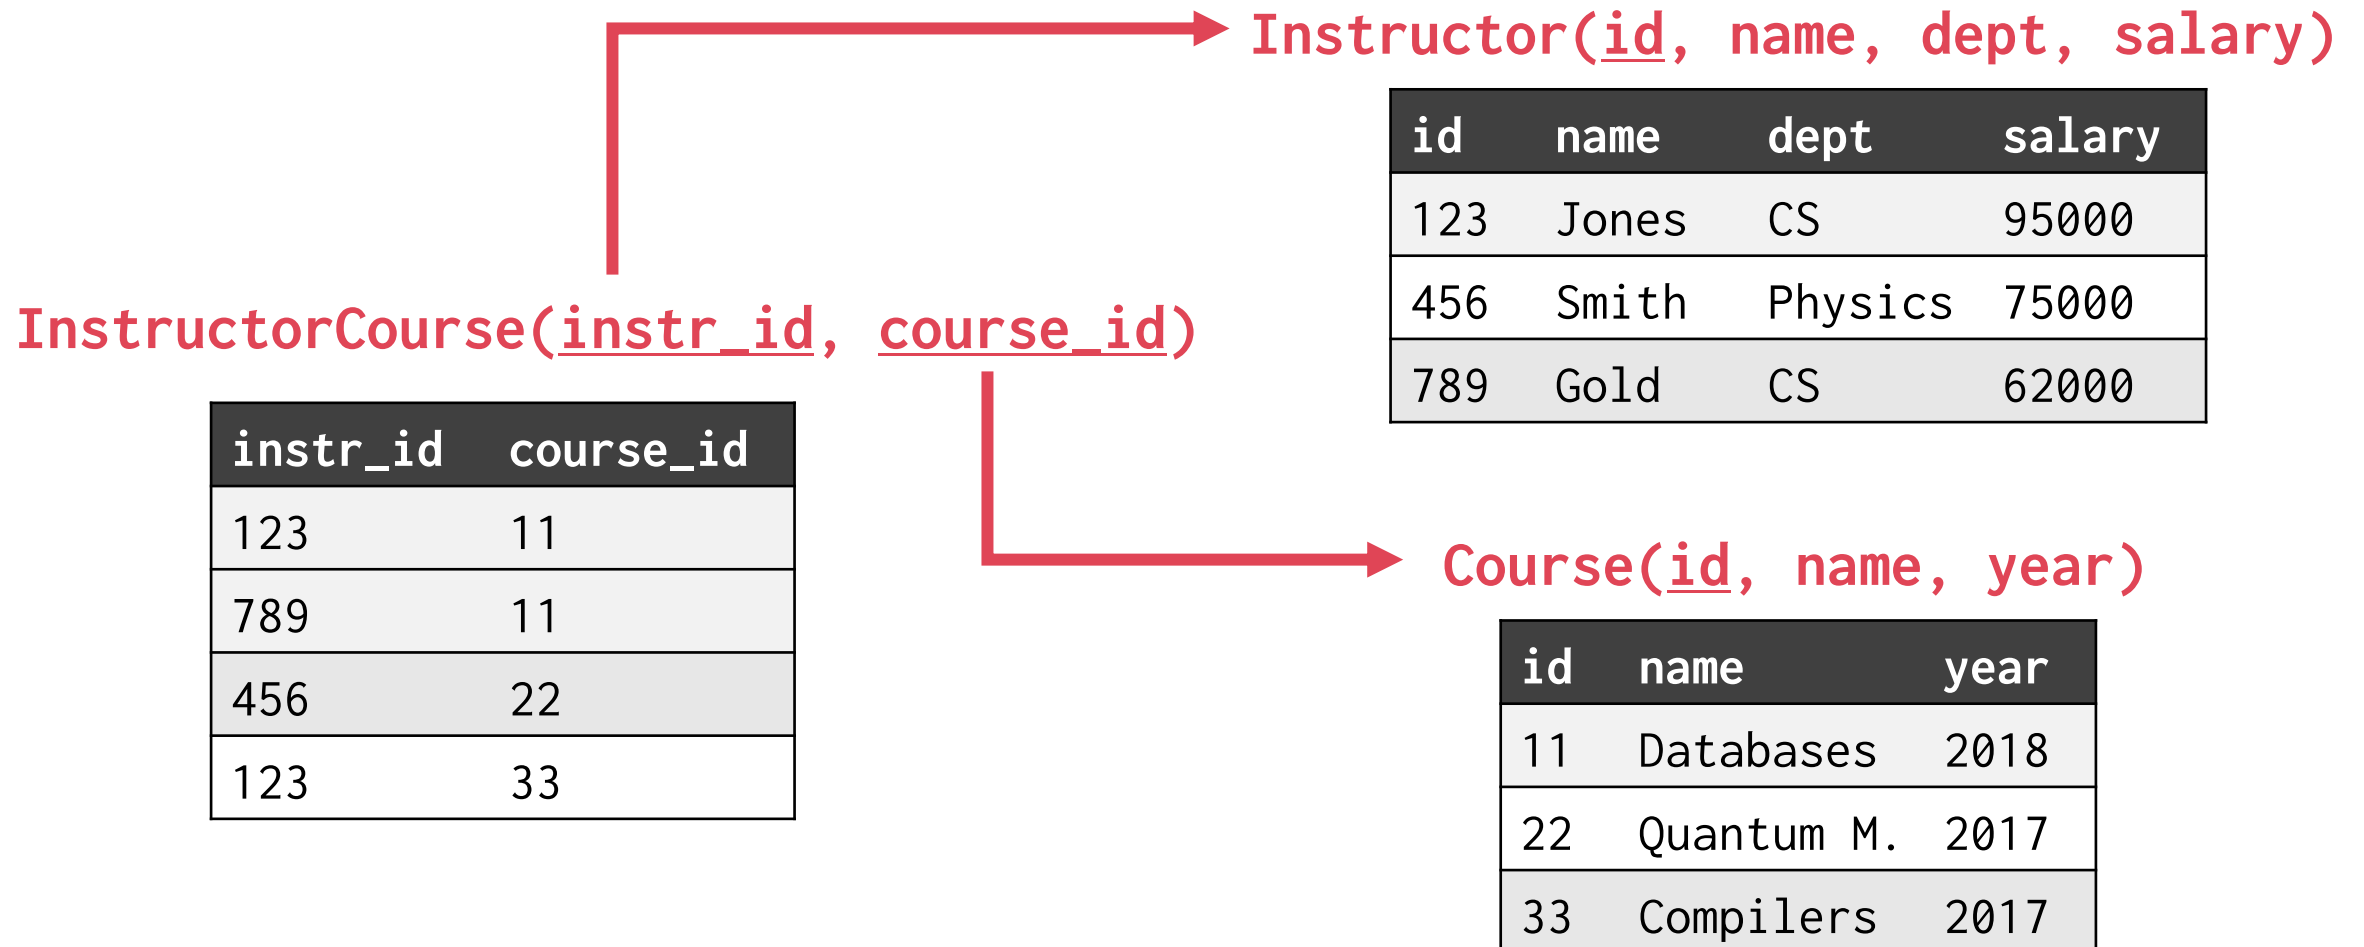
id	name	dept	salary
123	Jones	CS	95000
456	Smith	Physics	75000
789	Gold	CS	62000

**Course(id, name, instr, year)**

id	name	instr	year
11	Databases	123	2018
22	Quantum M.	456	2017
33	Compilers	123	2017



# RELATIONAL MODEL: FOREIGN KEYS



# RELATIONAL QUERY LANGUAGES

How to retrieve information from a database

## Procedural

The query specifies the (high-level) strategy the DBMS should use to find the desired result

← **Relational Algebra**

## Declarative

The query specifies only what data is wanted and not how to find it

← **Relational Calculus**

# RELATIONAL ALGEBRA

A relational algebra expression  
takes as input one or more relations  
applies a **sequence of operations**  
returns a relation as output

Example:  $\pi_{\text{name}}(\sigma_{\text{dept} = \text{'CS'}}(\text{Instructor}))$

RA is based on set semantics  
but can also be defined over bags (multisets)

$\sigma$	Selection
$\pi$	Projection
$\cup$	Union
$\cap$	Intersection
$-$	Difference
$\times$	Product
$\bowtie$	Natural Join
$\rho$	Renaming

# RELATIONAL ALGEBRA: SELECTION

Choose a subset of the tuples from a relation that satisfy a selection predicate

Can combine multiple predicates using conjunctions / disjunctions

$R(\text{aid}, \text{bid})$

aid	bid
a1	101
a2	102
a2	103
a3	104

Syntax:  $\sigma_{\text{predicate}}(R)$

$\sigma_{\text{aid}='a2'}(R)$

aid	bid
a2	102
a2	103

$\sigma_{\text{aid}='a2' \wedge \text{bid} > 102}(R)$

aid	bid
a2	103

# RELATIONAL ALGEBRA: PROJECTION

Generate a relation with tuples that contains only the specified attributes

Can rearrange attributes' ordering

Can manipulate the values

Syntax:  $\pi_{A_1, A_2, \dots, A_n}(R)$

$R(aid, bid)$

aid	bid
a1	101
a2	102
a2	103
a3	104

$\pi_{bid-100, aid}(\sigma_{aid='a2'}(R))$

bid-100	aid
2	a2
3	a2

# RELATIONAL ALGEBRA: UNION

Generate a relation that contains all tuples that appear in either only one or both input relations.

Relations must have the same set of attributes

Same holds for intersection and difference

Syntax: **R U S**

**R(aid, bid)**

aid	bid
a1	101
a2	102
a3	103

**S(aid, bid)**

aid	bid
a3	103
a4	104
a5	105

**R U S**

aid	bid
a1	101
a2	102
a3	103
a4	104
a5	105

# RELATIONAL ALGEBRA: PRODUCT

Generate a relation that contains all possible combinations of tuples from the input relations

Syntax:  $R \times S$

$R(aid, bid)$

aid	bid
a1	101
a2	102
a3	103

$S(bid, cid)$

bid	cid
b3	23
b4	24

$R \times S$

aid	R.bid	S.bid	cid
a1	101	b3	23
a1	101	b4	24
a2	102	b3	23
a2	102	b4	24
a3	103	b3	23
a3	103	b4	24

# RELATIONAL ALGEBRA: NATURAL JOIN

Generate a relation that contains all tuples that are a combination of two tuples (one from each input relation) with a common value(s) for one or more attributes

Syntax:  $R \bowtie S$

$R(\text{aid}, \text{bid})$

aid	bid
a1	101
a2	102
a3	103

$S(\text{bid}, \text{cid})$

bid	cid
101	c3
101	c4
105	c5

$R \bowtie S$

aid	bid	cid
a1	101	c3
a1	101	c3



# OBSERVATION

Relational algebra still defines the high-level steps of how to compute a query

$\sigma_{\text{bid} = 102} (R \bowtie S)$  vs.  $(R \bowtie (\sigma_{\text{bid} = 102} (S)))$

A better approach is to state the high-level answer that you want the DBMS to compute

Retrieve the joined tuples from **R** and **S** where **bid** equals 102

# RELATIONAL MODEL: QUERIES

**SQL** is the *de facto* standard

```
for line in file:  
    record = parse(line)  
    if "CS" == record[1]:  
        print record[0]
```

```
SELECT name FROM instructor  
WHERE dept = 'CS';
```

DBMS needs to figure out HOW to compute a query

The passage from **WHAT** to **HOW** goes through relational algebra

# NEXT LECTURE

# DATABASE ARCHITECTURE

