



THE UNIVERSITY  
*of* EDINBURGH

# Advanced Databases

Spring 2020

---

Lecture #11:

## Query Optimisation II

Milos Nikolic

# FINDING THE “BEST” QUERY PLAN

*Holy grail of any DBMS implementation*

**Challenge:** There may be more than one way to answer a given query

Which one of the join operators should we pick?

With which parameters (block size, buffer allocation, ...)?

Which join ordering?

# FINDING THE “BEST” QUERY PLAN

The query optimiser

1. **Enumerates** all possible query execution plans  
If this yields too many plans, at least enumerate the “promising” plan candidates
2. Determines the **cost** (quality) of each plan
3. Chooses the **best** one as the final execution plan

**Ideally:** Want to find the best plan. **Practically:** Avoid worst plans!

# JOIN OPTIMISATION: OVERVIEW

We have translated the query into a graph of query blocks

Query blocks are essentially a multi-way product of relations with projections on top

We can estimate the cost of a given execution plan

Use result size estimates in combination with the cost for individual join algorithms



**Task:** **enumerate** all possible execution plans

i.e., all possible 2-way join combinations for each query block

Example: three-way join

12 possible re-orderings

2 shown here



# ENORMOUS SEARCH SPACE

# of relations n	# of different join trees
2	2
3	12
4	120
5	1,680
6	30,240
7	665,280
8	17,297,280
10	17,643,225,600

We have not even considered *different join algorithms*!

We need to restrict search space!

# JOIN TREES: DYNAMIC PROGRAMMING

Use **dynamic programming** to reduce the number of cost estimations

Find the cheapest plan for an  $n$ -way join  $(R_1 \bowtie \dots \bowtie R_n)$  in  **$n$  passes**

In pass  $k$ , find the best plans for all  $k$ -relation sub-queries

Construct plans in pass  $k$  from best  $i$ -relation and  $(k-i)$ -relation subplans ( $1 \leq i < k$ )

Assumption: **“Principle of optimality”**

To find optimal global plan suffices to only consider optimal plans of sub-queries

Reduces considerably the search space, yet may lead to suboptimal plans

# EXAMPLE: DYNAMIC PROGRAMMING

**Pass #1** (best 1-relation plans): Find best access path to each relation (e.g., index, full table scans)




```
SELECT * FROM R, S, T
WHERE R.A = S.A
      AND S.B = T.B
```



# EXAMPLE: DYNAMIC PROGRAMMING

**Pass #2** (best 2-relation plans): determine best join order ( $R \bowtie S$  or  $S \bowtie R$ ), choose best candidate

```
SELECT * FROM R, S, T
WHERE R.A = S.A
      AND S.B = T.B
```




$R \bowtie S \bowtie T$



# EXAMPLE: DYNAMIC PROGRAMMING

**Pass #2** (best 2-relation plans): determine best join order ( $R \bowtie S$  or  $S \bowtie R$ ), choose best candidate

```
SELECT * FROM R, S, T
WHERE R.A = S.A
      AND S.B = T.B
```

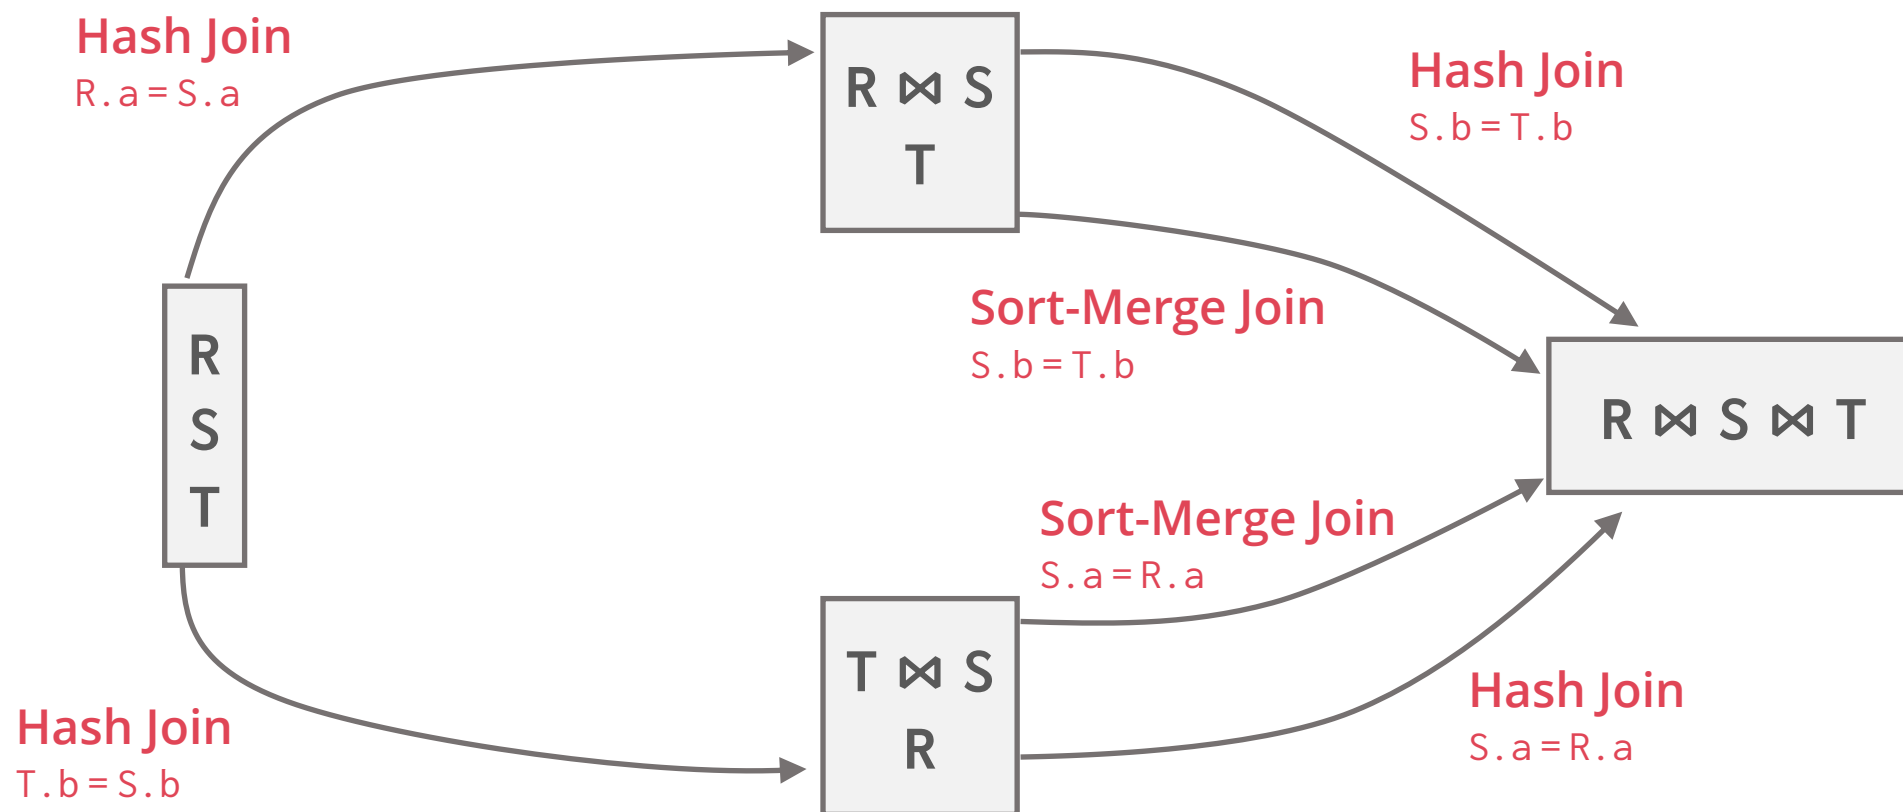


$R \bowtie S \bowtie T$

# EXAMPLE: DYNAMIC PROGRAMMING

**Pass #3** (best 3-relation plans):  
best 2-relation plans + one other relation


```
SELECT * FROM R, S, T
WHERE R.A = S.A
AND S.B = T.B
```



# EXAMPLE: DYNAMIC PROGRAMMING

**Pass #3** (best 3-relation plans):  
best 2-relation plans + one other relation

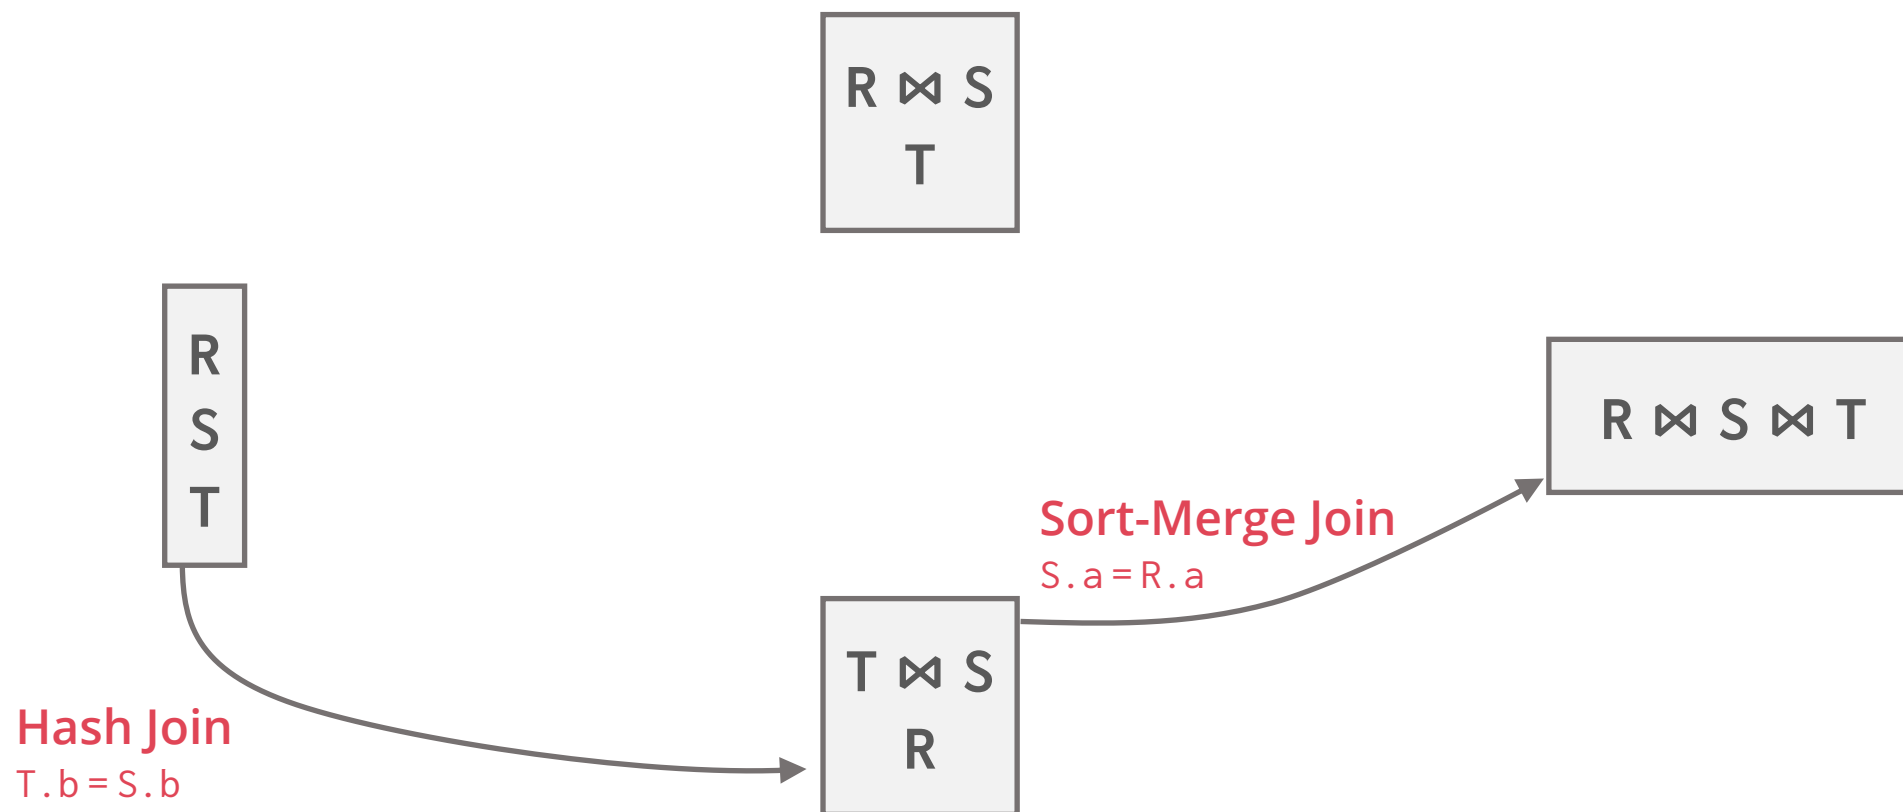
```
SELECT * FROM R, S, T
WHERE R.A = S.A
AND S.B = T.B
```



# EXAMPLE: DYNAMIC PROGRAMMING

**Pass #3** (best 3-relation plans):  
best 2-relation plans + one other relation

```
SELECT * FROM R, S, T
WHERE R.A = S.A
AND S.B = T.B
```



# DYNAMIC PROGRAMMING ALGORITHM

Function `find_join_tree_dp( $q(R_1, \dots, R_n)$ )`

```

for  $i = 1$  to  $n$  do
    optPlan[\{ $R_i$ \}] = access_plans( $R_i$ )
    prune_plans(optPlan[\{ $R_i$ \}])
for  $i = 2$  to  $n$  do
    foreach  $S \subseteq \{R_1, \dots, R_n\}$  such that  $|S| = i$  do
        optPlan[ $S$ ] =  $\emptyset$ 
        foreach  $T \subset S$  with  $T \neq \emptyset$  do
            optPlan[ $S$ ] = optPlan[ $S$ ]  $\cup$  possible_joins(optPlan[ $T$ ], optPlan[ $S \setminus T$ ])
        prune_plans(optPlan[ $S$ ])
return optPlan[\{ $R_1, \dots, R_n$ \}]

```


`possible_joins( $R, S$ )` enumerates the possible joins between  $R$  and  $S$  (e.g., NL join, SM join)

`prune_plans( $set$ )` discards all but the best plan from  $set$


# JOIN TREE SHAPES

Fundamental decision in IBM's **System R** (late 1970):


**Only consider left-deep join trees**



left-deep



bushy  
(everything else)



right-deep

# LEFT-DEEP JOIN TREES

DBMSs often prefer left-deep join trees

- The inner (rhs) relation always is a base relation

- Allows the use of index nested loops join


- Allows for **fully pipelined plans** where intermediate results are not written to temporary files

  - Should be factored into global cost calculation

  - Not all left-deep trees are fully pipelined (e.g., sort-merge join)

  - Pipelining requires **non-blocking** operators

Modern DBMSs may also consider non left-deep join trees



# MULTI-RELATION QUERY PLANNING

**System R-style join order enumeration**

Left-deep tree #1, Left-deep tree #2...

Eliminate plans with cross products immediately

**Enumerate the plans for each operator**

Hash, Sort-Merge, Nested Loop...

**Enumerate the access paths for each table**

Index #1, Index #2, Sequential scan...



# INTERESTING ORDERS

System R-style query optimisers also consider **interesting orders**

Sorting orders of the input tables that may be beneficial later in the query plan

e.g., for a sort-merge join, projection with duplicate removal, order-by clause

Determined by ORDER BY and GROUP BY clauses in the input query or join attributes of subsequent joins (to facilitate merging)

For each subset of relations, retain only:

Cheapest plan overall, plus

Cheapest plan for each **interesting order** of the tuples

# EXAMPLE: INTERESTING ORDERS

Consider the join query:

```
SELECT * FROM Orders O, Lineitem L
WHERE O.o_orderkey = L.l_orderkey
```

Orders has an unclustered index OK\_IDX on column o\_orderkey

Possible table access plans (1-relation plans) are:

Orders	<b>Full table scan</b> (arbitrary order). Cost: #pages(Orders) <b>Index scan</b> (sorted on o_orderkey). Cost = #pages(OK_IDX) + $k * \text{\#pages(Orders)}$ , $k$ is #index entries/page
Lineitem	<b>Full table scan</b> . Cost = #pages(Lineitem)

# EXAMPLE: INTERESTING ORDERS

The full table scan is the cheapest access method for both tables

Although more expensive than a full table scan, the use of the index (**order enforcement**) may pay off later on in the overall plan

- An index scan of OK\_IDX yields the scan output in o\_orderkey order

- This is beneficial for sort-merge join as we need to sort only Lineitem

- This could turn out to be the best 2-relation plan!

System R-style optimisers would keep both full table scan and index scan as best 1-relation plans for Orders

# CONCLUSION

Query optimization is an important task in a relational DBMS

Explores a set of alternative plans

- Must prune search space; typically, left-deep plans only

- Uses dynamic programming for join orderings

Must estimate cost of each plan that is considered

- Must estimate the size of result and cost for each plan node

Query optimizer is the most complex part of database systems!