



THE UNIVERSITY  
*of* EDINBURGH

# Advanced Databases

Spring 2020

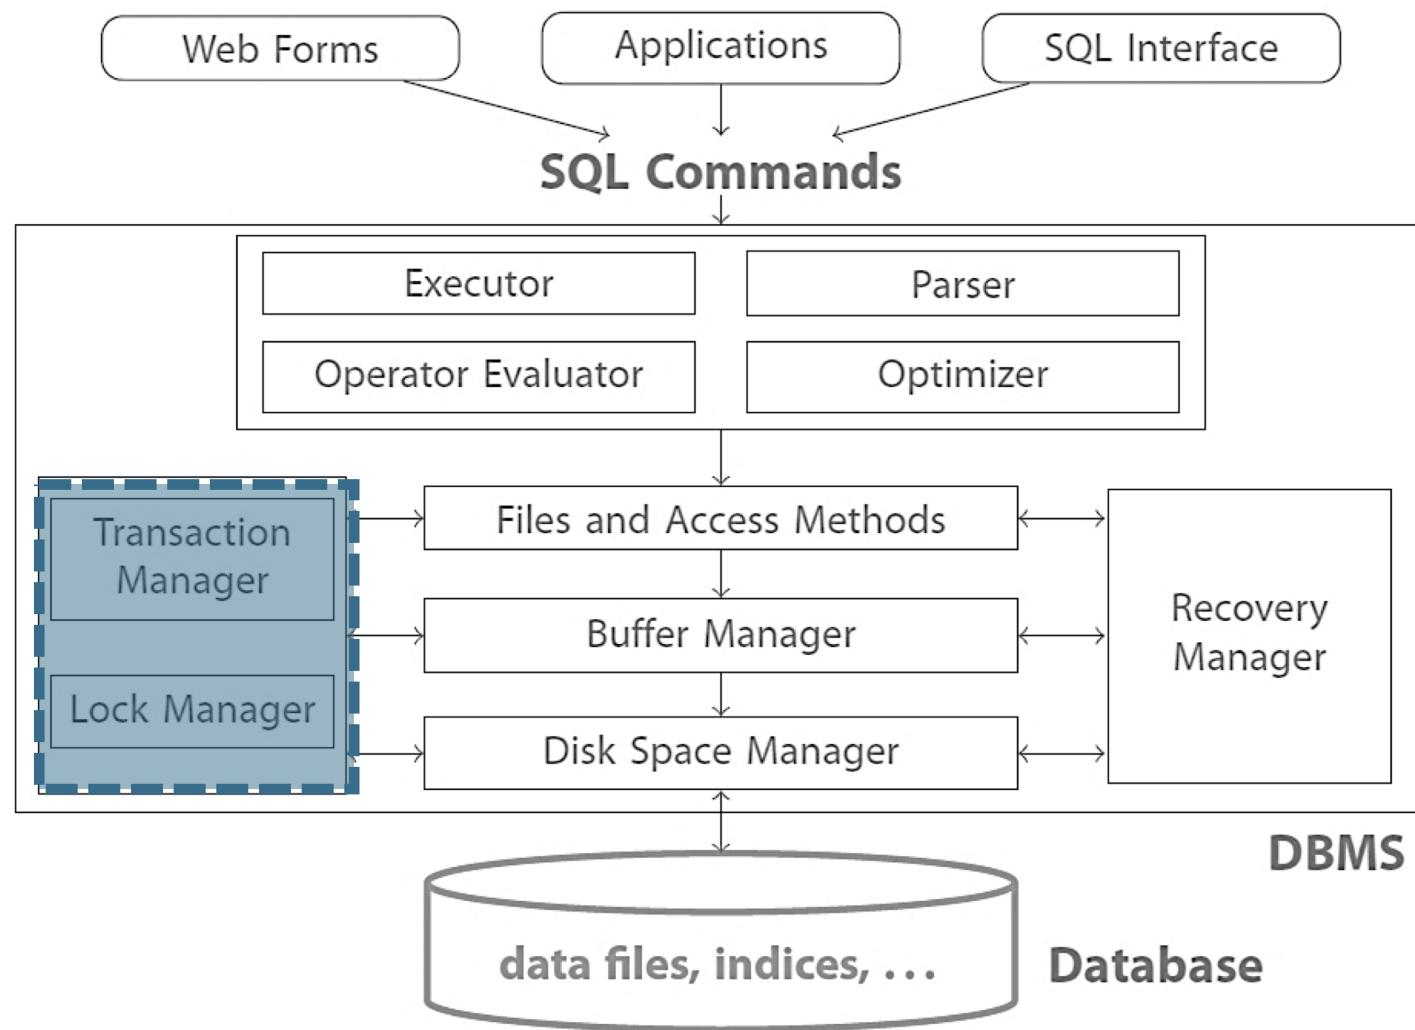
---

Lecture #12:

## Transactions

Milos Nikolic

# DATABASE ARCHITECTURE



# MOTIVATION

We both change the same record in a table at the same time.

How to avoid race condition?

You transfer £100 between bank accounts but there is a power failure.

What is the correct database state?



Concurrency Control



Recovery


Both concurrency control and recovery are based on a concept of transactions with **ACID** properties

# TRANSACTIONS

A **transaction** is the execution of a sequence of operations (e.g., SQL queries) on a shared database to perform some higher-level function

Basic unit of change in a DBMS

**Partial transactions are not allowed!**



# USER PERSPECTIVE: TRANSACTIONS

Transaction (abbr. txn) = **group of operations** the user wants the DBMS to treat “as one”

A new transaction starts with the **BEGIN** command

The transaction stops with either **COMMIT** or **ABORT (ROLLBACK)**

- If commit, all changes are saved

- If abort, all changes are undone (as if the txn never executed at all)

- Abort can be either self-inflicted or caused by DBMS

# TRANSACTION EXAMPLE

Transfer £100 from Checking to Savings account of user 1904

**BEGIN**

*// check if Checking balance > 100*

**UPDATE** Accounts

**SET** balance = balance - 100

**WHERE** customer\_id = 1904

**AND** account\_type = 'Checking';

**UPDATE** Accounts

**SET** balance = balance + 100

**WHERE** customer\_id = 1904

**AND** account\_type = 'Savings';

**COMMIT**

Consistent DB

Temporary inconsistent DB

Consistent DB

# TRANSACTION EXAMPLE

Transfer £100 from Checking to Savings account of user 1904

```
BEGIN
  // check if Checking balance > 100
  UPDATE Accounts
    SET balance = balance - 100
  WHERE customer_id = 1904
    AND account_type = 'Checking';
  UPDATE Accounts
    SET balance = balance + 100
  WHERE customer_id = 1904
    AND account_type = 'Savings';
COMMIT
```

How to check if balance > 100?

Outside DBMS using another language

E.g., in Java or PHP code

Inside DBMS using **stored procedures**  
expressed in PL/SQL

PL/SQL = SQL + procedural constructs such as  
if-then-else, loops, variables, functions...

# DATABASE PERSPECTIVE

A transaction may carry out many operations on the data retrieved from the database

However, the DBMS is only concerned about what data is read/written from/to the database

Changes to the “outside world” are beyond scope of the DBMS



# TRANSACTIONS: FORMAL DEFINITION

**Database** = fixed set of named data objects ( $A, B, C, \dots$ )

Transactions access object  $A$  using read  $A$  and write  $A$ , for short  $R(A)$  and  $W(A)$

In a relational DBMS, an object can be an attribute, record, or table

**Transaction** = sequence of read and write operations

$T = \langle R(A), W(A), W(B), \dots \rangle$

DBMS's abstract view of a user program

# STRAWMAN EXECUTION

Execute each txn **one-by-one** (serial order) as they arrive in the DBMS

One and only one txn can be running at the same time in the DBMS

Before a txn starts, **copy** the entire database to a new file and make all changes to that file

If the txn completes successfully, overwrite the original file with the new one

If the txn fails, just remove the dirty copy

SQLite executes transactions in serial order

# CONCURRENT EXECUTION


A (potentially) better approach is to allow **concurrent execution** of independent transactions

*Why do we want that?*

Better resource utilization (CPU) and throughput

Decreased response times to users

But we also would like **correctness** and **fairness**



# TRANSACTION GUARANTEES: ACID

**A**tomicity: All actions in the txn happen, or none happen

*"all or nothing"*

**C**onsistency: If each txn is consistent and the DB starts consistent, then it ends up consistent

*"it looks correct to me"*

**I**solation: Execution of one txn is isolated from that of other txns

*"as if alone"*

**D**urability: If a txn commits, its effects persist

*"survive failures"*

# ACID PROPERTIES: ATOMICITY

Two possible outcome of executing a transaction:

- Commit after completing all actions

- Abort (or be aborted by the DBMS) after executing some actions

DBMS guarantees that transactions are **atomic**

From user's point of view: a transaction always either executes all its actions or executes no actions at all

Example:

Take £100 from account A, but then a power failure happens before crediting account B

*When the DBMS comes back online, what should be the correct state of the database?*

# MECHANISMS FOR ENSURING ATOMICITY

## Approach #1: Logging

DBMS logs all actions so that it can undo the actions of aborted transactions

Write-ahead logging is used by almost all modern database systems

Audit trail & efficiency reasons

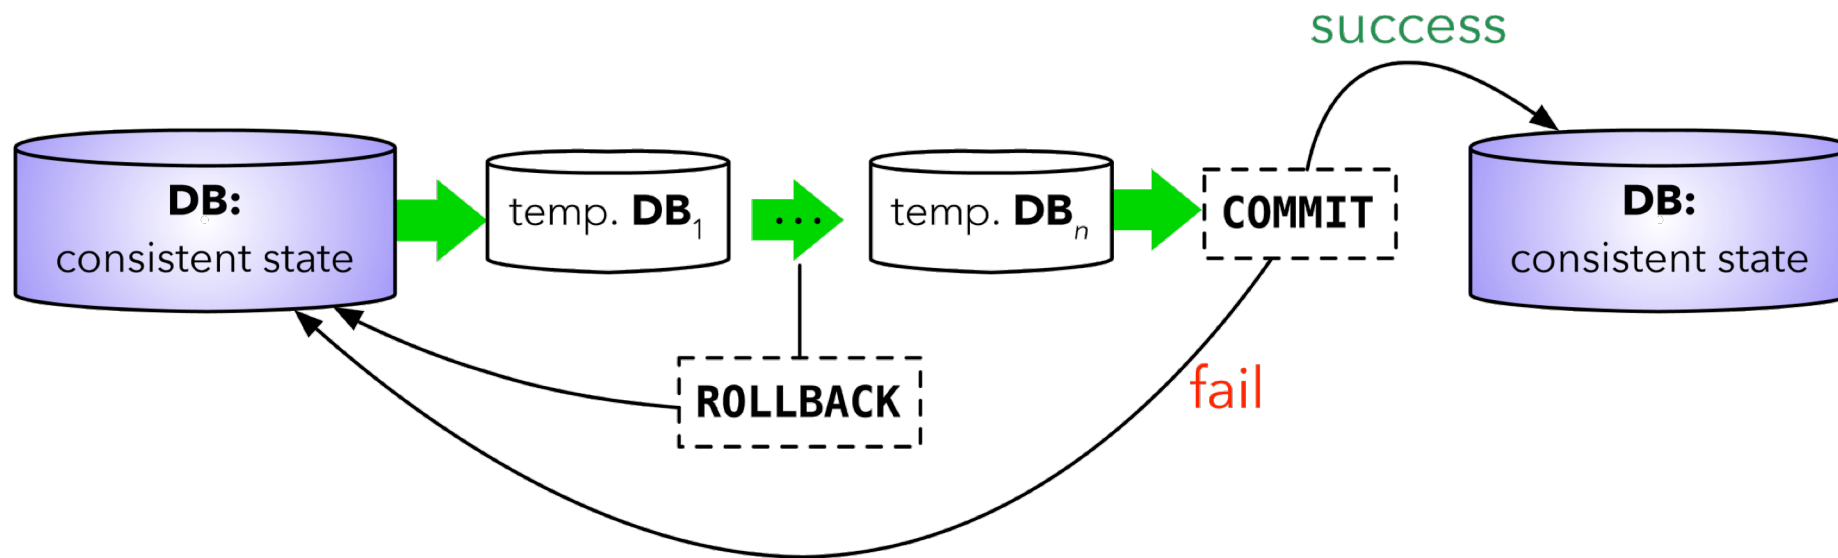
## Approach #2: Shadow Paging (copy-on-write)

DBMS makes copies of pages and transactions make changes to those copies

Only when the transaction commits is the page made visible to others

Few database systems do this (CouchDB, LMDB)

# ACID PROPERTIES: CONSISTENCY



## Database consistency

The database accurately models the real world and follows integrity constraints

Transactions in the future see the effects of transactions committed in the past

## Transaction consistency

If the database is consistent before the txn starts (running alone), it will be also consistent after

Transaction consistency is the application's responsibility!

# ACID PROPERTIES: ISOLATION

Users submit transactions, and each transaction executes as if it was running alone

DBMS achieves concurrency by interleaving actions (read/writes of database objects) of various transactions

*How do we achieve this?*



# MECHANISMS FOR ENSURING ISOLATION

A **concurrency control** protocol is how the DBMS decides the proper interleaving of operations from multiple transactions

Two main categories:

**Pessimistic:** Don't let problems arise in the first place

**Optimistic:** Assume conflicts are rare, deal with them after they happen

# EXAMPLE

Assume at first accounts **A** and **B** each have £1000

**T<sub>1</sub>** transfers £100 from **A** to **B**

**T<sub>2</sub>** credits both accounts with 6% interest

**T<sub>1</sub>**

**BEGIN**

$A = A - 100$

$B = B + 100$

**END**

**T<sub>2</sub>**

**BEGIN**

$A = A * 1.06$

$B = B * 1.06$

**END**

# EXAMPLE

Assume at first accounts **A** and **B** each have £1000

*What are the possible outcomes of running **T<sub>1</sub>** and **T<sub>2</sub>**?*

**T<sub>1</sub>**

**BEGIN**

A = A - 100

B = B + 100

**END**

**T<sub>2</sub>**

**BEGIN**

A = A \* 1.06

B = B \* 1.06

**END**

# EXAMPLE

Assume at first accounts **A** and **B** each have £1000

*What are the possible outcomes of running **T<sub>1</sub>** and **T<sub>2</sub>**?*

Many! But **A+B** should be  **$2000 * 1.06 = 2120$**

There is no guarantee that **T<sub>1</sub>** will execute before **T<sub>2</sub>** or vice versa, if both are submitted together

But the net effect must be equivalent to these two transactions running **serially** in some order

# EXAMPLE

Assume at first accounts **A** and **B** each have £1000

Legal outcomes:

$$\mathbf{A} = 954, \mathbf{B} = 1166 \rightarrow \mathbf{A+B = 2120}$$

$$\mathbf{A} = 960, \mathbf{B} = 1160 \rightarrow \mathbf{A+B = 2120}$$

The outcome depends on whether **T<sub>1</sub>** executes before **T<sub>2</sub>** or vice versa

**T<sub>1</sub>**

**BEGIN**

A = A - 100

B = B + 100

**END**

**T<sub>2</sub>**


**BEGIN**

A = A \* 1.06

B = B \* 1.06

**END**

# EXAMPLE: SERIAL EXECUTION



# INTERLEAVING TRANSACTIONS

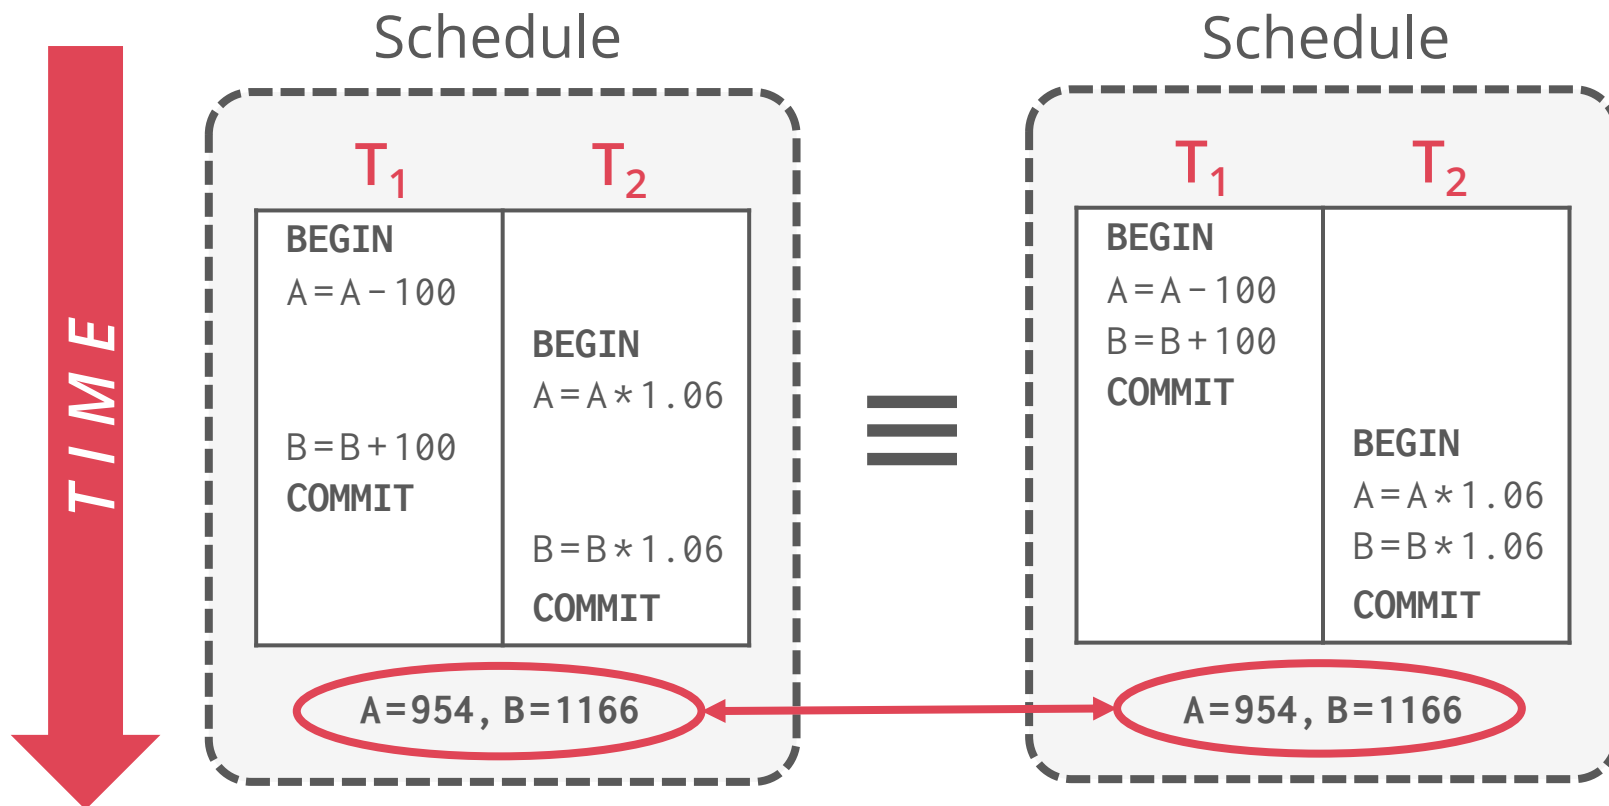
We interleave transactions to maximise concurrency

- Slow disk I/O or network

- Multi-core CPU

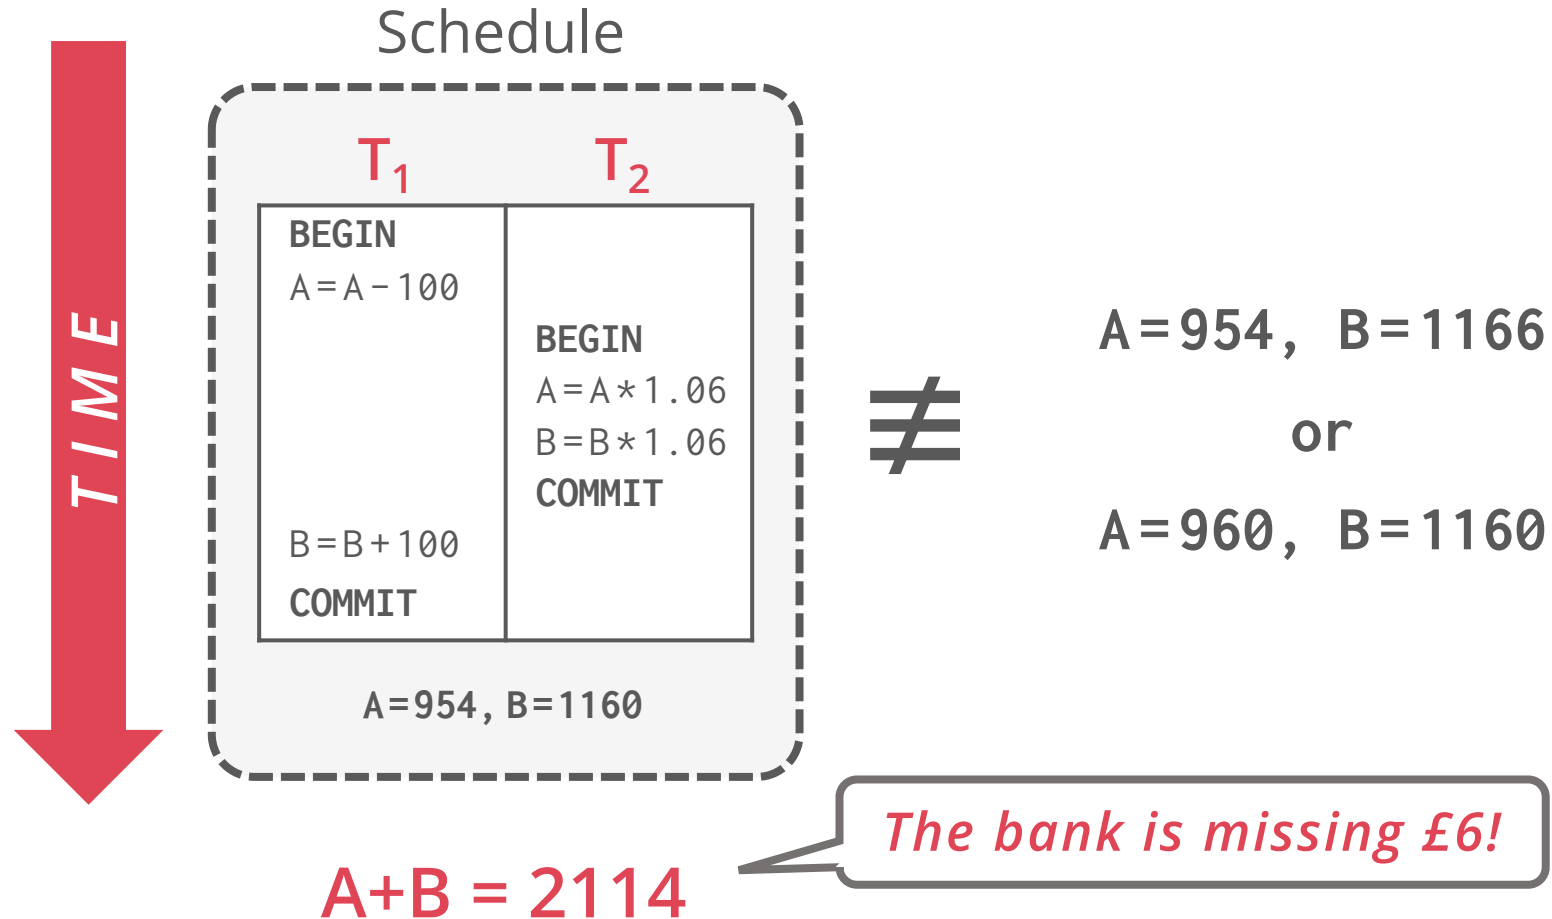
When one txn stalls because of a resource (e.g., page fault), another txn can continue executing and make forward progress

# EXAMPLE: INTERLEAVING (GOOD)

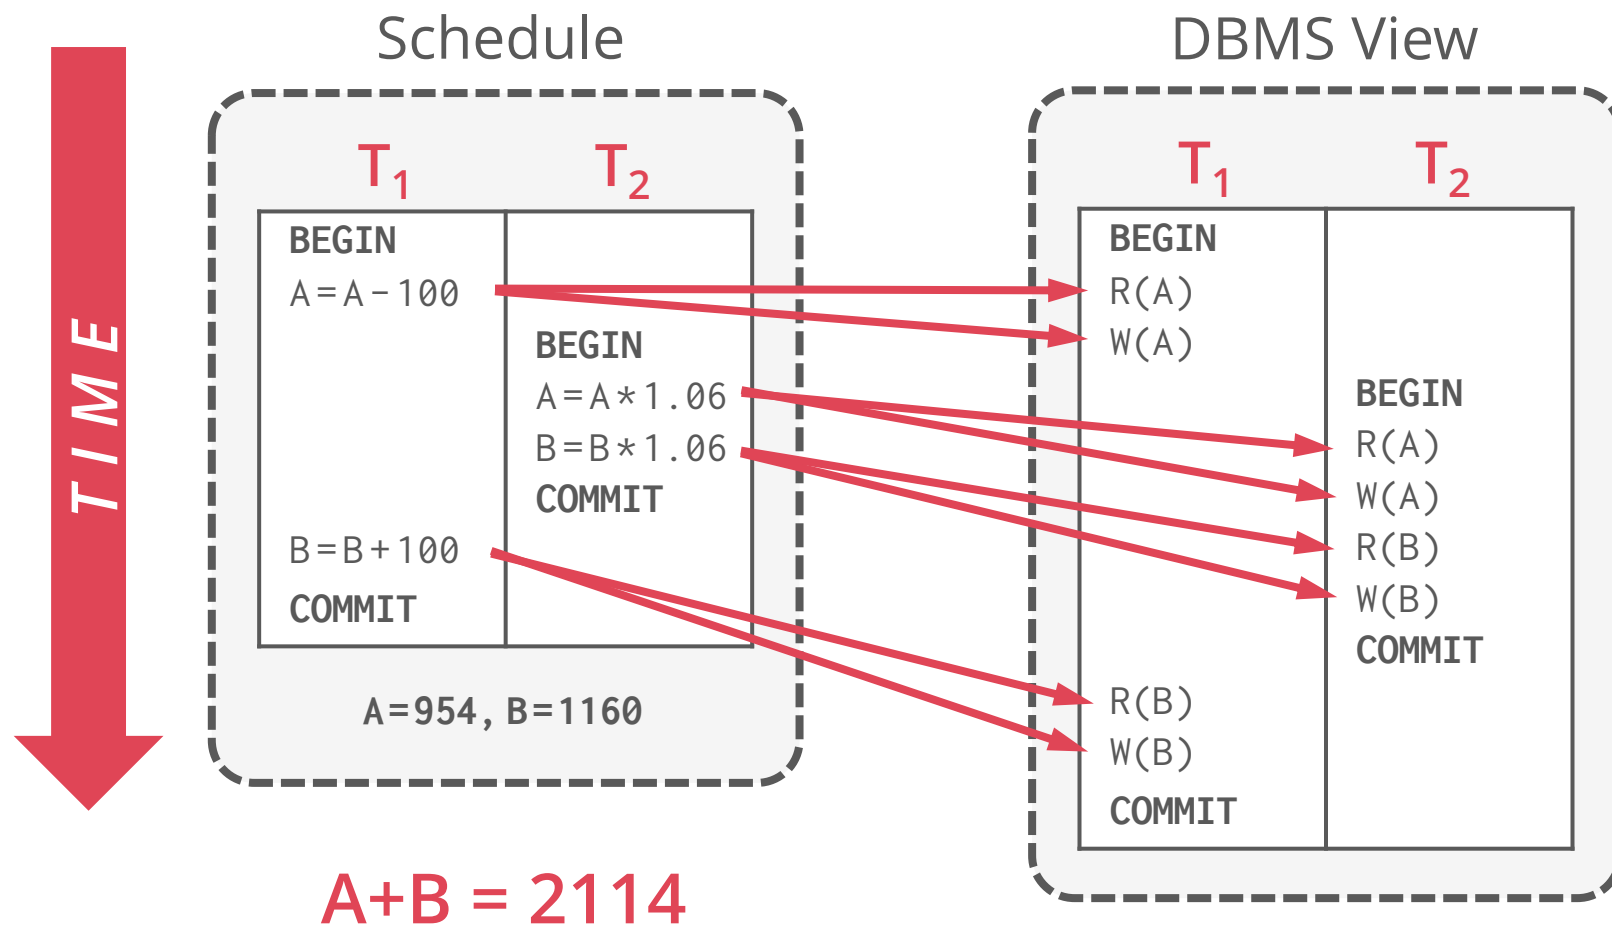




# EXAMPLE: INTERLEAVING (BAD)



# EXAMPLE: INTERLEAVING (BAD)



# CORRECTNESS

*How do we judge whether a schedule is correct?*

If the schedule is **equivalent** to some **serial execution**

**Schedule  $S$**  for a set of transactions  $\{ T_1, \dots, T_n \}$

$S$  contains *all* steps of all transactions and order among steps in each  $T_i$  is *preserved*

$S = \langle (T_1, \text{read } B), (T_2, \text{read } A), (T_2, \text{write } B), (T_1, \text{write } A) \rangle$

for short,  $S = \langle R_1(B), R_2(A), W_2(B), W_1(A) \rangle$

# FORMAL PROPERTIES OF SCHEDULES


## Equivalent schedules

For any database state, the effect of executing the first schedule is identical to the effect of executing the second schedule

Does not matter what the higher-level operations are!

## Serial schedule (no concurrency)

A schedule that does not interleave the actions of different transactions



# FORMAL PROPERTIES OF SCHEDULES

## Serializable schedule

A schedule that is equivalent to some serial execution of the transactions

If each transaction preserves consistency, every serializable schedule preserves consistency

## Serializability

Less intuitive notion of correctness compared to transaction initiation time or commit order

But it provides the DBMS with flexibility in scheduling operations

More flexibility means **better parallelism**

# CONFLICTING OPERATIONS

We need a formal notion of equivalence that can be implemented efficiently based on the notion of “conflicting” operations

Two operations **conflict** if

- They are by different transactions

- They are on the same object and at least one of them is a write

**Interleaved execution anomalies:**

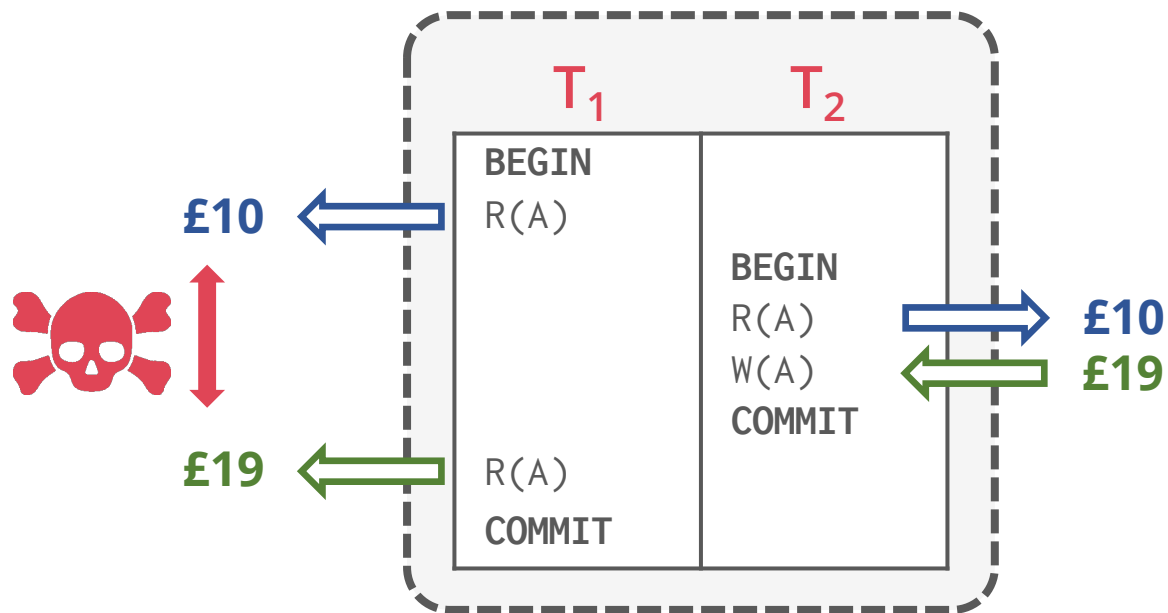
- Read-Write conflicts (R-W)

- Write-Read conflicts (W-R)

- Write-Write conflicts (W-W)


# READ-WRITE CONFLICTS

## Unrepeatable Reads



# WRITE-READ CONFLICTS

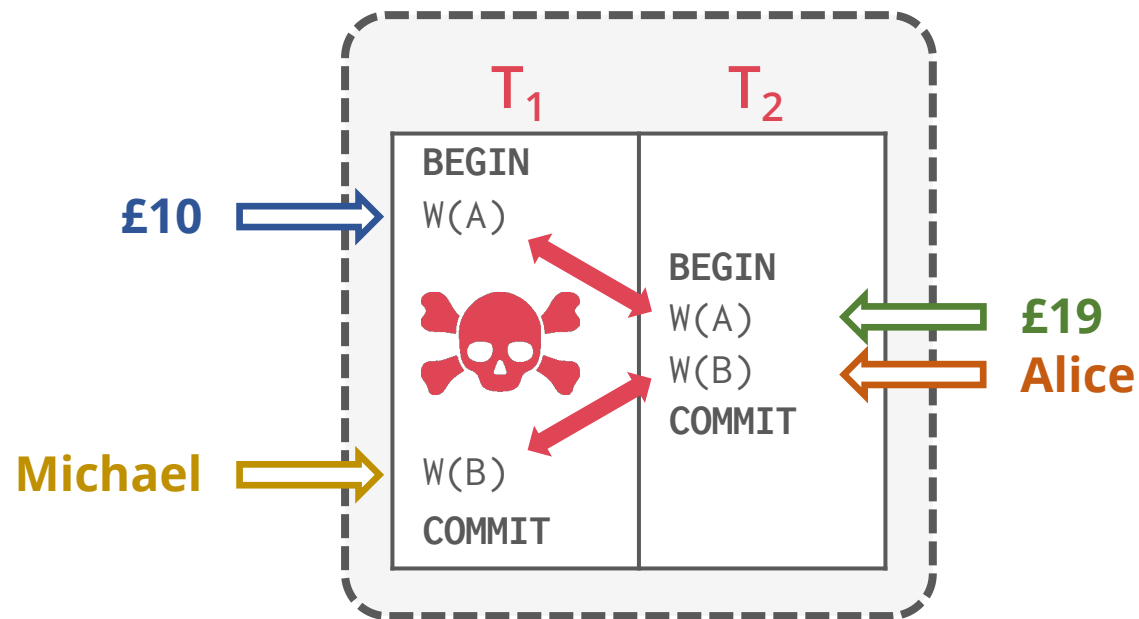
Reading Uncommitted Data ("Dirty Reads")





# WRITE-WRITE CONFLICTS

Overwriting Uncommitted Data ("Lost Update")



# FORMAL PROPERTIES OF SCHEDULES

Given these conflicts, we can now understand what it means for a schedule to be serializable

This is to check whether schedules are correct

This is not how to generate a correct schedule

There are levels of serializability

Conflict Serializability

*Most DBMS try to support this*

View Serializability

*No DBMS supports this*

# CONFLICT SERIALIZABLE SCHEDULES

Two schedules are **conflict equivalent** iff

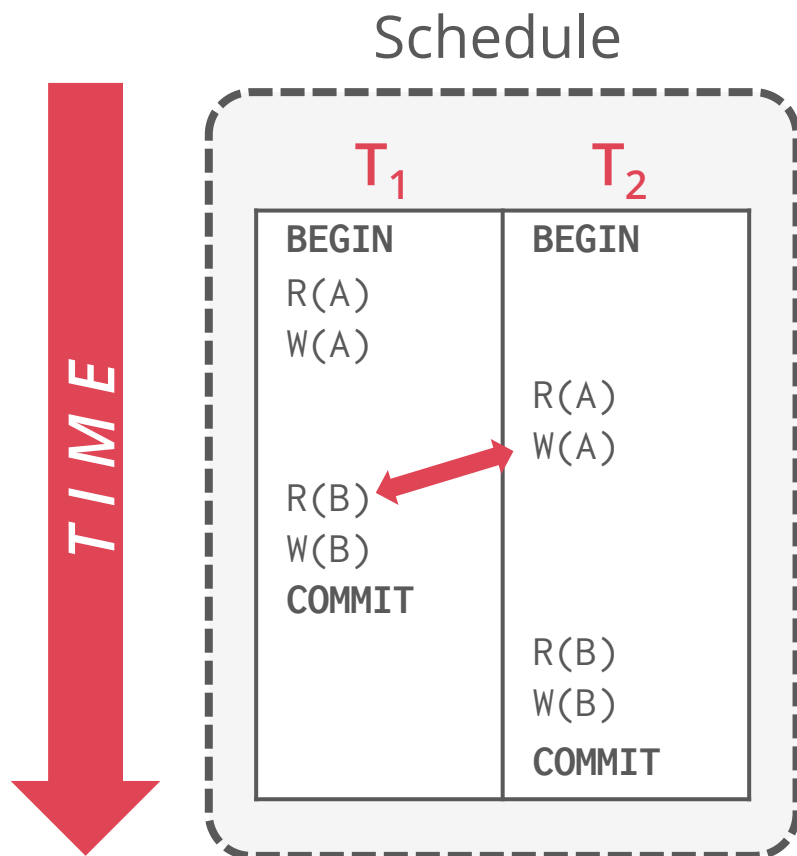
- They involve the same actions of the same transactions

- Every pair of conflicting actions is ordered in the same way

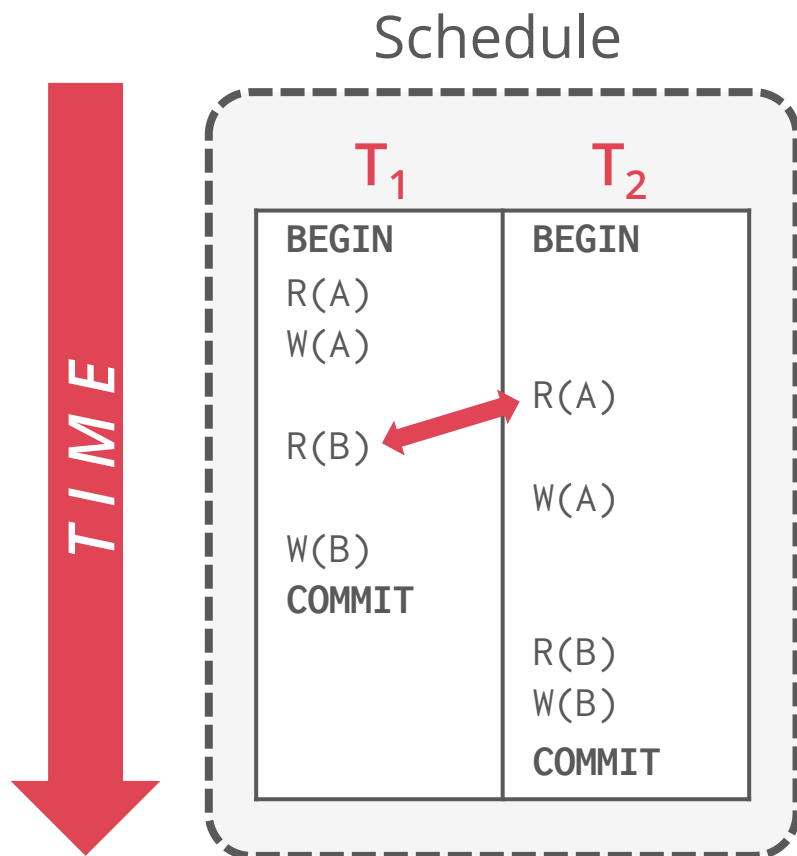
Schedule **S** is **conflict serializable** if **S** is conflict equivalent to some serial schedule

*Intuition: Schedule **S** is conflict serializable if you can transform **S** into a serial schedule by swapping consecutive non-conflicting operations of different txns*

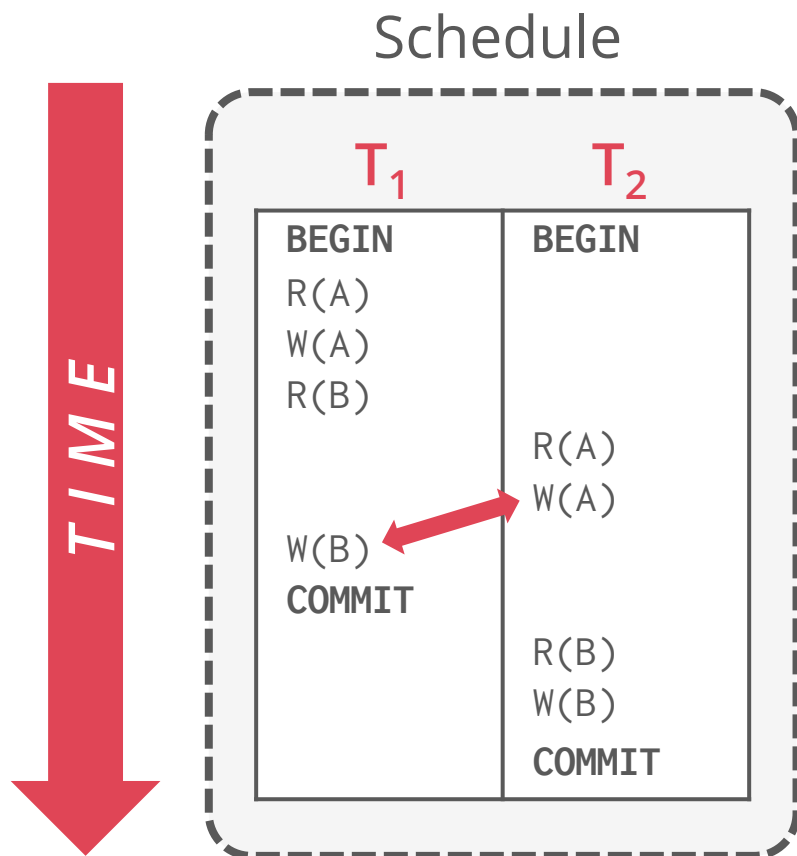
# CONFLICT SERIALIZABILITY: INTUITION



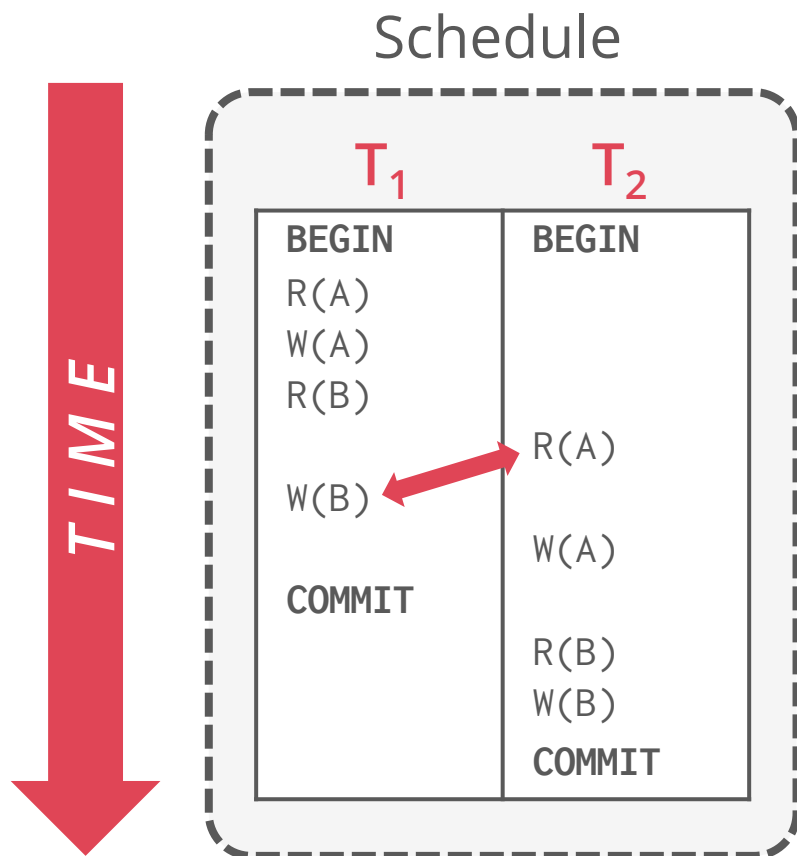
# CONFLICT SERIALIZABILITY: INTUITION



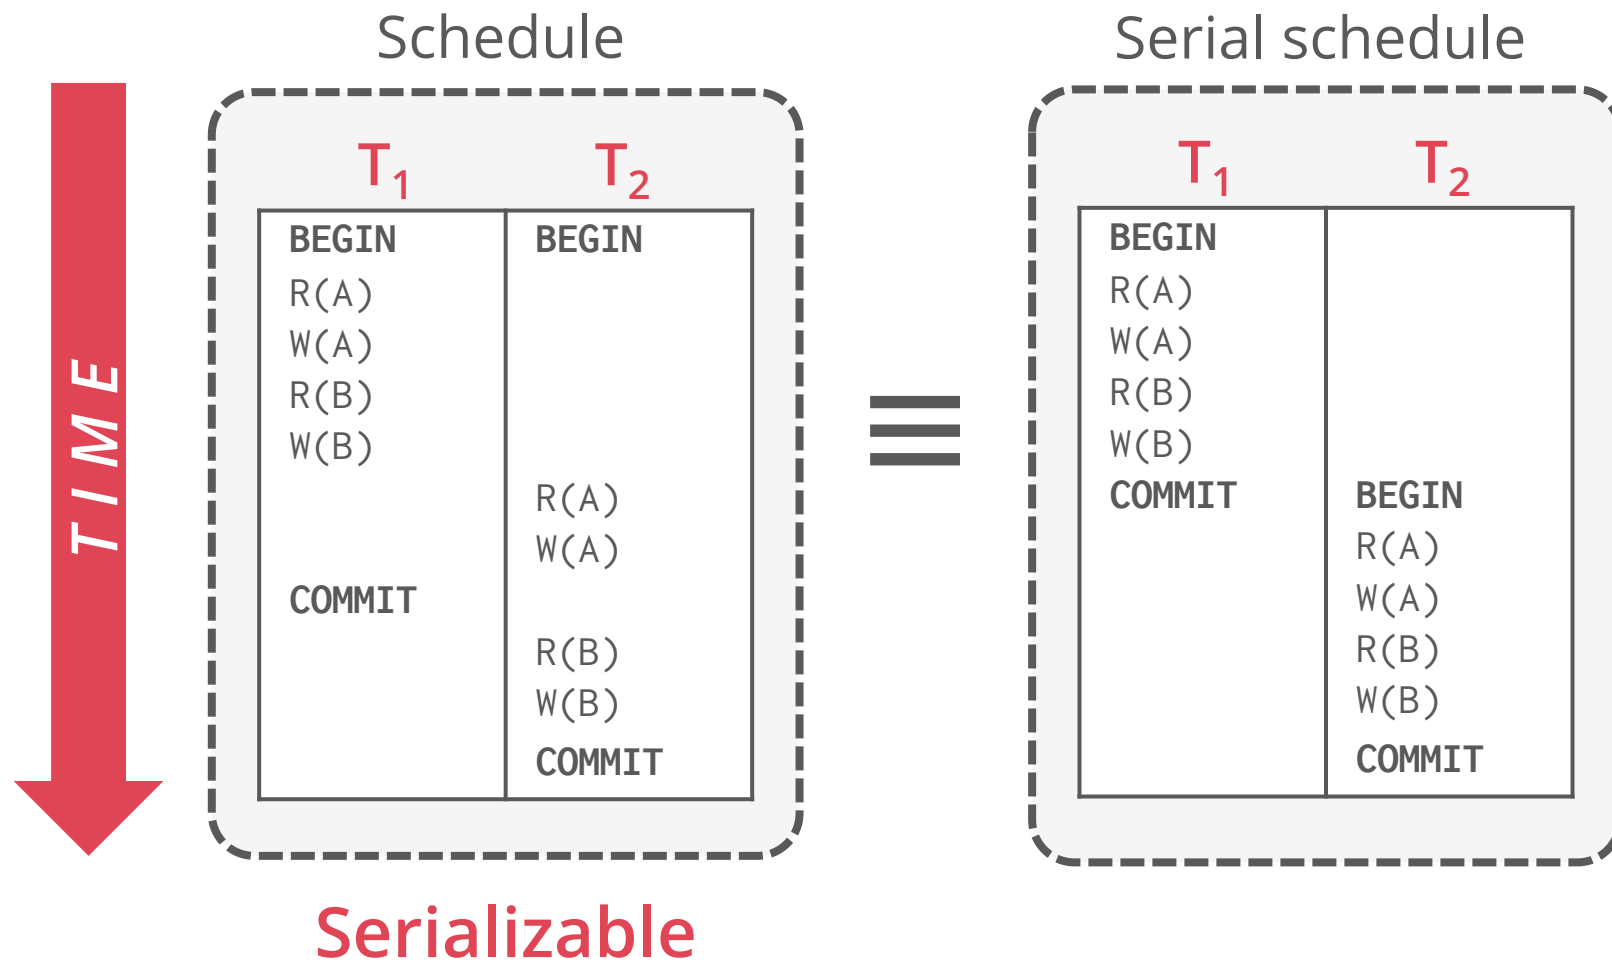
# CONFLICT SERIALIZABILITY: INTUITION



# CONFLICT SERIALIZABILITY: INTUITION

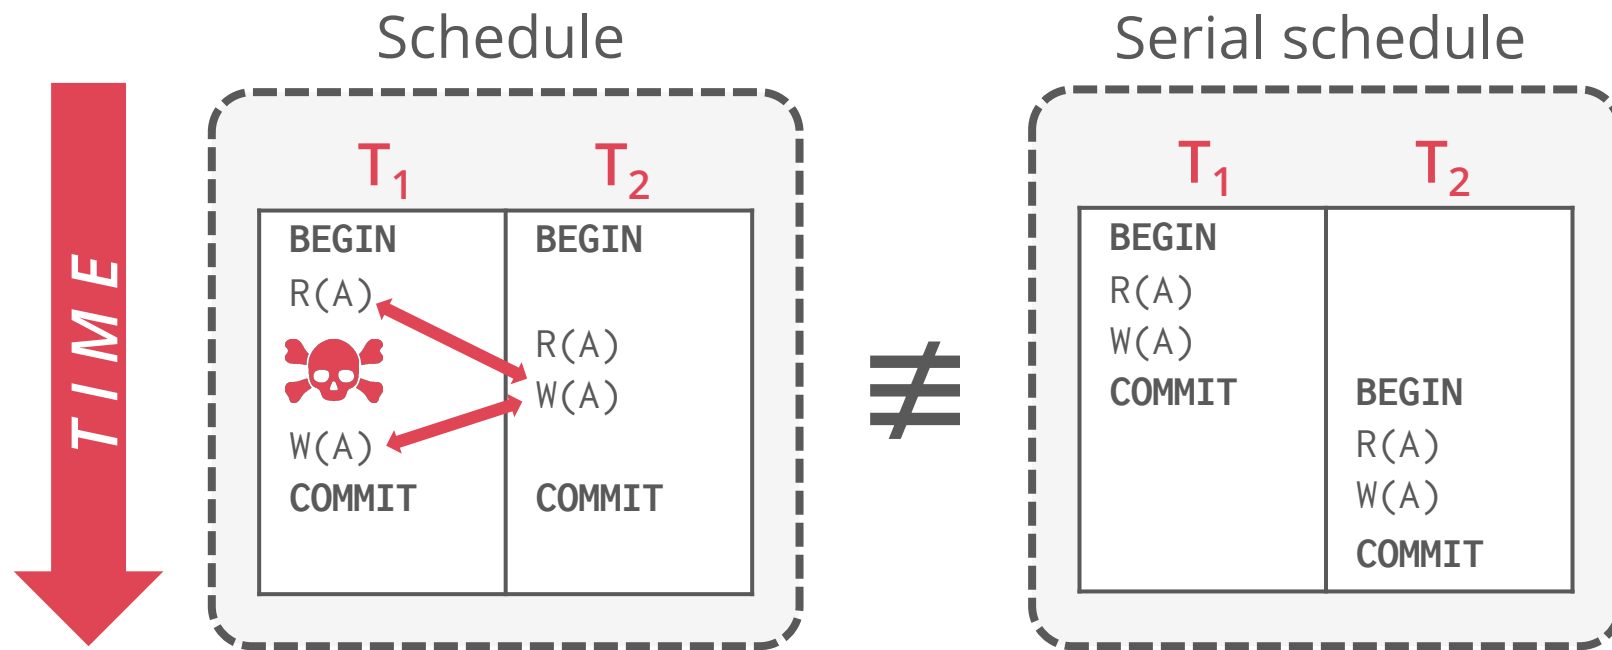


# CONFLICT SERIALIZABILITY: INTUITION





# CONFLICT SERIALIZABILITY: INTUITION



**Not conflict-serializable**

# SERIALIZABILITY

Swapping operations is easy when there are only two txns in the schedule. It's cumbersome when there are many txns

*Are there any faster algorithms to figure this out other than transposing operations?*

# DEPENDENCY GRAPHS


**Dependency graph** for a schedule

One node per transaction

Edge from  $T_i$  to  $T_j$  if:

Operation  $O_i$  of  $T_i$  conflicts with an operation  $O_j$  of  $T_j$  and  
 $O_i$  appears earlier in the schedule than  $O_j$

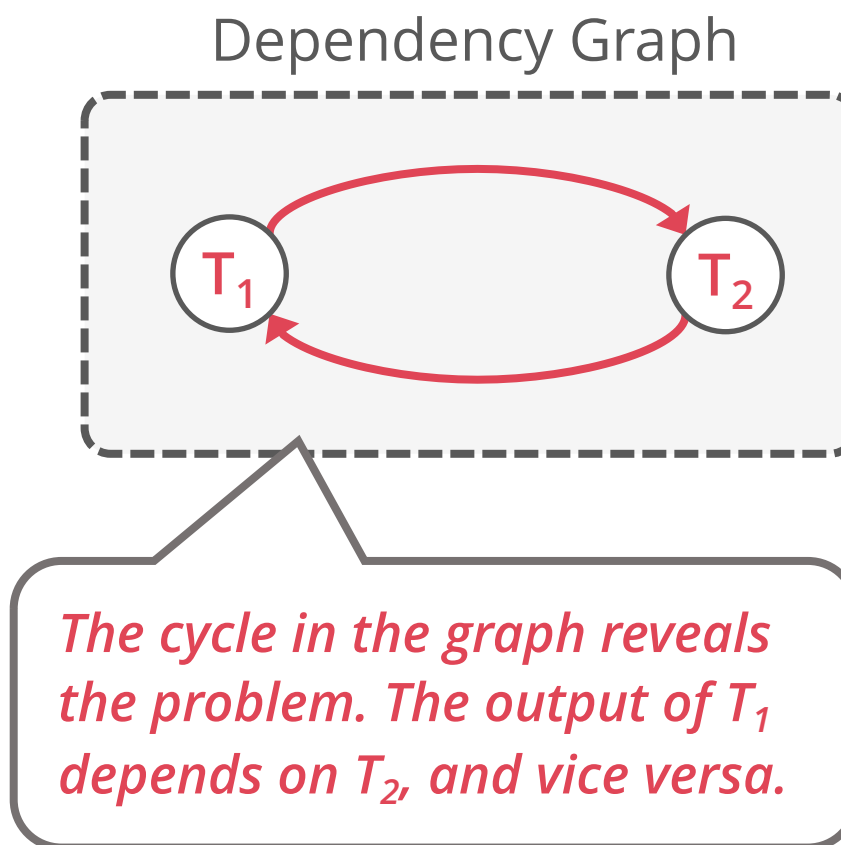
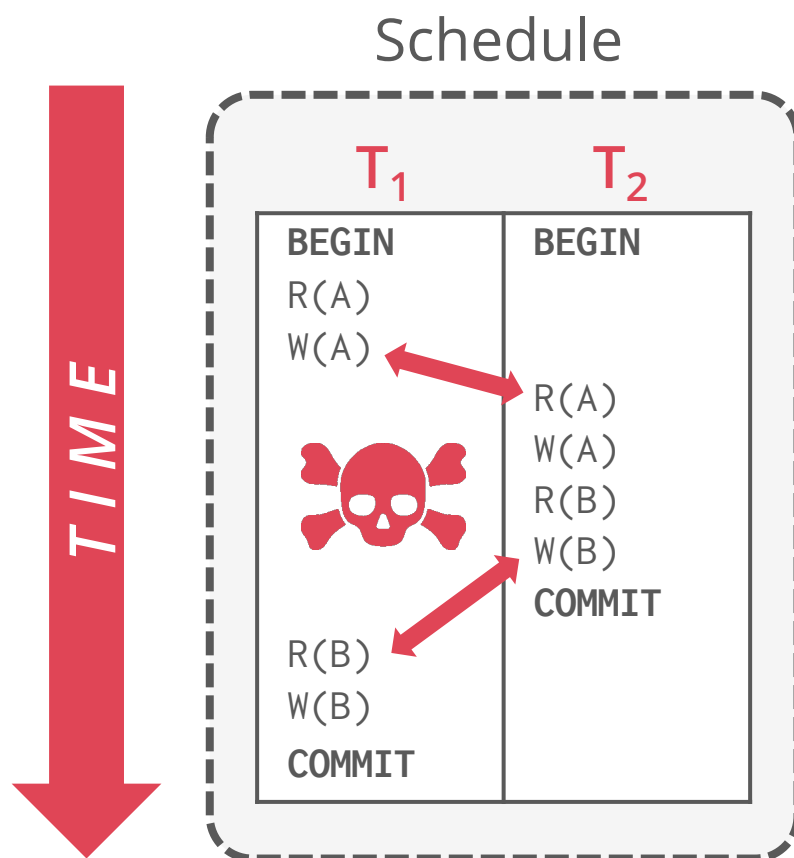
Also known as a **conflict graph** or **precedence graph**




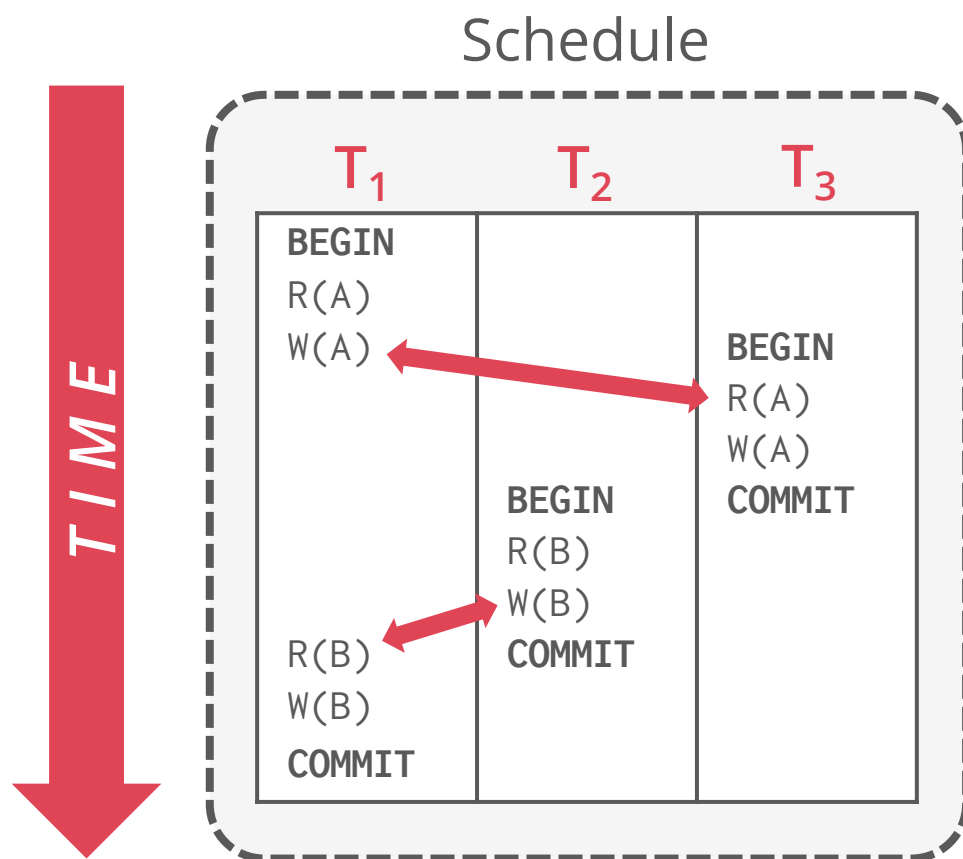
A schedule is conflict-serializable if and only if its dependency graph is acyclic.

Equivalent **serial schedule** can be obtained by sorting the graph **topologically**

# EXAMPLE #1



# EXAMPLE #2 - THREESOME



*Is this equivalent to a serial schedule?*

Yes, ( $T_2, T_1, T_3$ )

Notice that  $T_3$  should go after  $T_2$  although  $T_3$  starts before  $T_2$ !

# VIEW SERIALIZABILITY

Alternative (weaker) notion of serializability


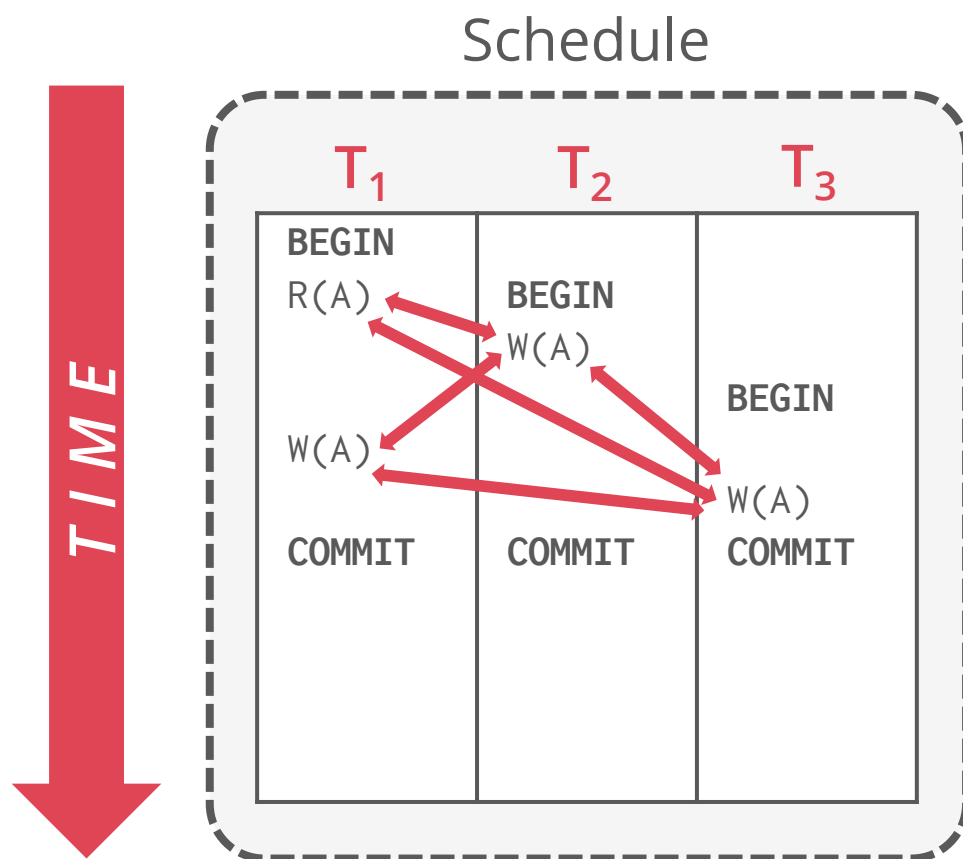
Schedule  $S_1$  and  $S_2$  are **view equivalent** iff

If  $T_1$  reads initial value of  $A$  in  $S_1$ , then  $T_1$  also reads initial value of  $A$  in  $S_2$

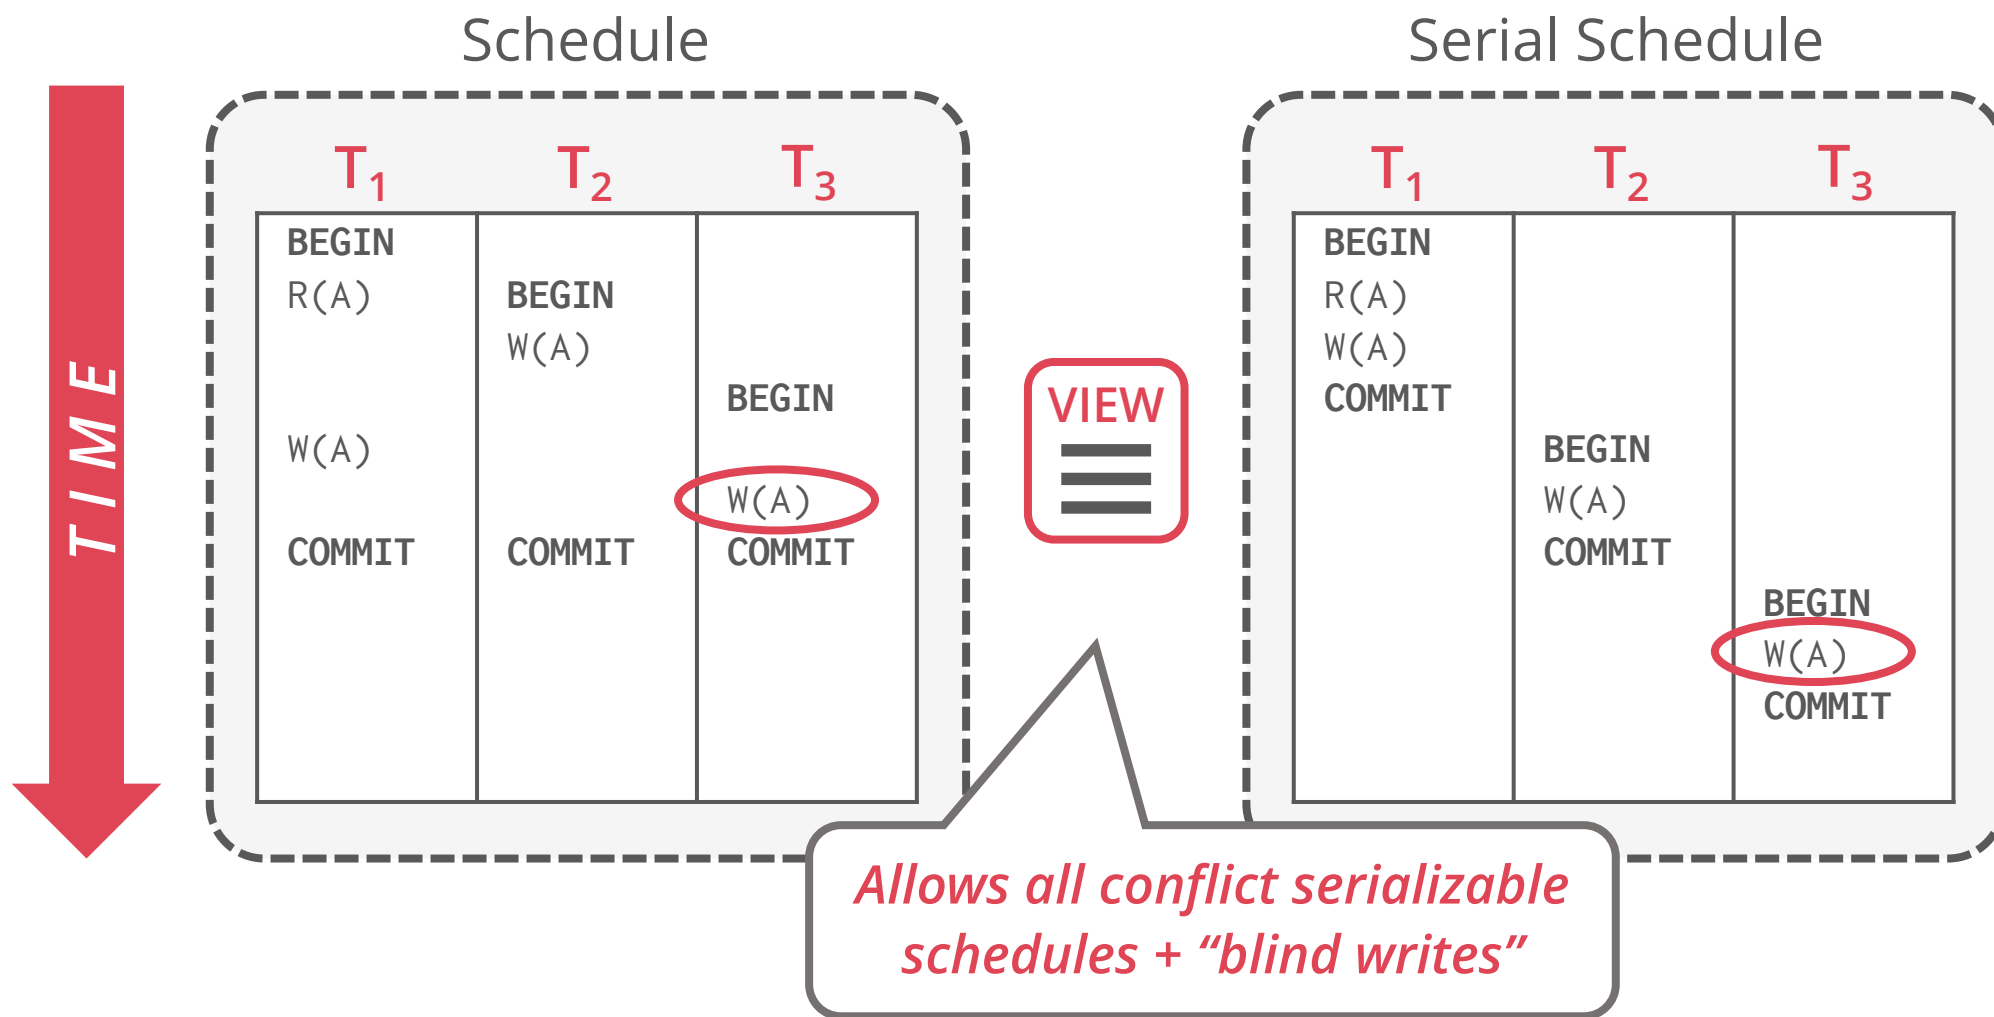
If  $T_1$  reads value of  $A$  written by  $T_2$  in  $S_1$ , then  $T_1$  also reads value of  $A$  written by  $T_2$  in  $S_2$

If  $T_1$  writes final value of  $A$  in  $S_1$ , then  $T_1$  also writes final value of  $A$  in  $S_2$

# VIEW SERIALIZABILITY



# VIEW SERIALIZABILITY





# SERIALIZABILITY

## Conflict serializability

Can enforced efficiently

All DBMSs support it

## View serializability

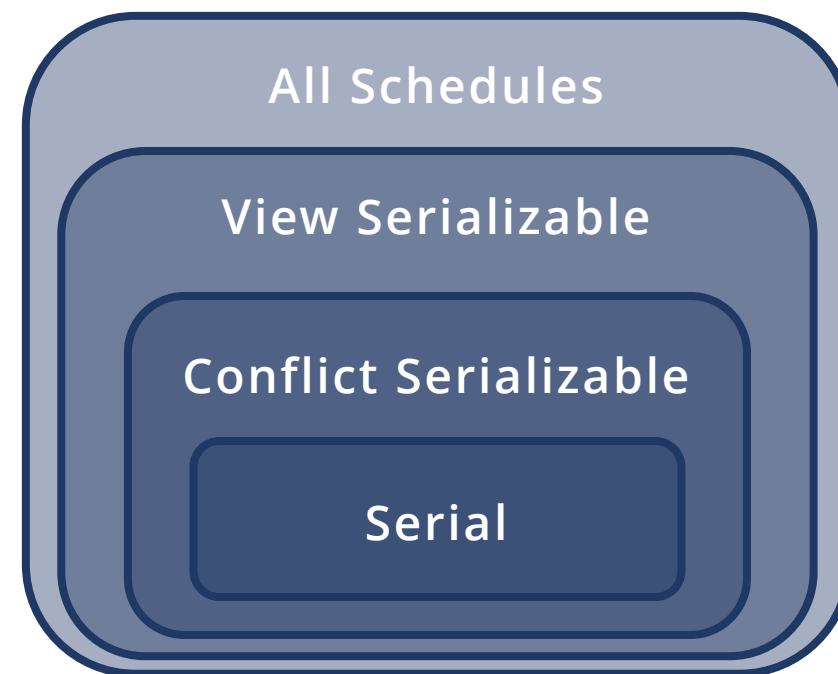
Admits (slightly) more schedules than CS

But it is difficult to enforce efficiently

No DBMS supports it

Neither definition allows all “serializable” schedules

They do not understand the meaning of the operations or the data



# ACID PROPERTIES: DURABILITY

All of the changes of committed transactions must be persistent

- No torn updates

- No changes from failed transactions

The DBMS uses either logging or shadow paging to ensure that all changes are durable

More about logging next week

# CONCLUSION

## ACID Transactions

**A**tomicity: All or nothing

**C**onsistency: Only valid data

**I**solation: No interference

**D**urability: Committed data persists

## Serializability

Serializable schedules

Conflict & view serializability

Checking for conflict serializability

Concurrency control and recovery are among the most important functions provided by a DBMS