



THE UNIVERSITY
of EDINBURGH

Advanced Databases

Spring 2020

Lecture #13:

Two-Phase Locking

Milos Nikolic

QUERY SCHEDULER

How to guarantee only serializable schedules in DBMS?

Problem: user does not need to specify the full transaction at once

Goal: build a query scheduler that always emits serializable schedules

Pessimistic (locking)

Use locks to protect database objects

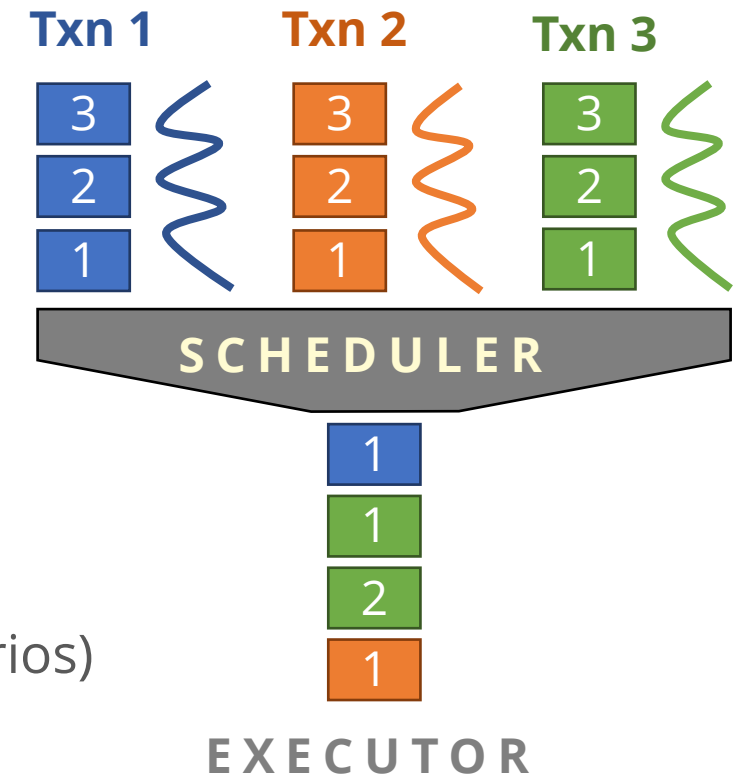
Standard approach if conflicts are frequent

Optimistic (versioning)

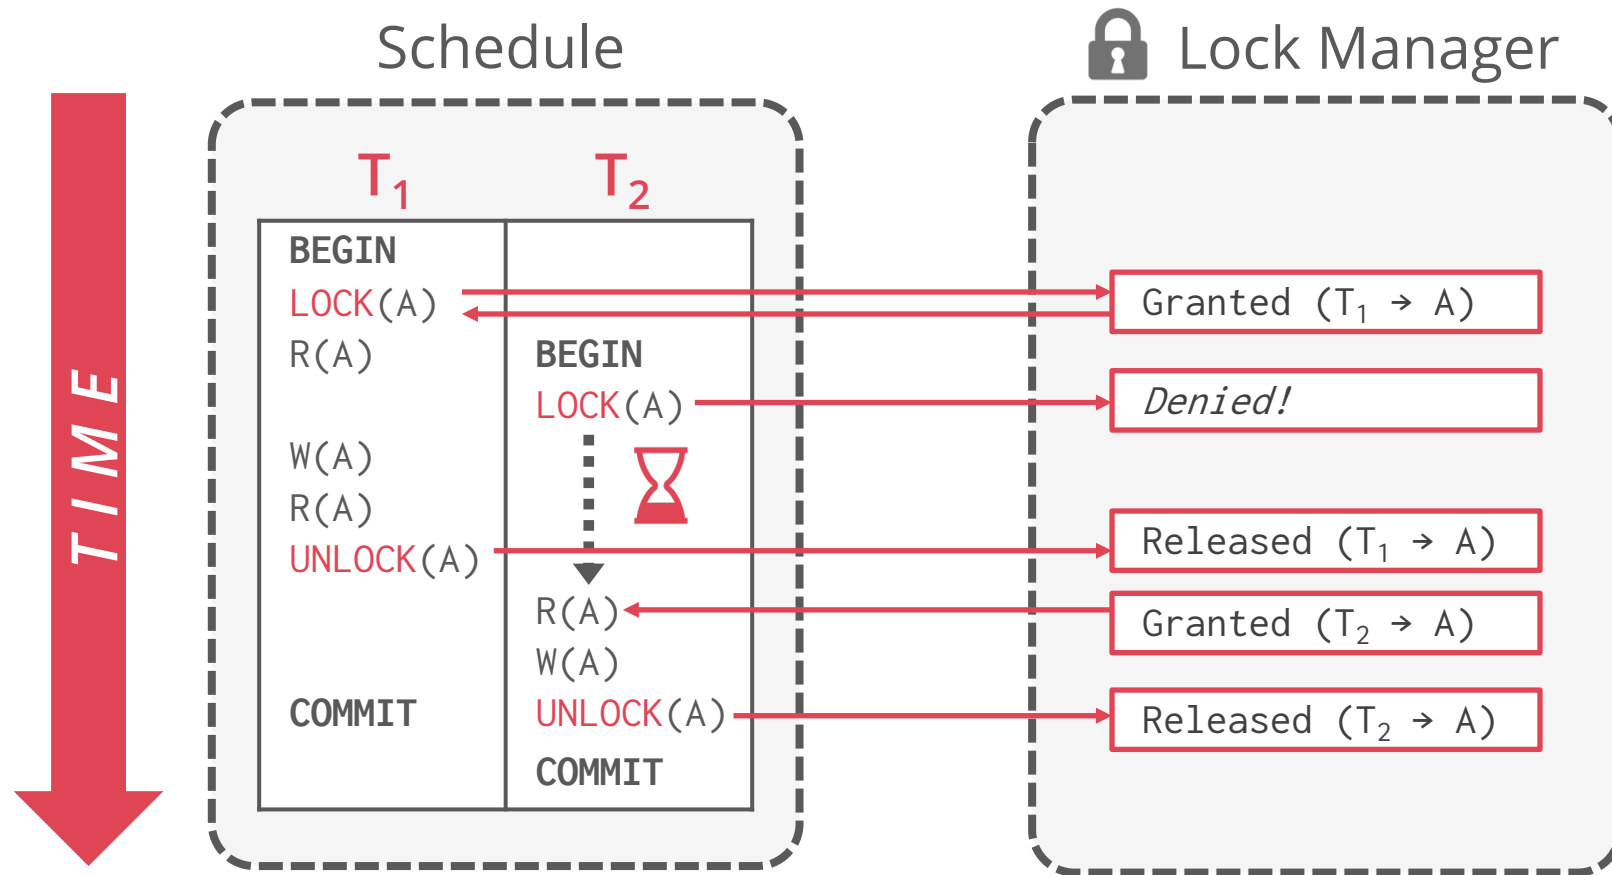
Record changes for each txn individually

Validate and possibly rollback on commit

Used if conflicts are rare (e.g., write-once-read-many scenarios)



EXECUTING WITH LOCKS



EXECUTING WITH LOCKS

Basic lock types:

S-LOCK: Shared locks for reads

X-LOCK: Exclusive locks for writes

Steps:

Transactions request locks (or upgrades) before accessing objects

Lock manager grants or blocks requests

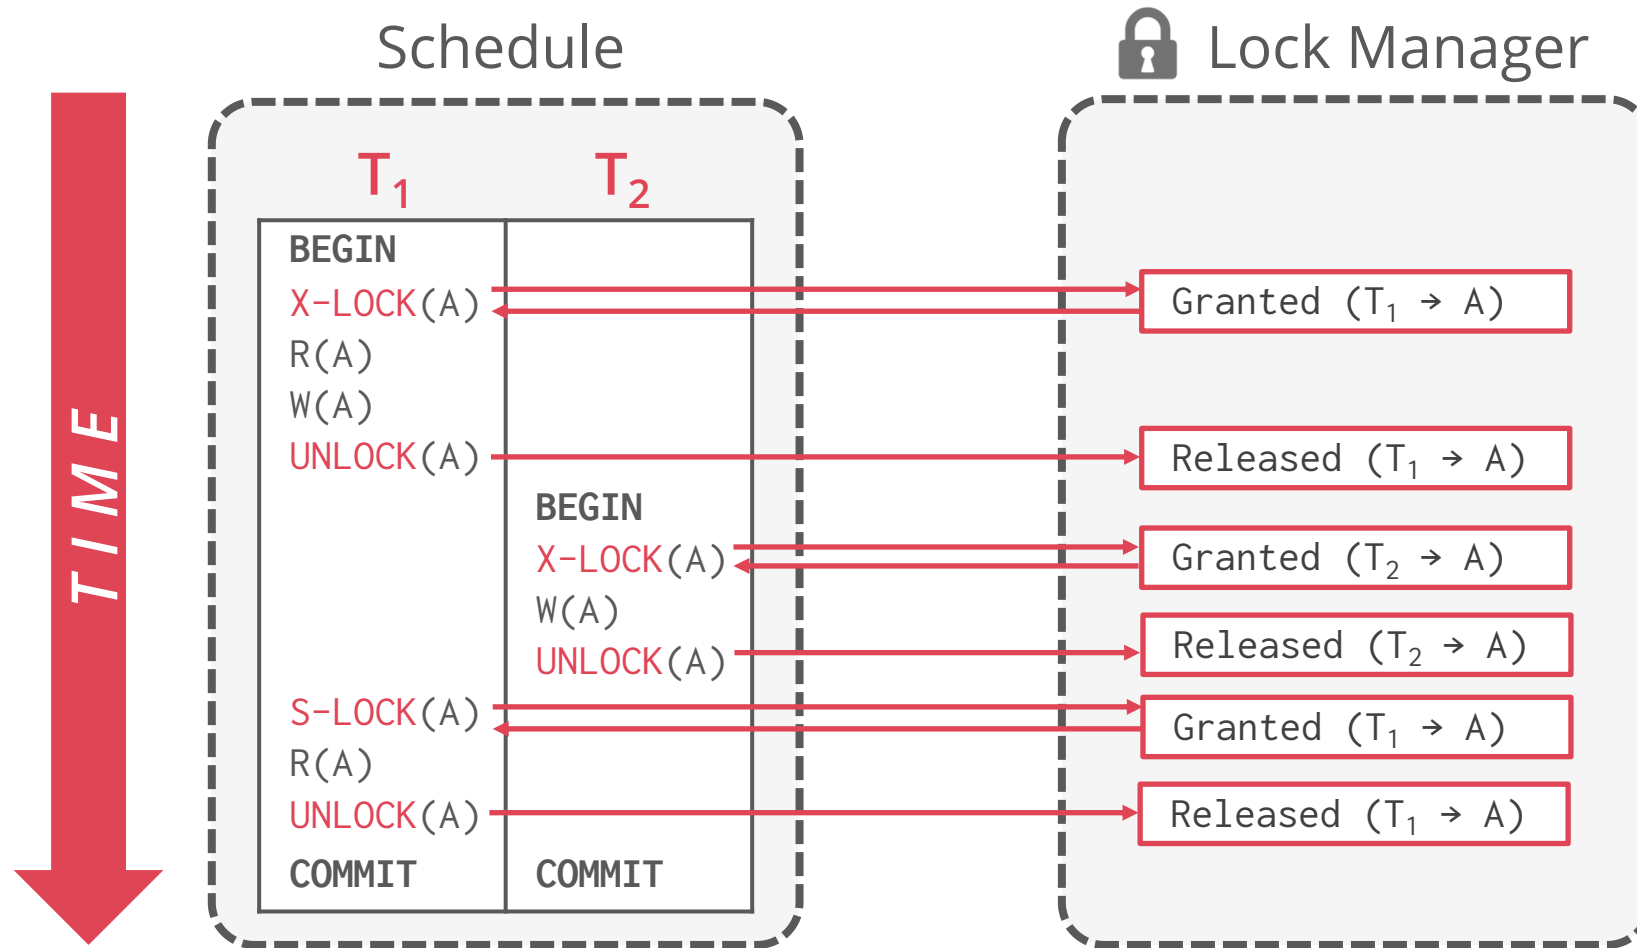
Transactions release locks

Lock manager updates its internal lock-table

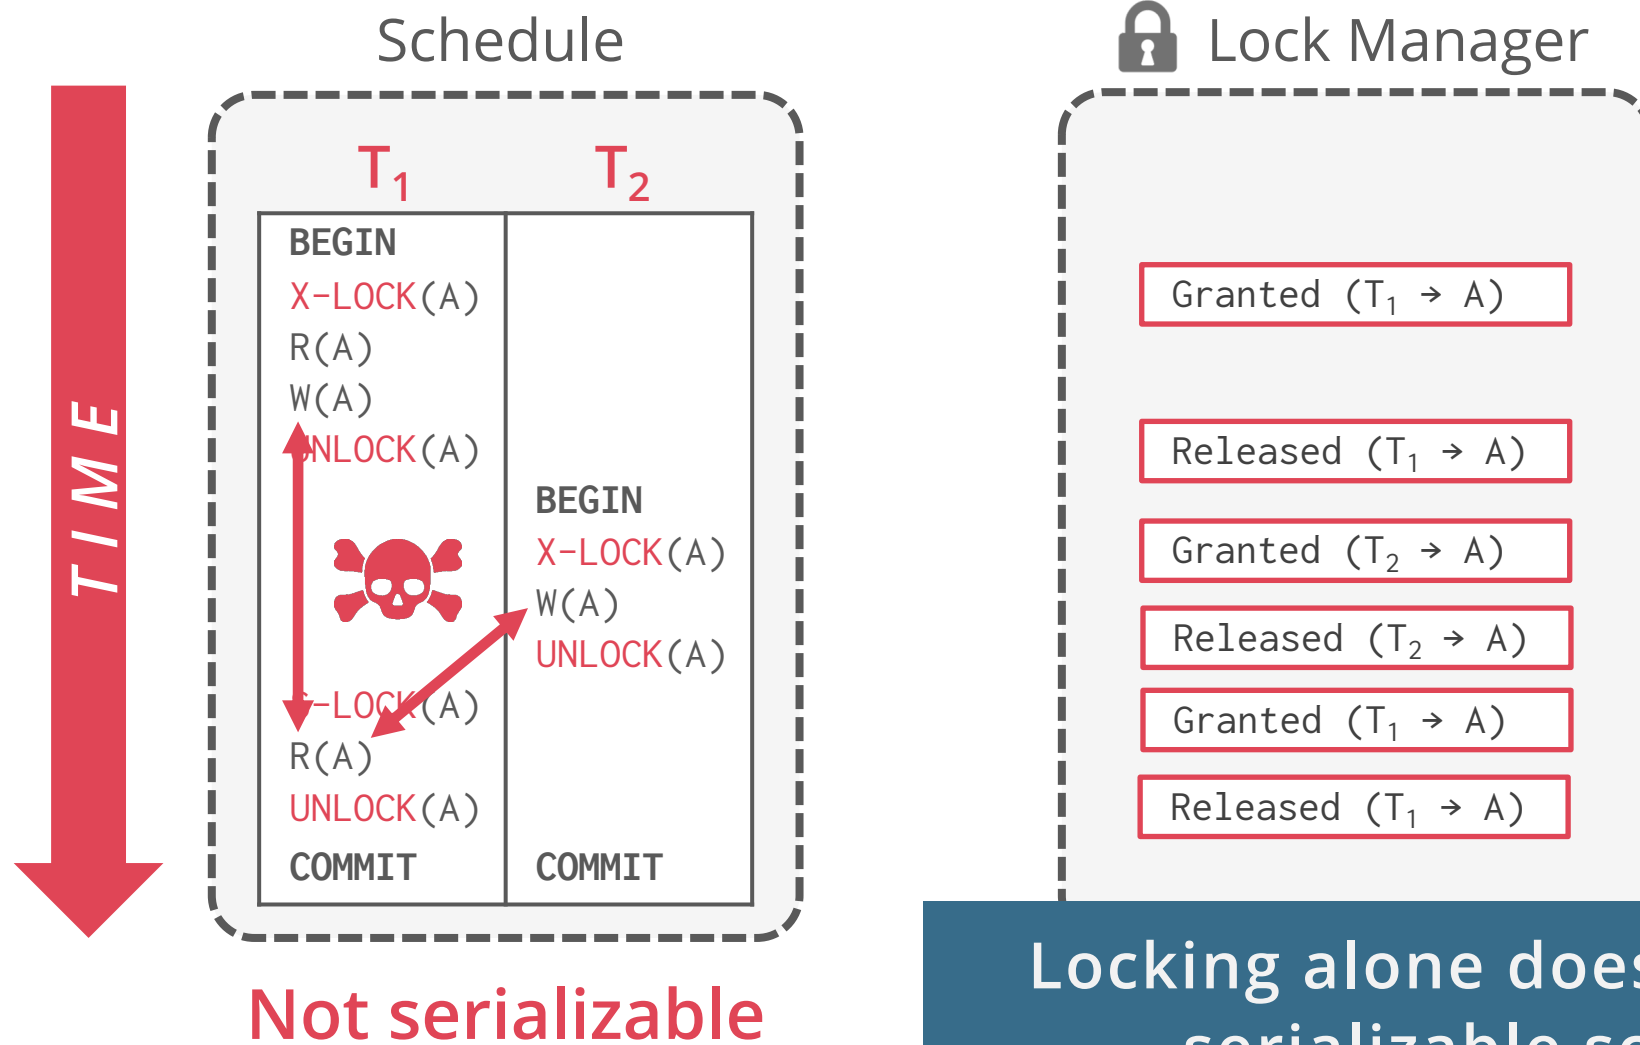
Compatibility Matrix

	Shared	Exclusive
Shared	✓	✗
Exclusive	✗	✗

EXECUTING WITH LOCKS



EXECUTING WITH LOCKS



TWO-PHASE LOCKING

Locks + concurrency control protocol

Determines if a txn is allowed to access an object in the database on the fly

Does not need to know all of the queries that a txn will execute ahead of time

Phase 1: Growing

Each txn requests the locks that it needs from the lock manager

The lock manager grants/denies lock requests


Phase 2: Shrinking

The txn is allowed to only release locks that it previously acquired

It cannot acquire new locks


TWO-PHASE LOCKING

The transaction is not allowed to acquire/upgrade locks after the growing phase finishes

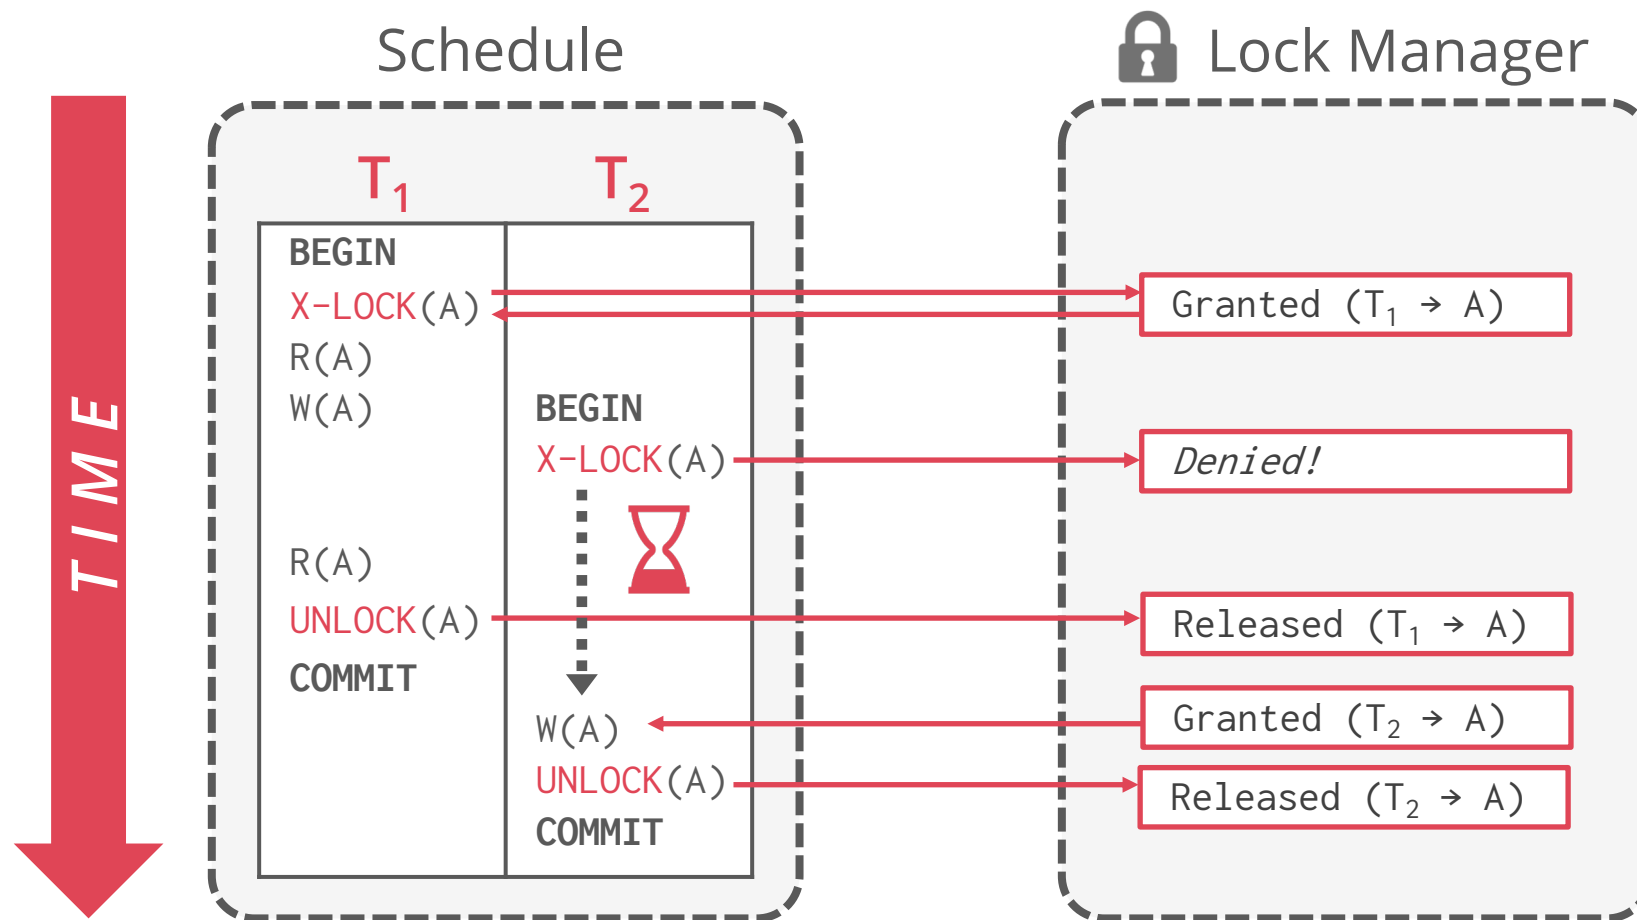


TWO-PHASE LOCKING

The transaction is not allowed to acquire/upgrade locks after the growing phase finishes

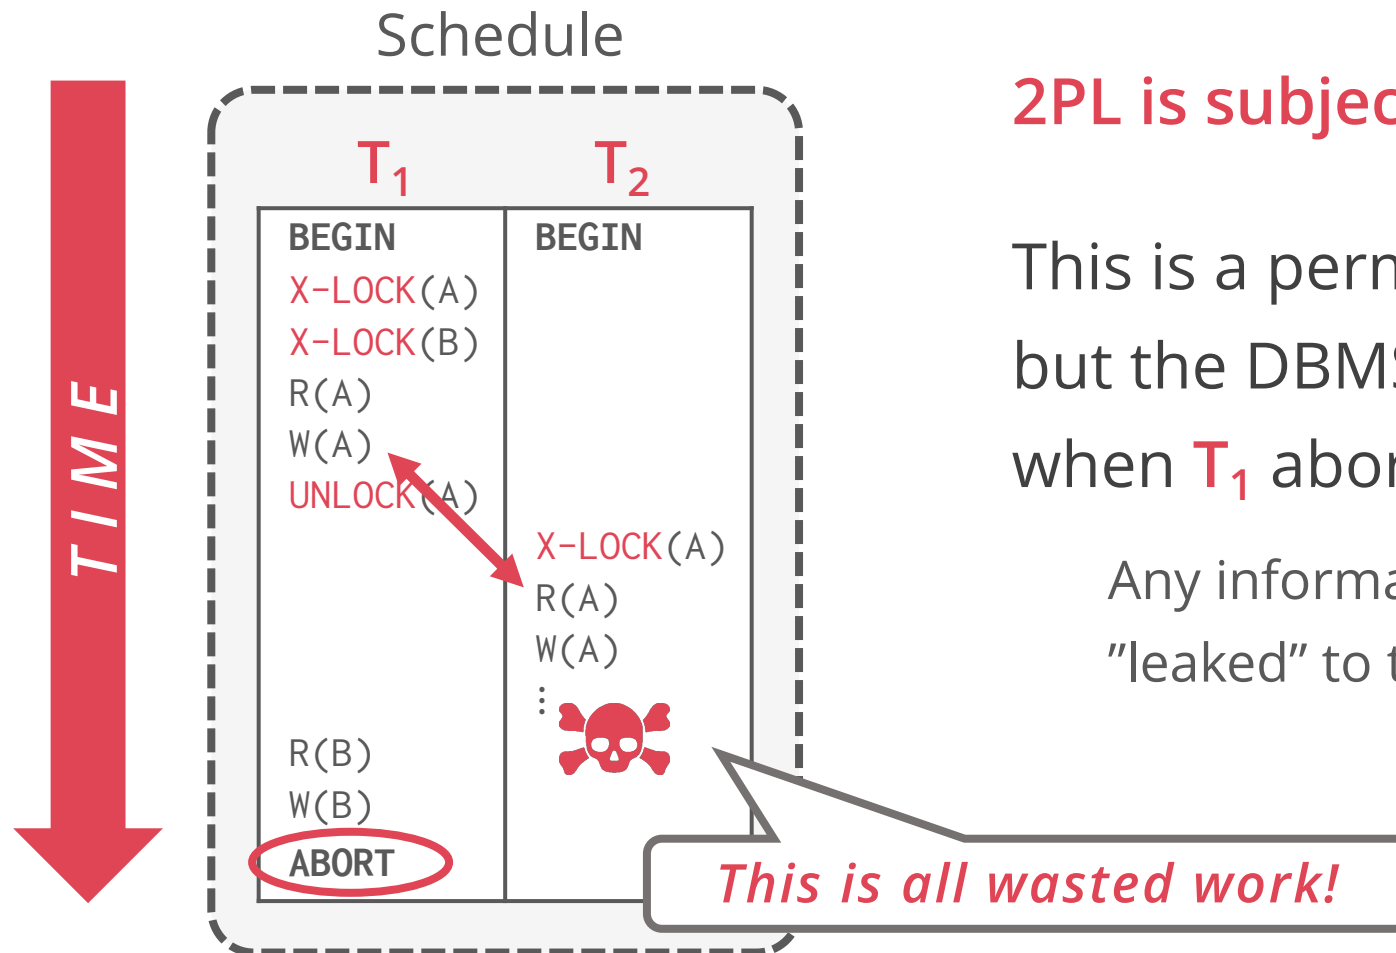


EXECUTING WITH LOCKS



2PL is sufficient to guarantee conflict-serializability
 (generates schedules whose precedence graph is acyclic)

2PL – CASCADING ABORTS



2PL is subject to cascading aborts

This is a permissible schedule in 2PL but the DBMS has to also abort T_2 when T_1 aborts

Any information about T_1 cannot be "leaked" to the outside world

2PL OBSERVATIONS

There are schedules that are serializable but not be allowed by 2PL

Locking limits concurrency

May require **cascading aborts**

Solution: Strict 2PL

May still have **"dirty reads"**

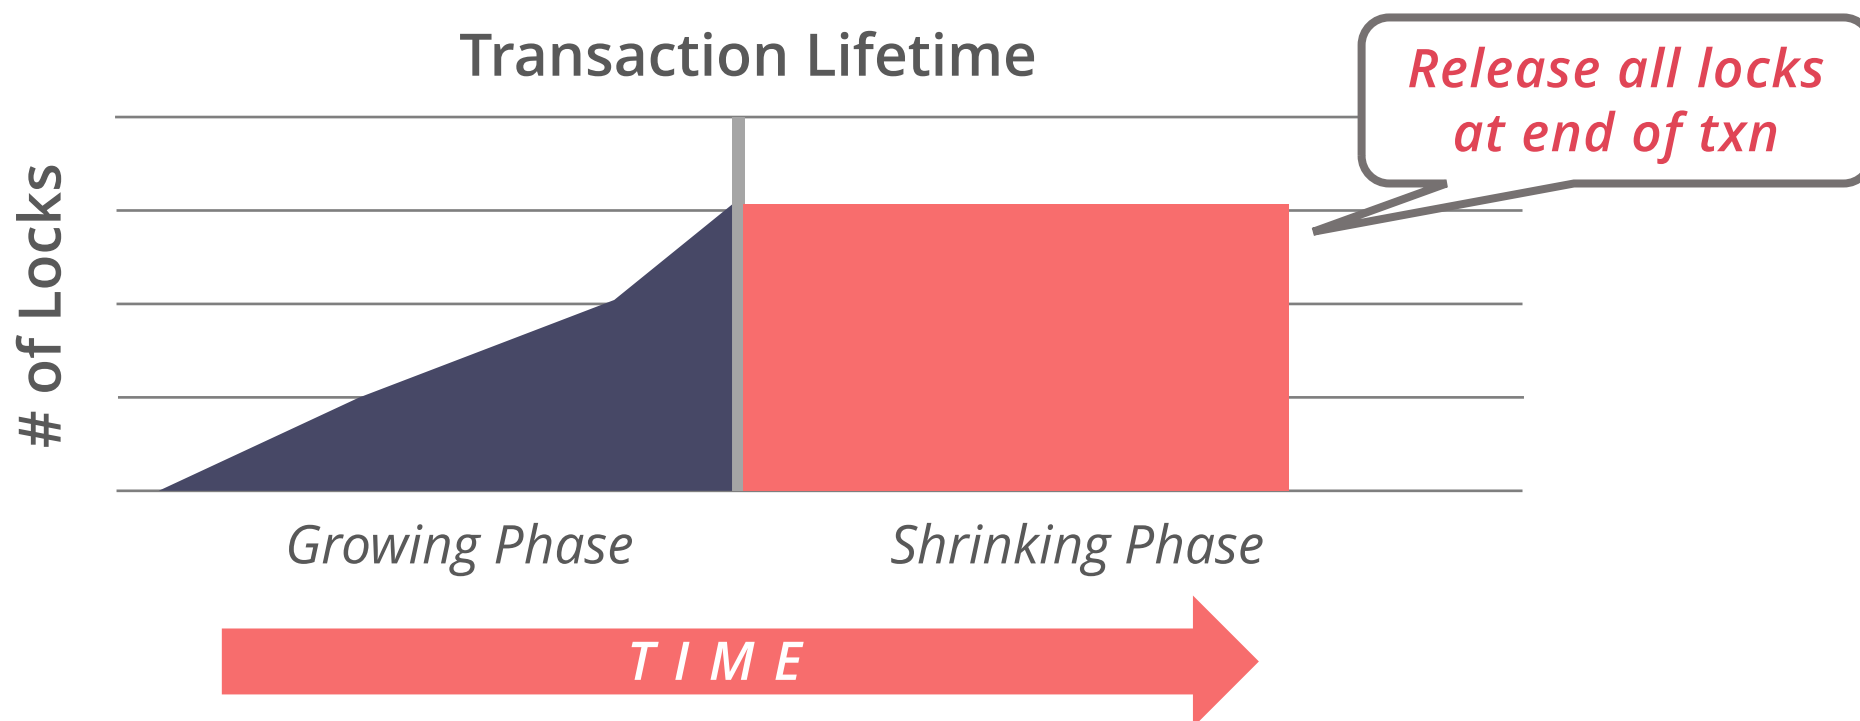
Solution: Strict 2PL

May lead to **deadlocks**

Solution: Detection or Prevention

STRICT TWO-PHASE LOCKING

The txn is not allowed to acquire/upgrade locks after the growing phase finishes
Allows only conflict-serializable schedules, but it is often stronger than needed for some applications

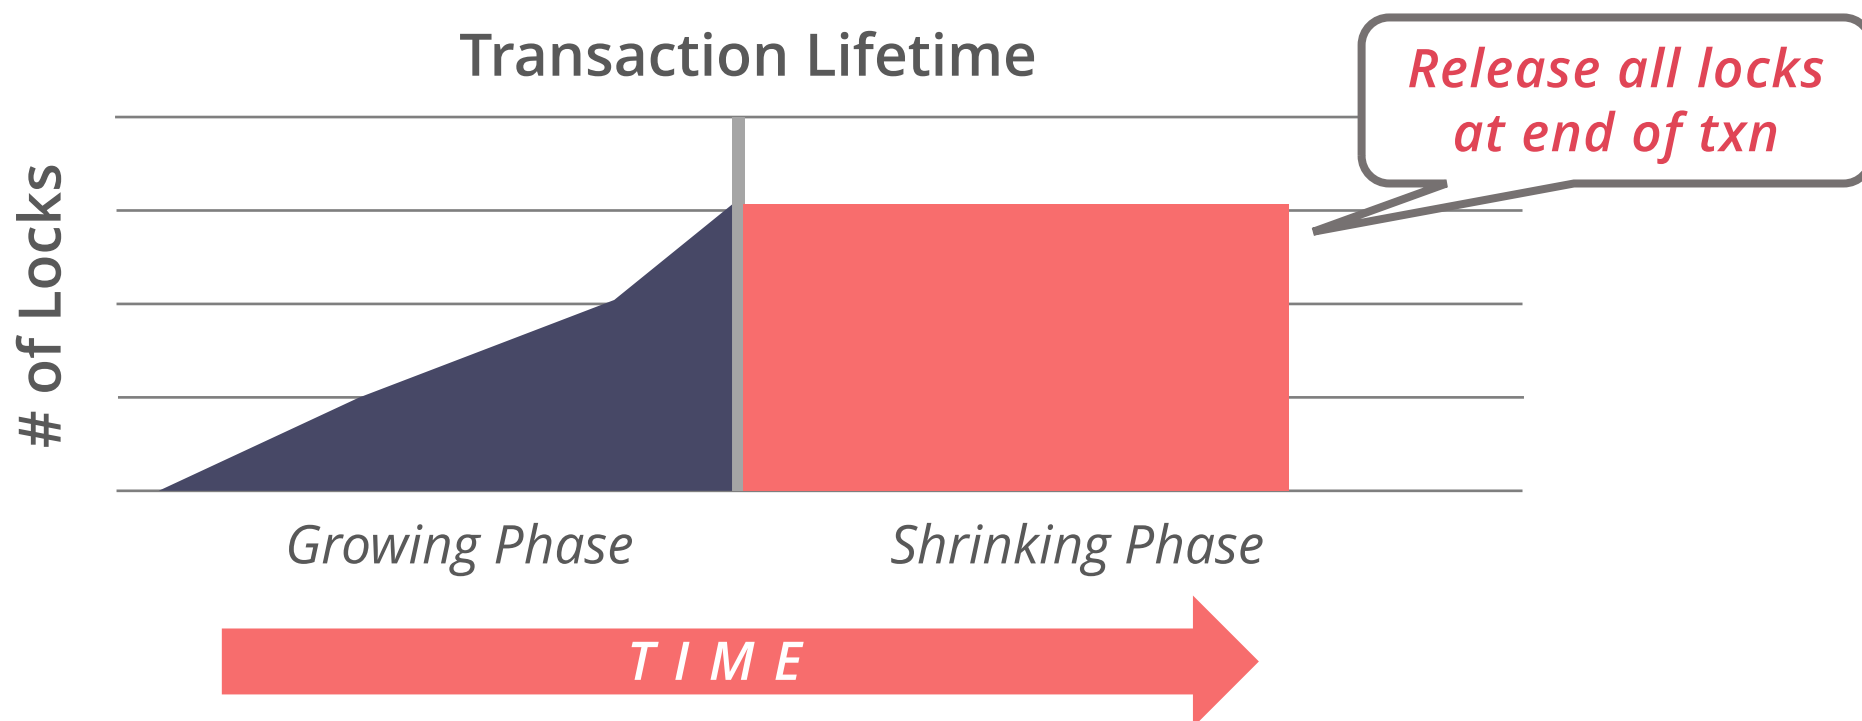


STRICT TWO-PHASE LOCKING

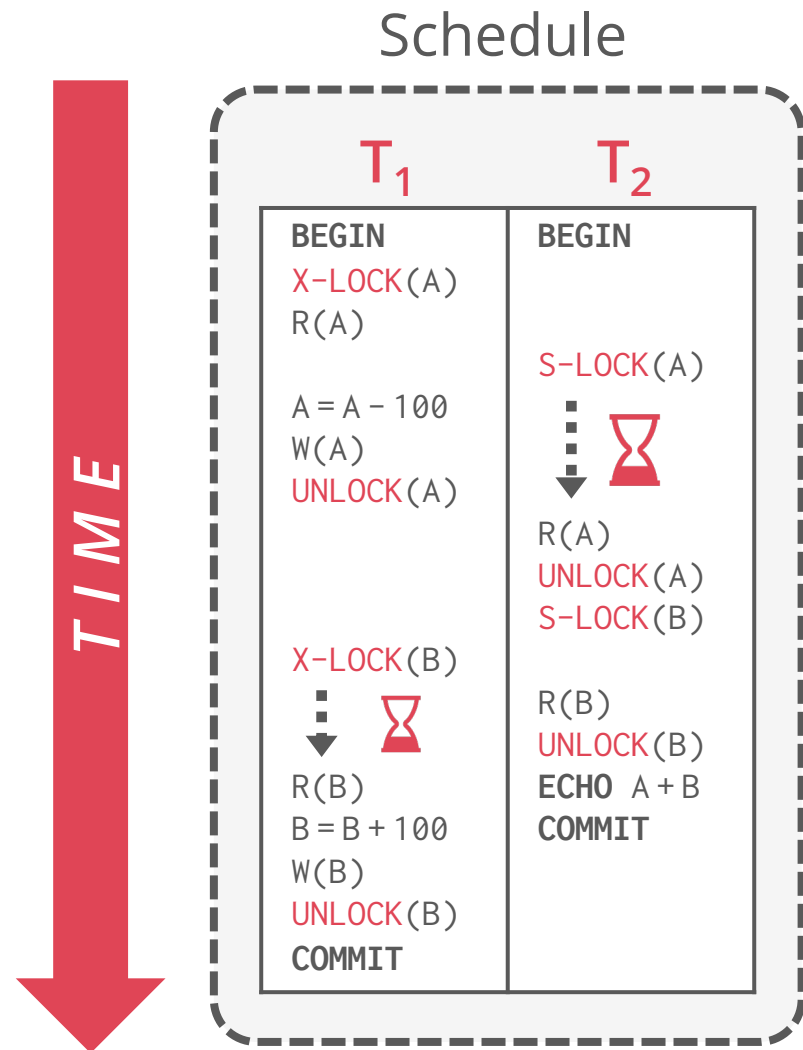
Advantages:

Does not incur cascading aborts

Aborted txns can be undone by just restoring original values of modified tuples



Non-2PL Example



T_1 – move £100 from account A to account B

T_2 – compute the total amount in all accounts and return it to the application

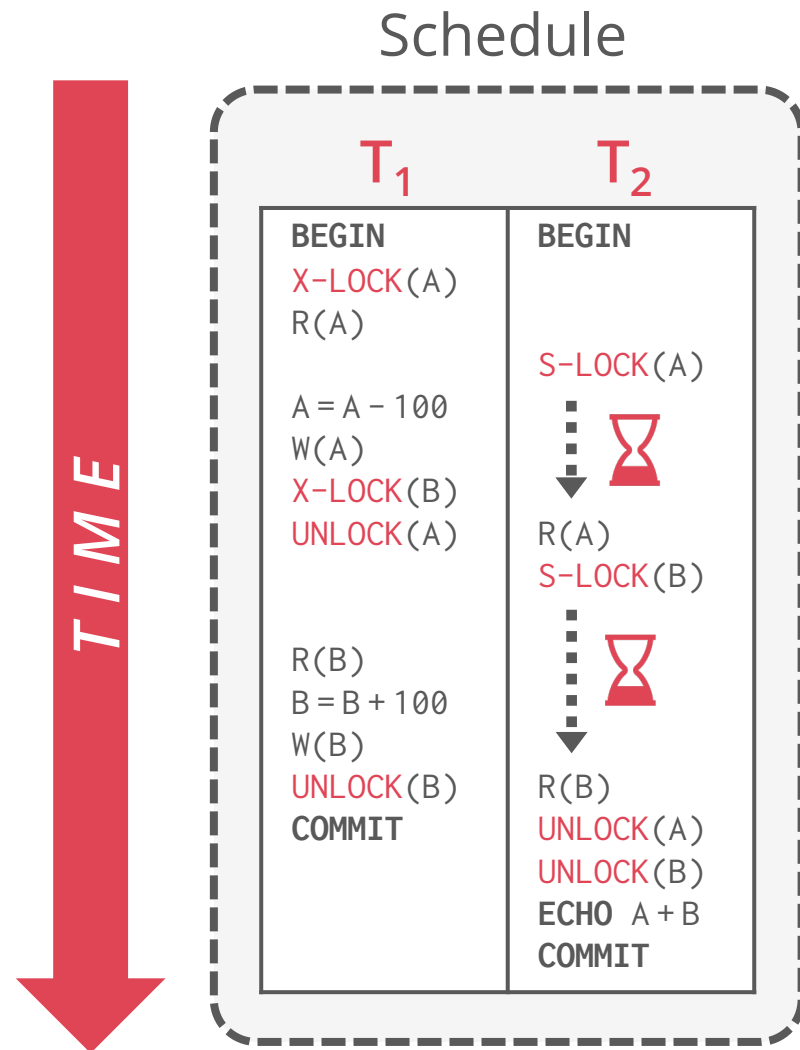
Initial Database State

$A = 1000, B = 1000$

T_2 Output

$A + B = 1900$

2PL EXAMPLE



T_1 – move £100 from account A to account B

T_2 – compute the total amount in all accounts and return it to the application

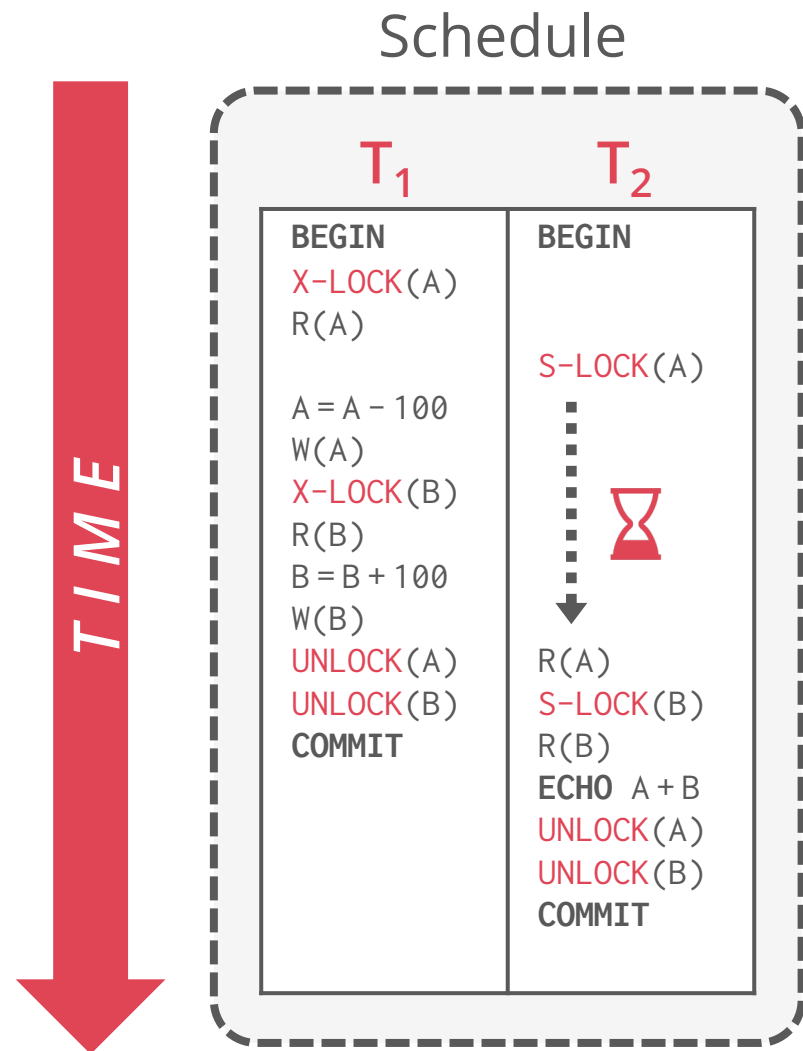
Initial Database State

$A = 1000, B = 1000$

T_2 Output

$A + B = 2000$

STRICT 2PL EXAMPLE



T_1 – move £100 from account A to account B

T_2 – compute the total amount in all accounts and return it to the application

Initial Database State

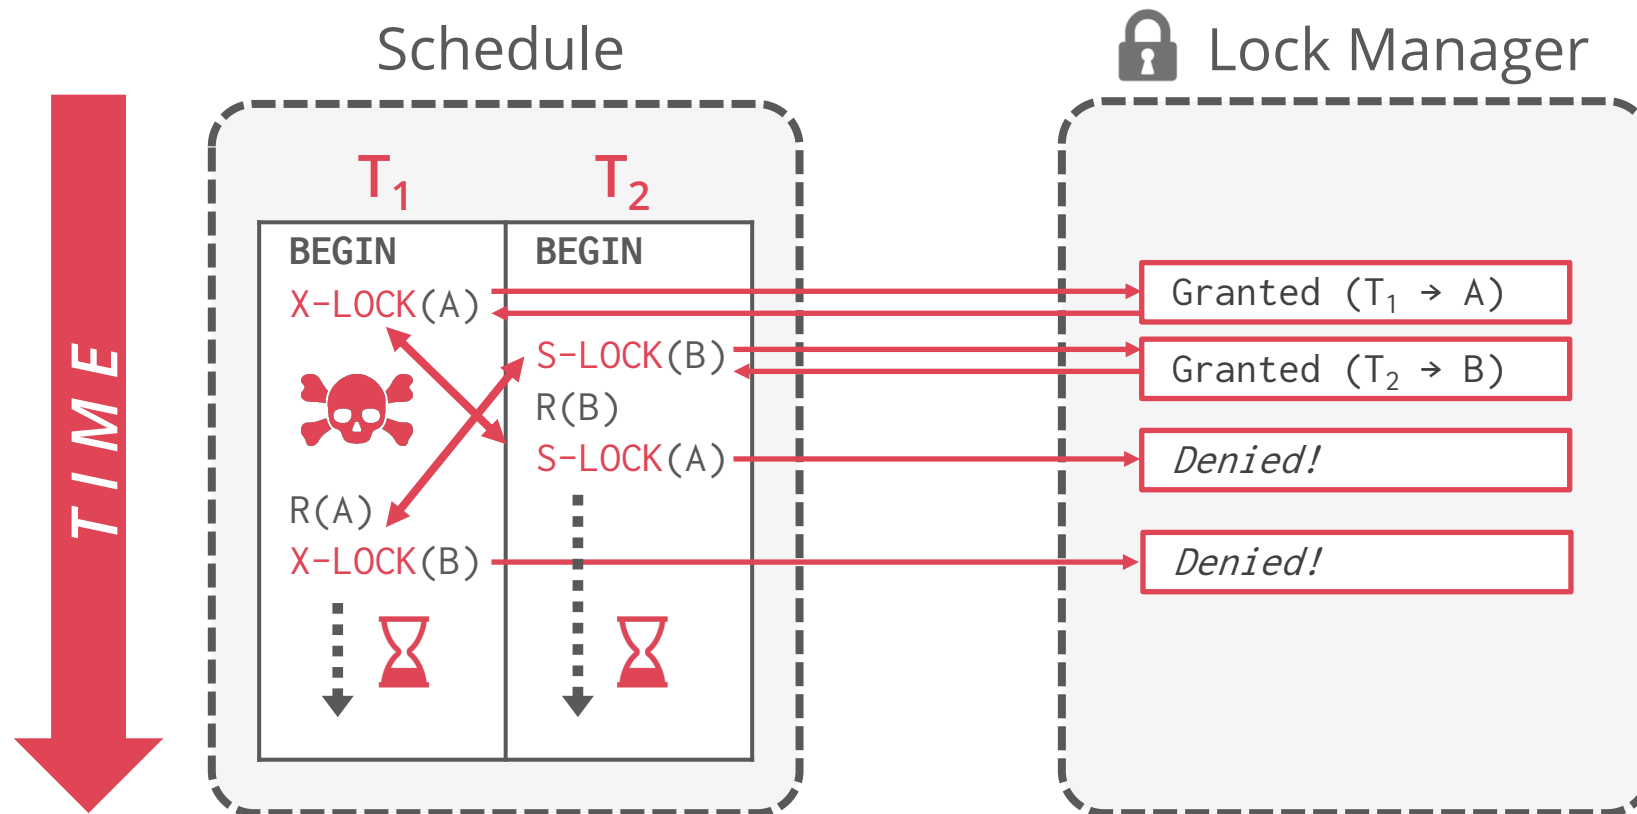
$A = 1000, B = 1000$

T_2 Output

$A + B = 2000$

SCHEDULING: DEADLOCKS

Two-phase locking has the risk of **deadlock situations**



2PL DEADLOCKS


Deadlock = a cycle of txns waiting for locks to be released by each other

Two ways of dealing with deadlocks:

Deadlock Detection

Deadlock Prevention

Conservative (or “preclaiming”) 2PL also prevents deadlocks. Why?



DEADLOCK DETECTION



The DBMS creates a **waits-for** graph to keep track of what locks each transaction is waiting to acquire:

Nodes are transactions

Edge from T_i to T_j if T_i is waiting for T_j to release a lock

The system periodically checks for cycles in waits-for graph and then make a decision on how to break it

DEADLOCK DETECTION



DEADLOCK HANDLING

Upon detecting a deadlock, the DBMS selects a “**victim**” transaction to rollback to break the cycle

Selecting a “victim” transaction might depend on:

- age (lowest timestamp)

- progress (least/most executed queries)

- # of items already locked

- # of txns that we have to rollback with it

- # of previous restarts (to prevent starvation)

There is a trade-off between the frequency of checking for deadlocks and how long transactions have to wait before deadlocks are broken

DEADLOCK PREVENTION

When a transaction tries to acquire a lock that is held by another transaction, kill one of them to prevent a deadlock

No waits-for graph or detection algorithm

Assign **priorities** based on timestamps

Older \Rightarrow higher priority (e.g., $T_1 > T_2$)

Two deadlock prevention policies:

Wait-Die ("Old Waits for Young")

Wound-Wait ("Young Waits for Old")

DEADLOCK PREVENTION

Wait-Die (“Old Waits for Young”)

If *requesting* txn has higher priority than *holding* txn

Then *requesting* txn **waits** for *holding* txn

Else *requesting* txn **aborts**

Wound-Wait (“Young Waits for Old”)

If *requesting* txn has higher priority than *holding* txn

Then *holding* txn **aborts** and releases locks

Else *requesting* txn **waits**

DEADLOCK PREVENTION

$T_{\text{req}} > T_{\text{hold}}$?



Wait

:



Die

$T_{\text{req}} > T_{\text{hold}}$?




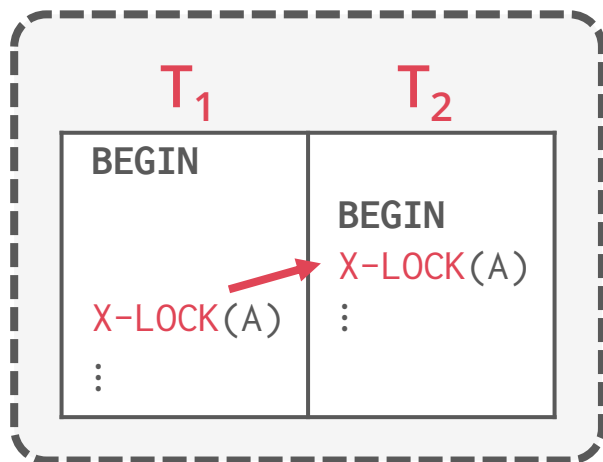
Wound

:



Wait

DEADLOCK PREVENTION




Wait-Die

T_1 waits

Wound-Wait

T_2 aborts



Wait-Die

T_2 aborts

Wound-Wait

T_2 waits

DEADLOCK PREVENTION

Why do these schemes guarantee no deadlocks?

Only one “type” of direction allowed when waiting for a lock

When a transaction restarts, what is its (new) priority?

Its original timestamp. Why?

CONCLUSION

ACID Transactions

Atomicity: All or nothing

Consistency: Only valid data

Isolation: No interference

Durability: Committed data persists

Serializability

Serializable schedules

Conflict & view serializability

Checking for conflict serializability

Concurrency Control

Prevent anomalous schedules

Locks + protocol (2PL, Strict 2PL) guarantees conflict serializability

Deadlock detection and deadlock prevention