



THE UNIVERSITY  
*of* EDINBURGH

# Advanced Databases

Spring 2020

---

Lecture #15:

## Crash Recovery

Milos Nikolic

# REVIEW: THE ACID PROPERTIES

**A**tomicity: All actions in the txn happen, or none happen

**C**onsistency: If each txn is consistent and the DB starts consistent, then it ends up consistent

**I**solation: Execution of one txn is isolated from that of other txns

**D**urability: If a txn commits, its effects persist

The **recovery manager** ensures atomicity, DB consistency, and durability

# MOTIVATION

Atomicity:

Transactions may abort (“rollback”)


Durability:

What if the DBMS stops running?

Desired behaviour after system restarts:

**T1, T2 & T3** should be durable

T4 & T5 should be aborted (effects not seen)



# TYPES OF FAILURES

## Logical Errors

Txn cannot complete due to an internal error condition (e.g., integrity constraint violation)

## Internal State Errors

DBMS must terminate an active transaction due to an error condition (e.g., deadlock)

---

## Software Failures

Problem with the DBMS implementation (e.g., uncaught divide-by-zero exception)

## Hardware Failures

The computer hosting the DBMS crashes (e.g., power plug gets pulled)

Fail-stop assumption: Non-volatile storage contents are not corrupted by system crash

---

## Non-Repairable Hardware Failure

A head crash or similar disk failure destroys all or part of non-volatile storage

Destruction is assumed to be detectable (e.g., disk controller use checksums to detect failures)

*No DBMS can recover from this! Database must be restored from an archived version (replica).*

*Transaction  
Failures*

*System  
Failures*

*Storage  
Media  
Failures*

# CRASH RECOVERY

Recovery algorithms are techniques to ensure database consistency, transaction atomicity, and durability despite failures

Recovery algorithms have two parts:

- Actions during normal txn processing to ensure that the DBMS can recover from a failure

- Actions after a failure to recover the database to a state that ensures atomicity, consistency, and durability

# OBSERVATION

The primary storage location of the database is on non-volatile storage, but this is much slower than volatile storage

Use volatile memory for faster access:

Bring pages into memory, perform writes in memory, write dirty pages back to disk

The DBMS needs to guarantee that:

The changes of any txn are durable once the DBMS has confirmed that it committed

No partial changes are durable if the txn aborted

How the DBMS supports this depends on how it manages the buffer pool...

# HANDLING THE BUFFER POOL

## Steal Policy

Whether the DBMS allows an uncommitted txn to overwrite the most recent committed value of an object in non-volatile storage

**STEAL**: Is allowed

**NO-STEAL**: Is not allowed

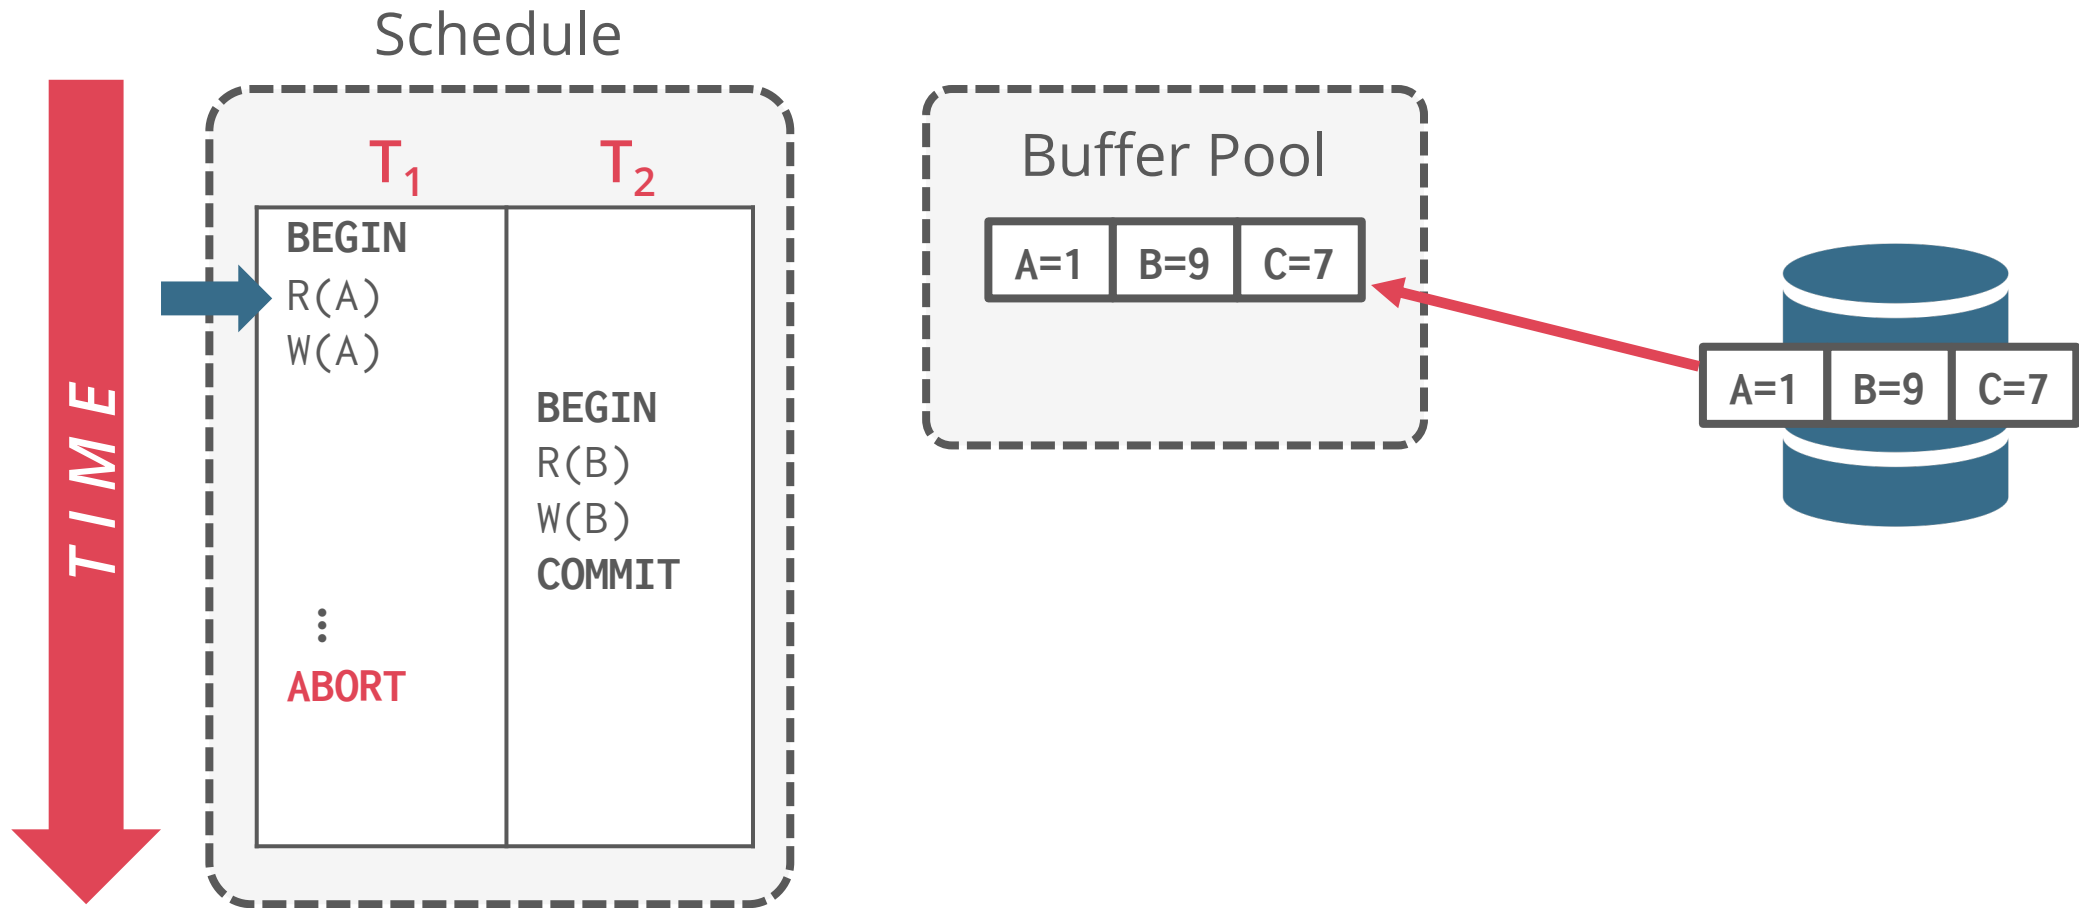
## Force Policy

Whether the DBMS requires that all updates made by a txn are reflected on non-volatile storage before the txn is allowed to commit

**FORCE**: Is enforced

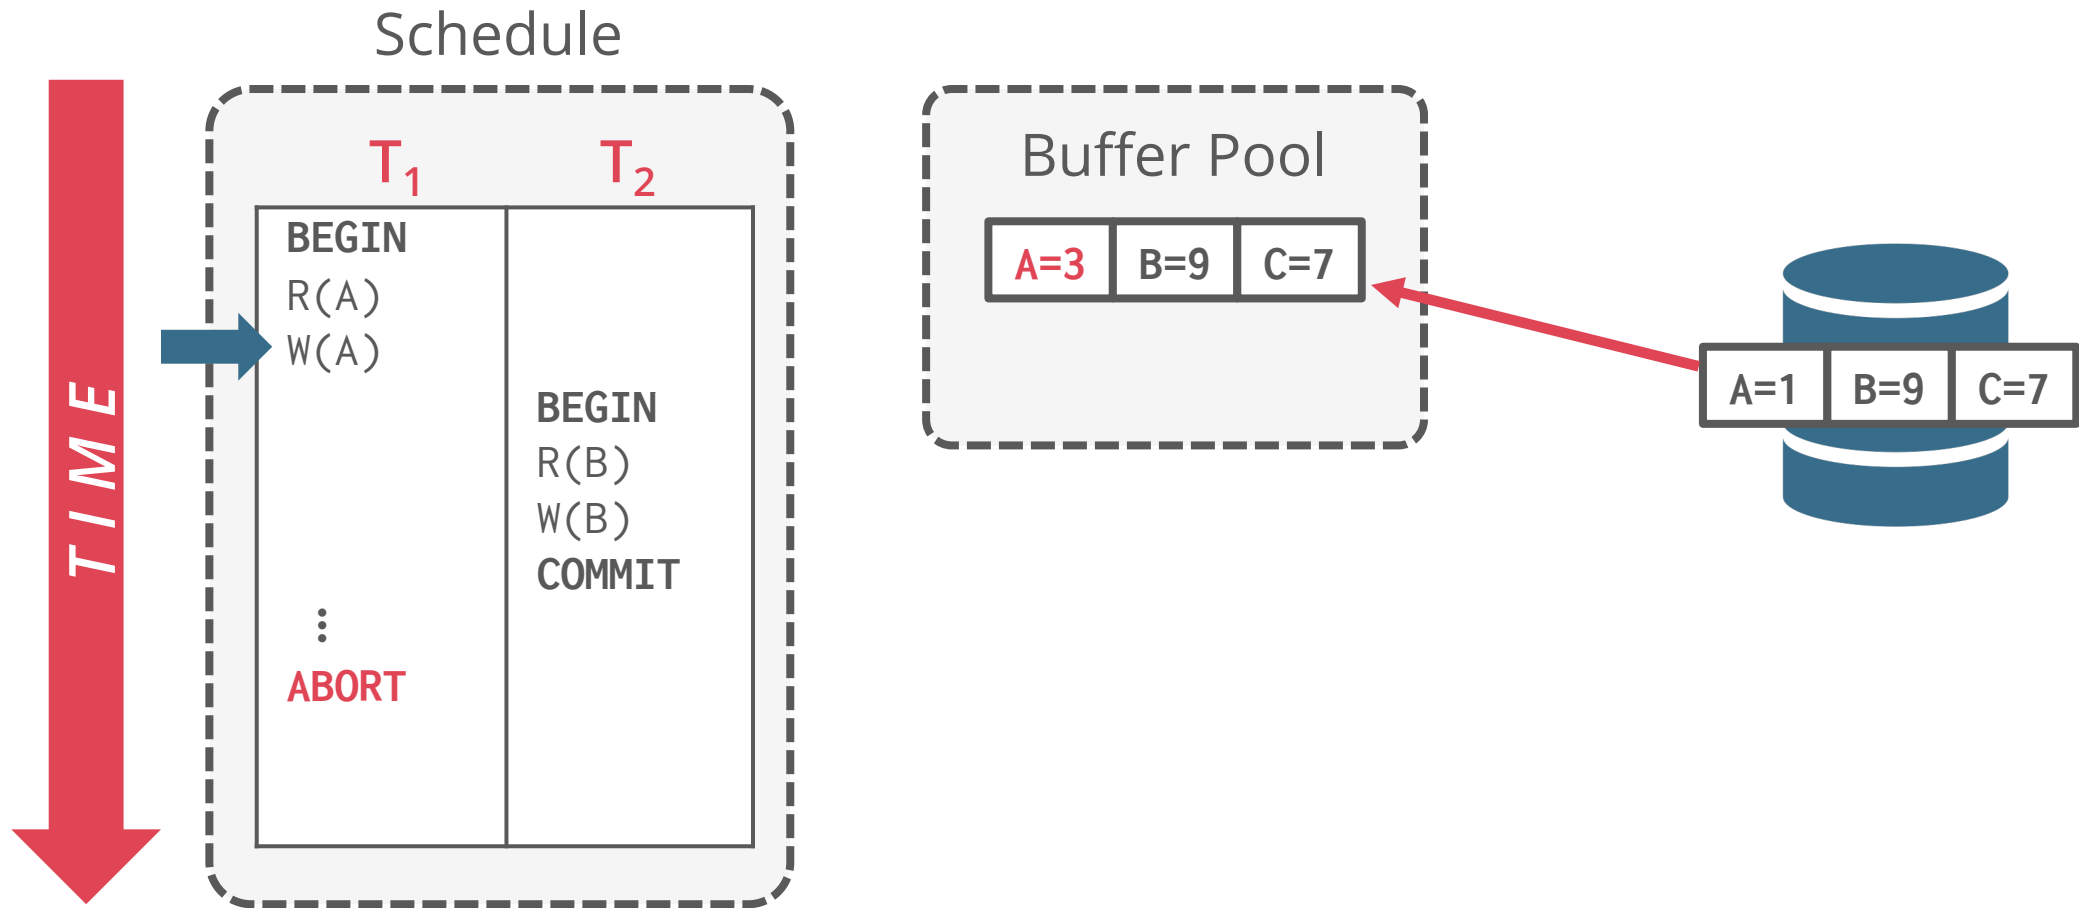
**NO-FORCE**: Is not enforced

# No-STEAL + FORCE

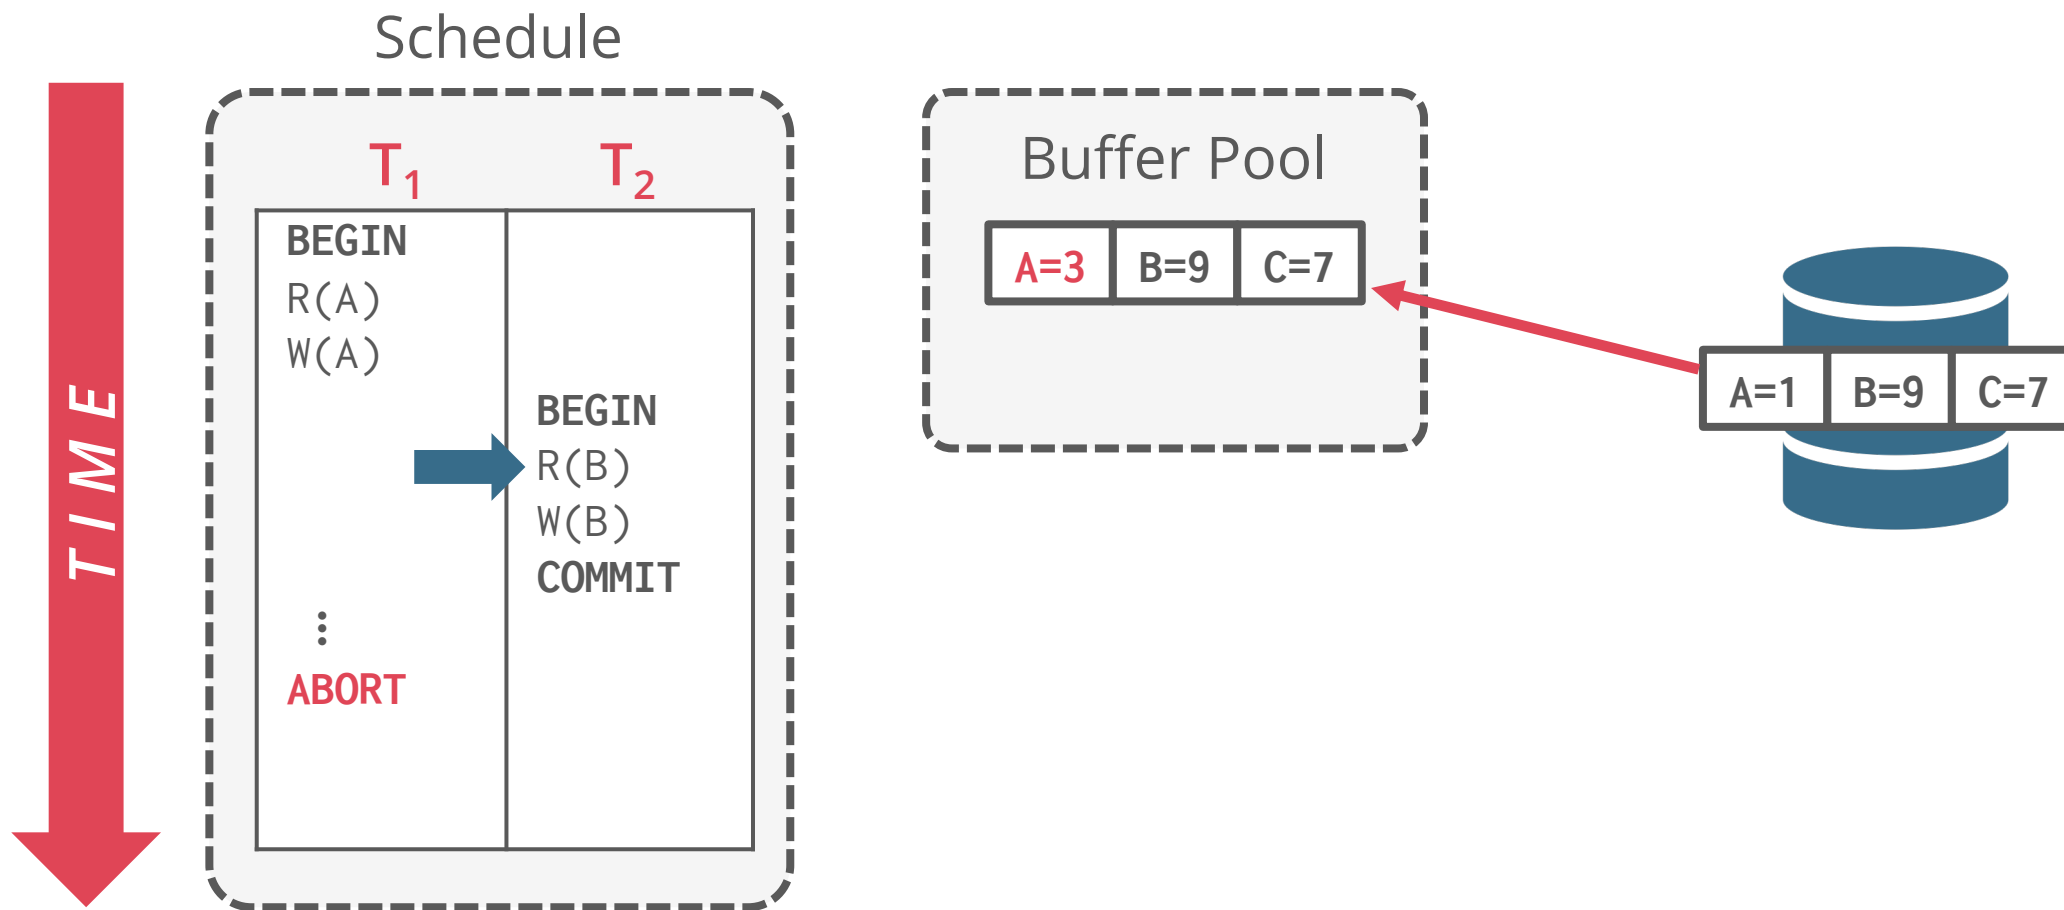




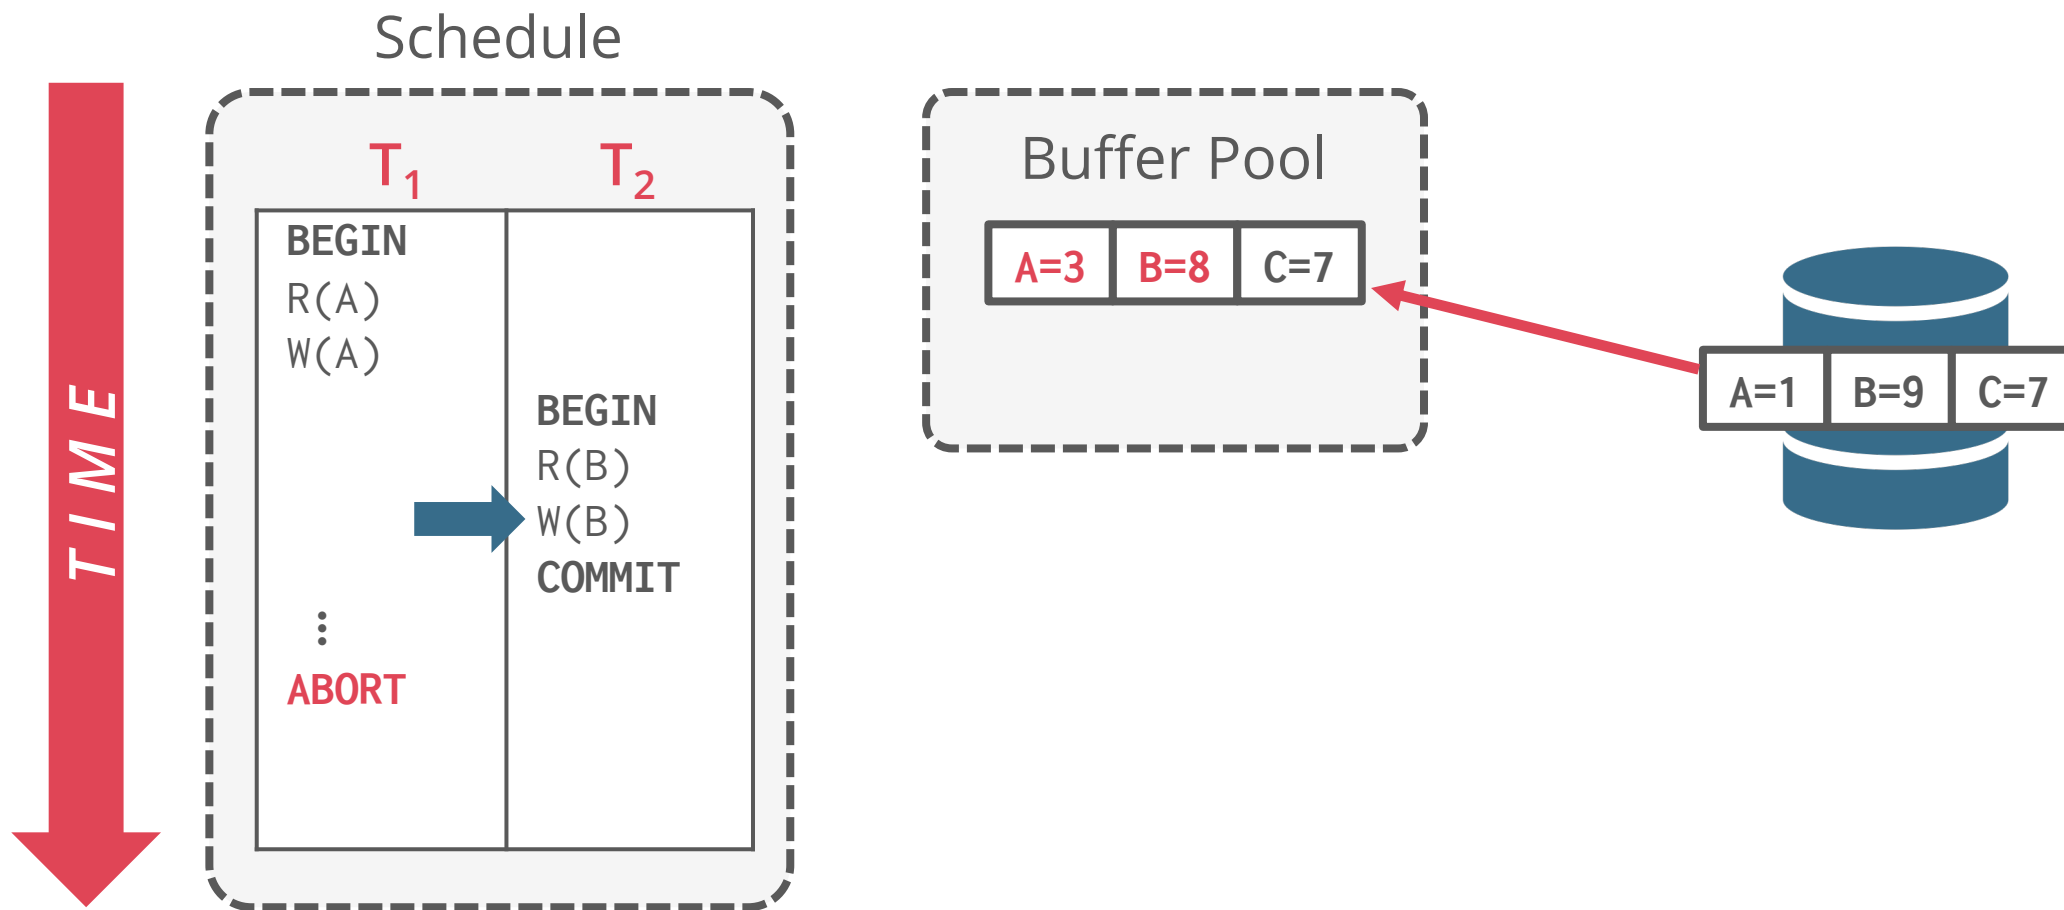
# No-STEAL + FORCE



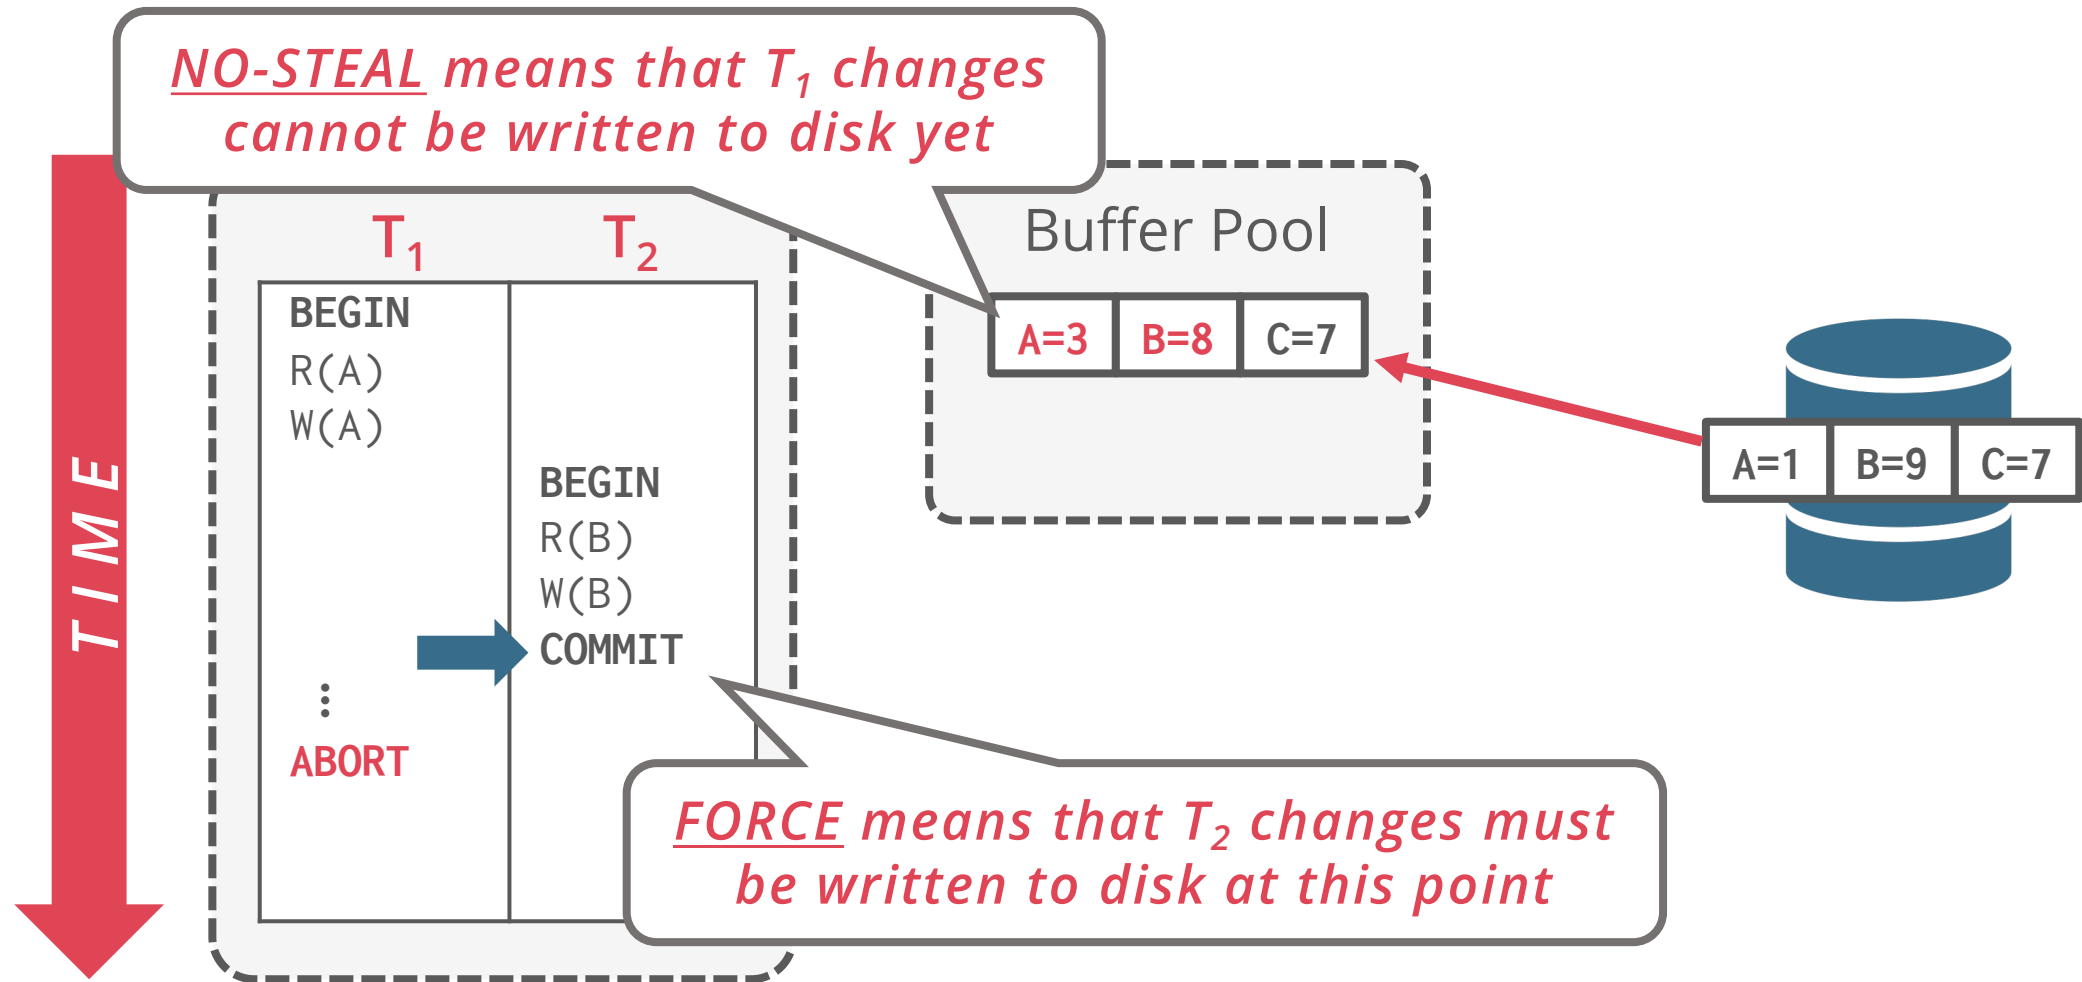
# No-STEAL + FORCE



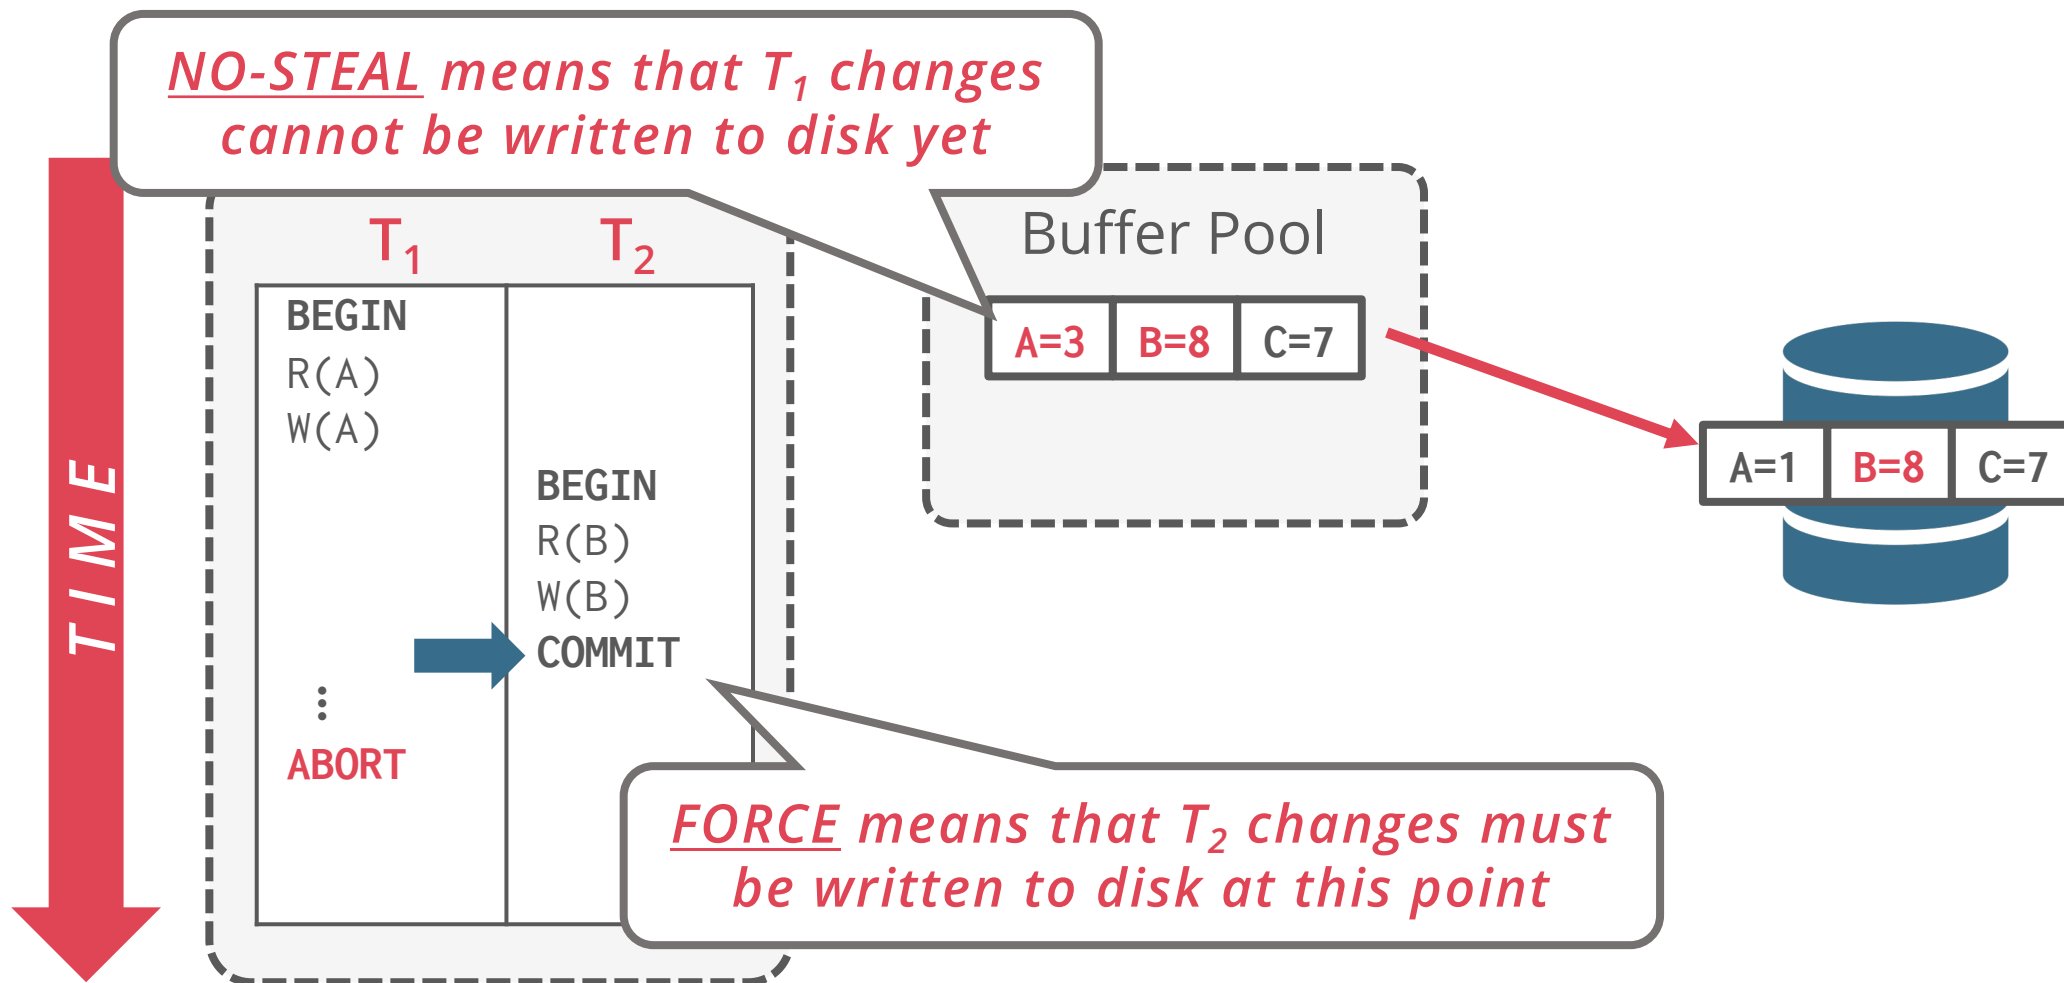
# No-STEAL + FORCE



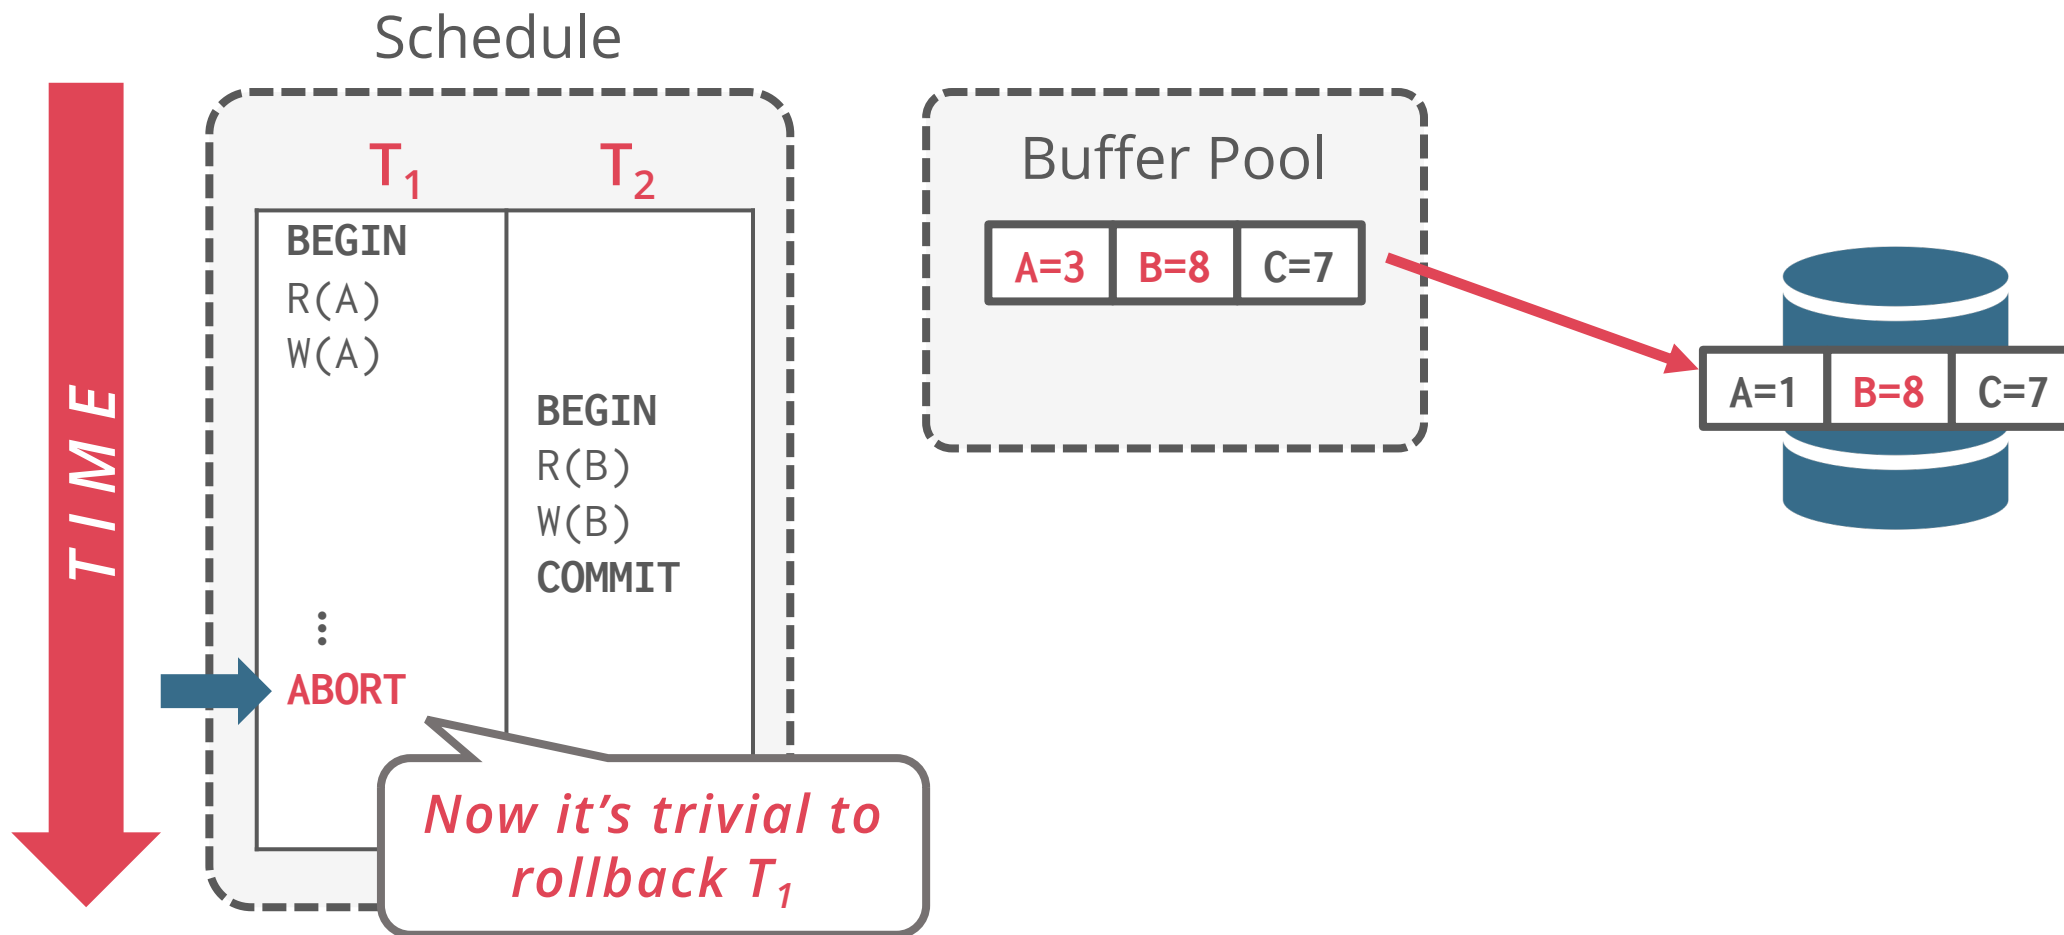
# No-STEAL + FORCE



# No-STEAL + FORCE



# No-STEAL + FORCE



# NO-STEAL + FORCE

This approach is the **easiest to implement**

**Never have to undo** changes of an aborted txn because the changes were not written to disk

**Never have to redo** changes of a committed txn because all the changes are guaranteed to be written to disk at commit time

But has **important drawbacks**

**Poor performance:** flushing non-contiguous pages (random writes) is slow

**Memory requirements:** **NO-STEAL** assumes that all pages modified by uncommitted transactions can be accommodated in the buffer pool

# MORE ON STEAL AND FORCE

**STEAL:** Why enforcing atomicity is hard?

Stealing frame  $F$ : Current page  $P$  in  $F$  is written to disk; some txn holds lock on  $P$

What if the txn with the lock on  $P$  aborts?

Must remember the old value of  $P$  at steal time to support **UNDO**ing the write to  $P$

**NO-FORCE:** Why enforcing durability is hard?

What if the DBMS crashes before a modified page is written to disk?

Write as little as possible, in a convenient place, at commit time, to support **REDO**ing modifications



# BUFFER POOL POLICIES

## Runtime Performance

	NO-STEAL	STEAL
NO-FORCE	–	Fastest
FORCE	Slowest	–

## Recovery Performance

	NO-STEAL	STEAL
NO-FORCE	–	Slowest
FORCE	Fastest	–

*Undo + Redo*

*No Undo + No Redo*

**Undo:** removing the effects of an incomplete or aborted txn

**Redo:** re-instating the effects of a committed txn for durability

Almost every DBMS uses **STEAL + NO-FORCE**

# BASIC IDEA: LOGGING

Record **UNDO** and **REDO** information, for every update, in a **log** file

- Assume that the log is on stable storage

- Log file is separated from actual data

- Sequential writes to the log

- Minimal info (diff) written to the log, so multiple updates fit in a single log page

Log contains sufficient information to perform the necessary undo and redo actions to restore the database after a crash

# WRITE-AHEAD LOGGING (WAL)





Record the changes made to the database in a log file before the change is made

The DBMS stages all of a txn's log records in volatile storage (usually backed by buffer pool)




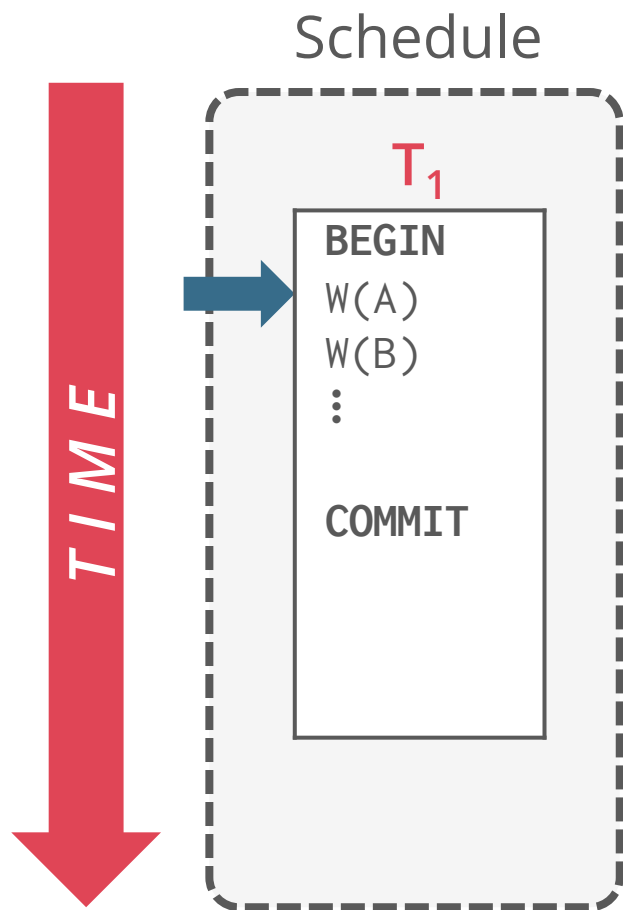
All log records pertaining to an updated page must be written to non-volatile storage before the page itself is overwritten in non-volatile storage

A txn is not considered committed until all its log records have been written to non-volatile storage





# WAL – EXAMPLE



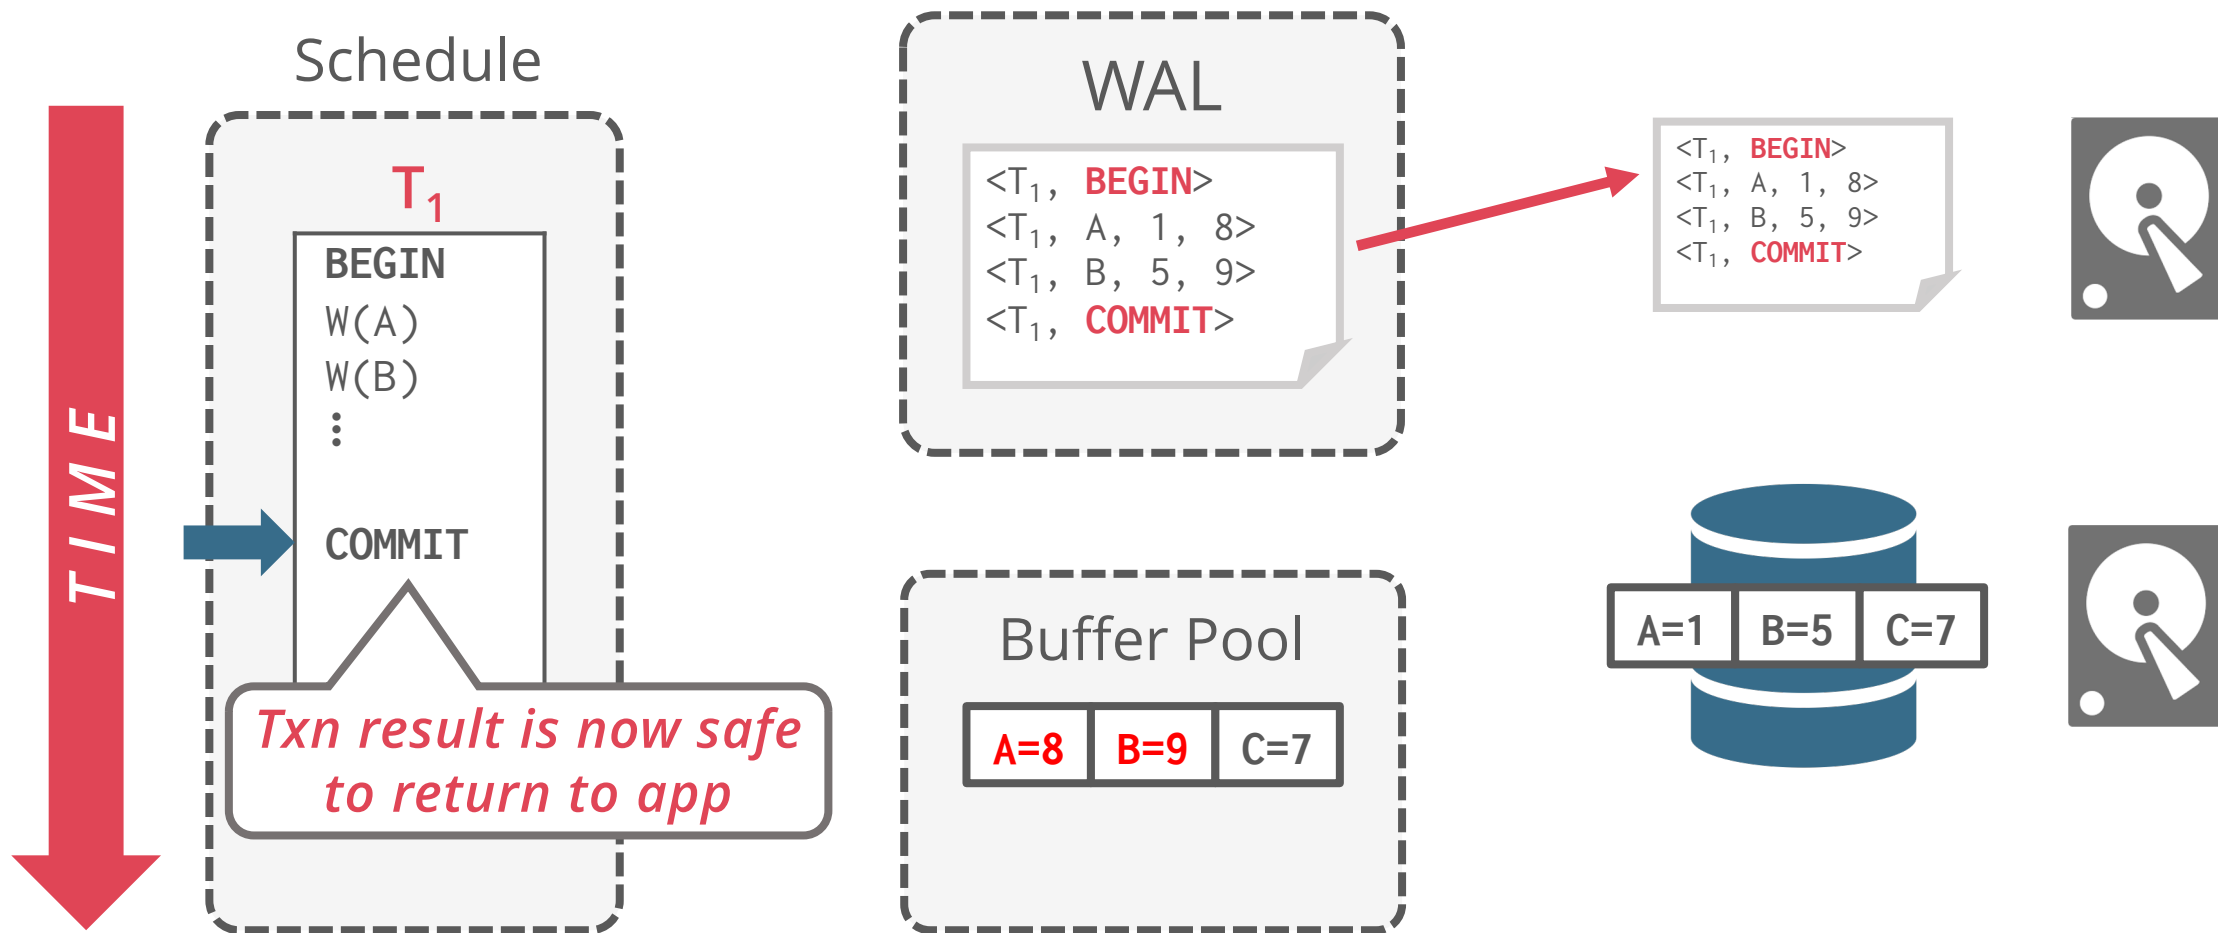
# WAL – EXAMPLE



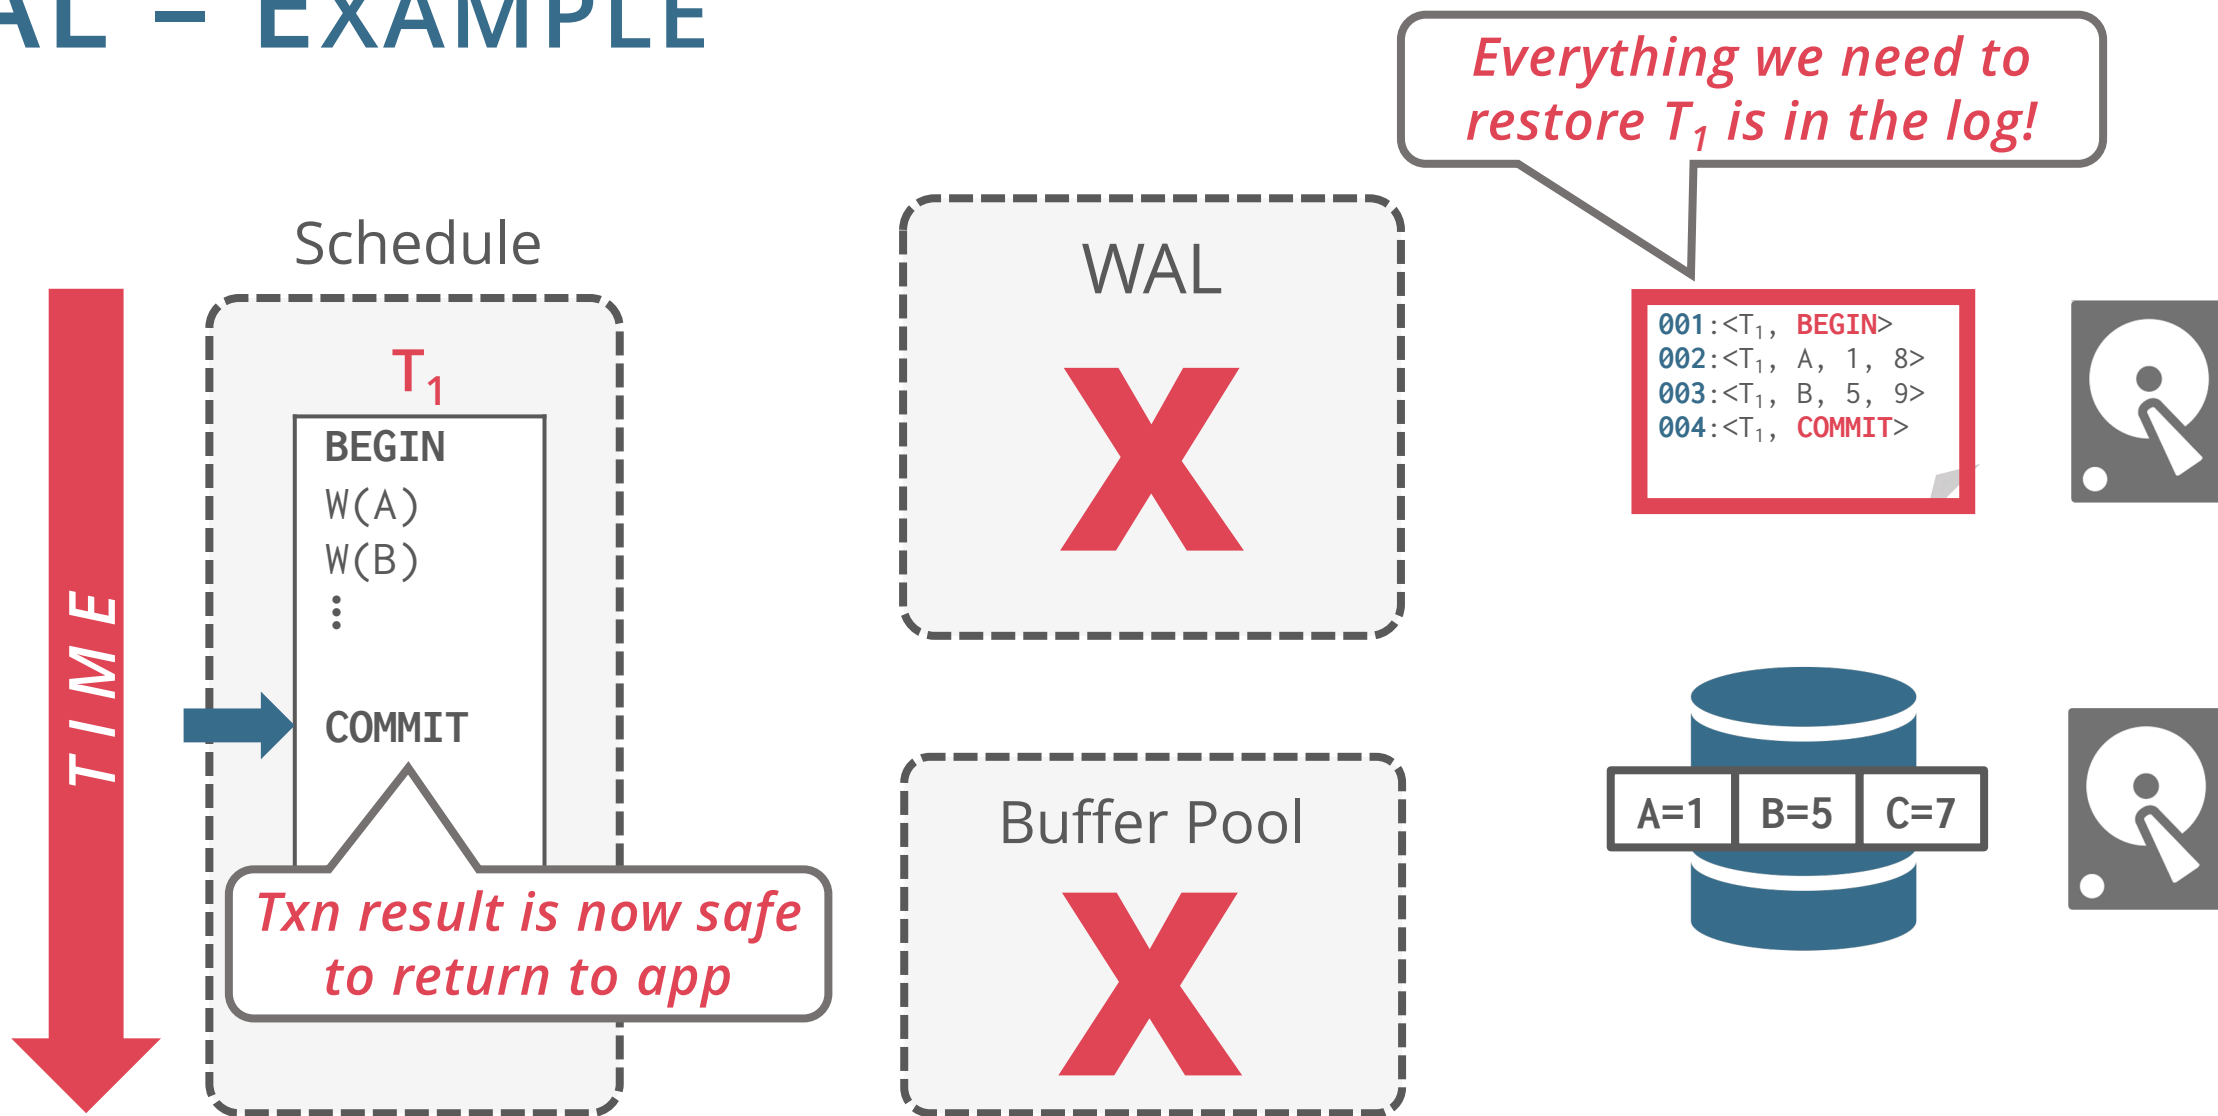
# WAL – EXAMPLE



# WAL – EXAMPLE



# WAL – EXAMPLE





# ARIES

Recovery algorithm developed at IBM Research in early 1990s

## Write-Ahead Logging

Any change is recorded in log on stable storage before the change is written to disk

Must use **STEAL** + **NO-FORCE** buffer pool policies

## Repeating History During Redo

On restart, retrace actions and restore database to exact state before crash

## Logging Changes During Undo

Log also undo actions to ensure action is not repeated in the event of repeated failures

# LOG RECORDS

Each log record has a unique **Log Sequence Number (LSN)**

LSNs are always increasing

Log record fields:

<i>LSN</i>	<i>Transaction ID</i>	<i>Type</i>
------------	-----------------------	-------------

Log Record Types:

**UPDATE** – consists of Object ID + Before Value (UNDO) + After Value (REDO)

**BEGIN / COMMIT / ABORT**

**TXN-END** – signifies the end of commit or abort

**CLR** (Compensation Log Record) – for UNDO actions

# LOG SEQUENCE NUMBERS

Name	Where	Definition
<b>flushedLSN</b>	Memory	Last LSN in log on disk
<b>pageLSN</b>	$page_x$	Newest update to $page_x$
<b>recLSN</b>	$page_x$	Oldest update to $page_x$ since it was last flushed
<b>lastLSN</b>	$T_i$	Latest record of txn $T_i$
<b>MasterRecord</b>	Disk	LSN of latest checkpoint

# WRITING LOG RECORDS

Each data page contains a **pageLSN**

The *LSN* of the most recent log record for an update to that page

Update the **pageLSN** every time a txn modifies a record in the page

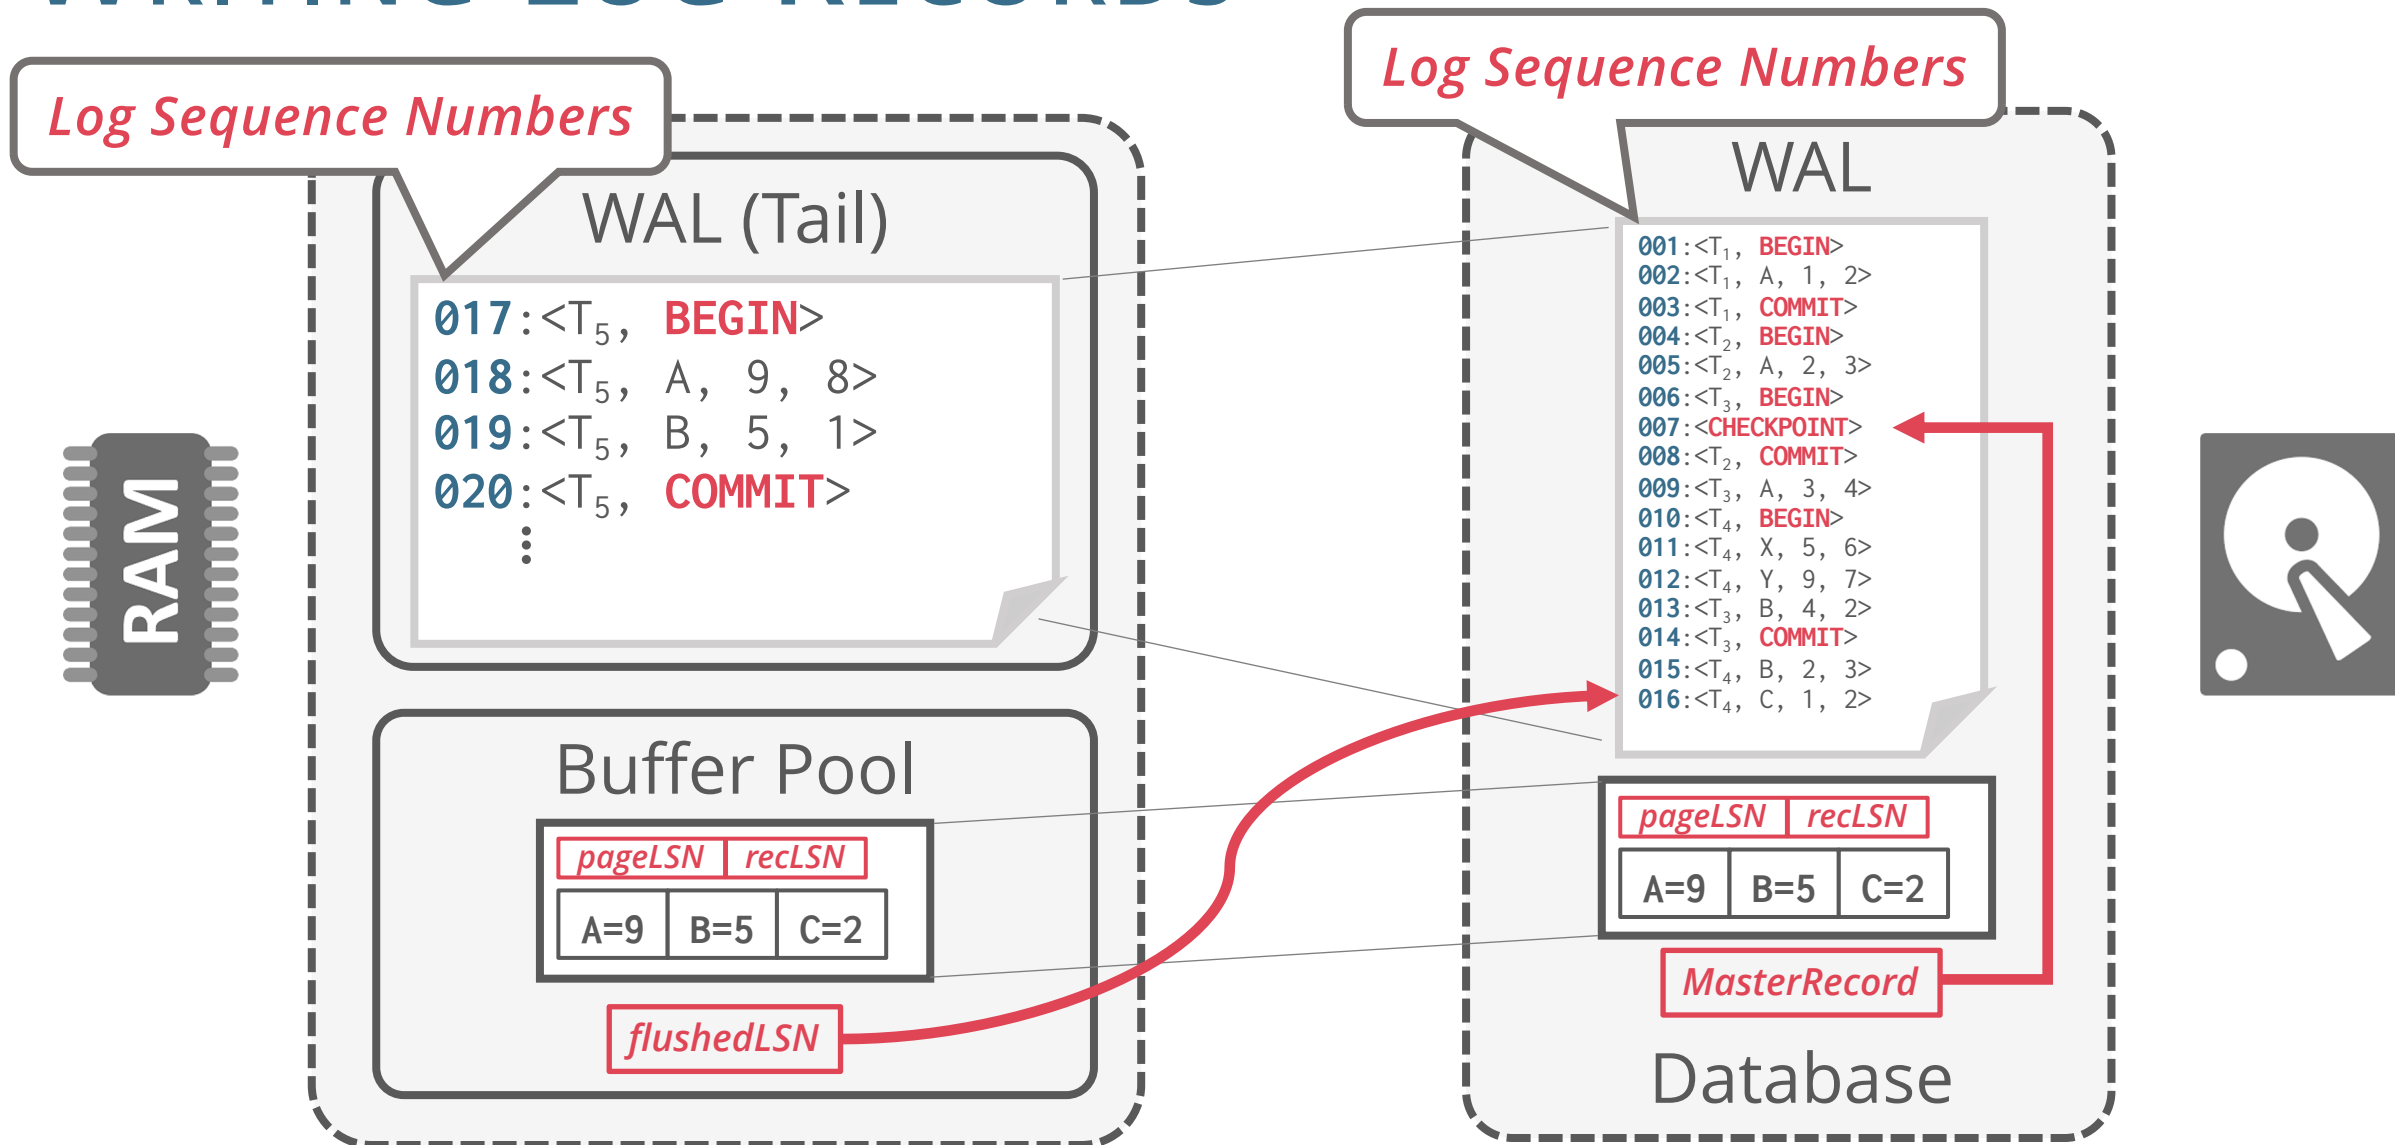
System keeps track of **flushedLSN**

The max *LSN* flushed so far

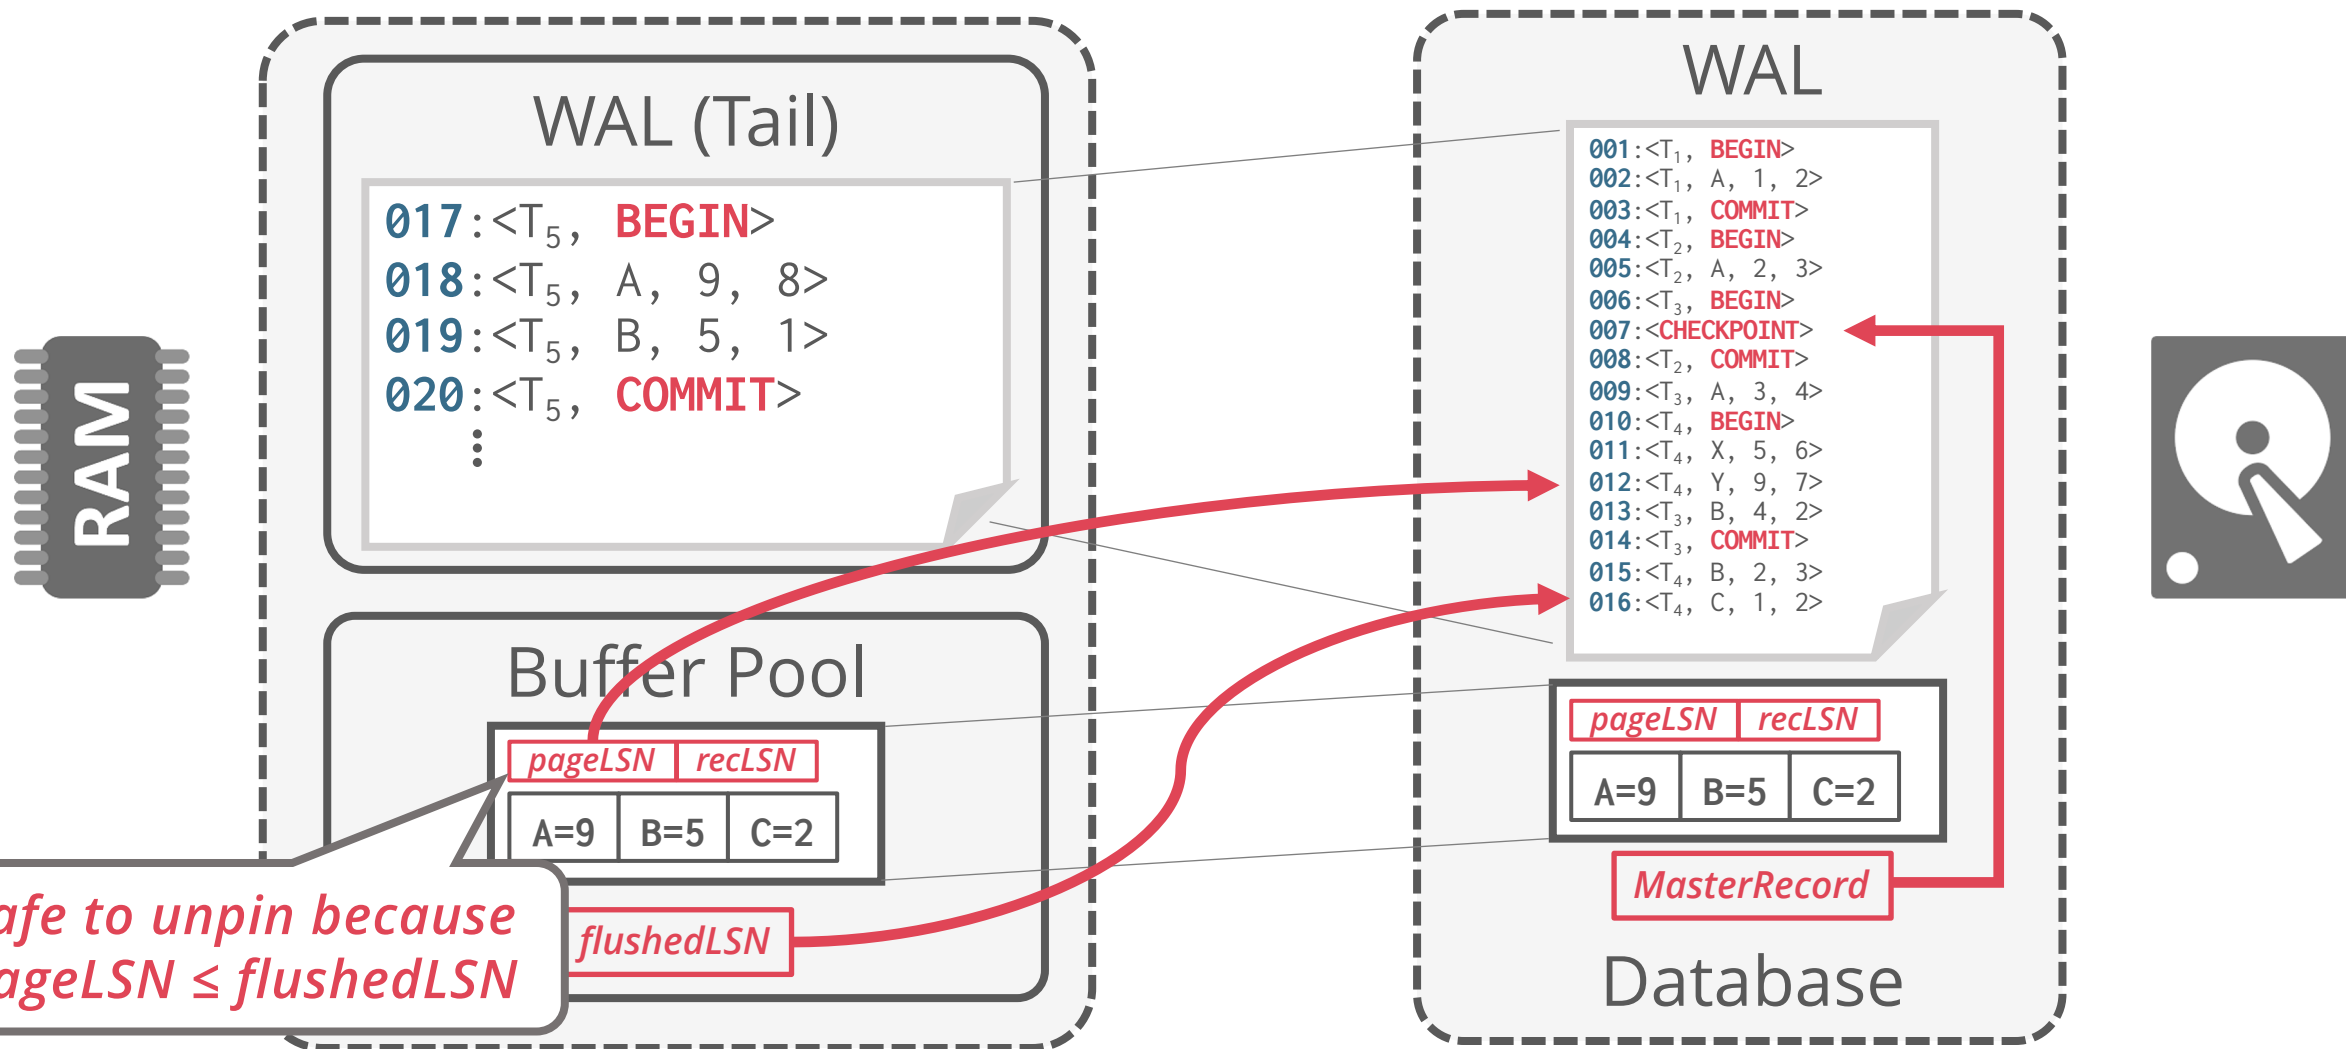
Update the **flushedLSN** in memory every time DBMS writes out the WAL buffer to disk

Before a page **X** can be written to disk, we must flush log  
at least to the point where  $\text{pageLSN}_x \leq \text{flushedLSN}$

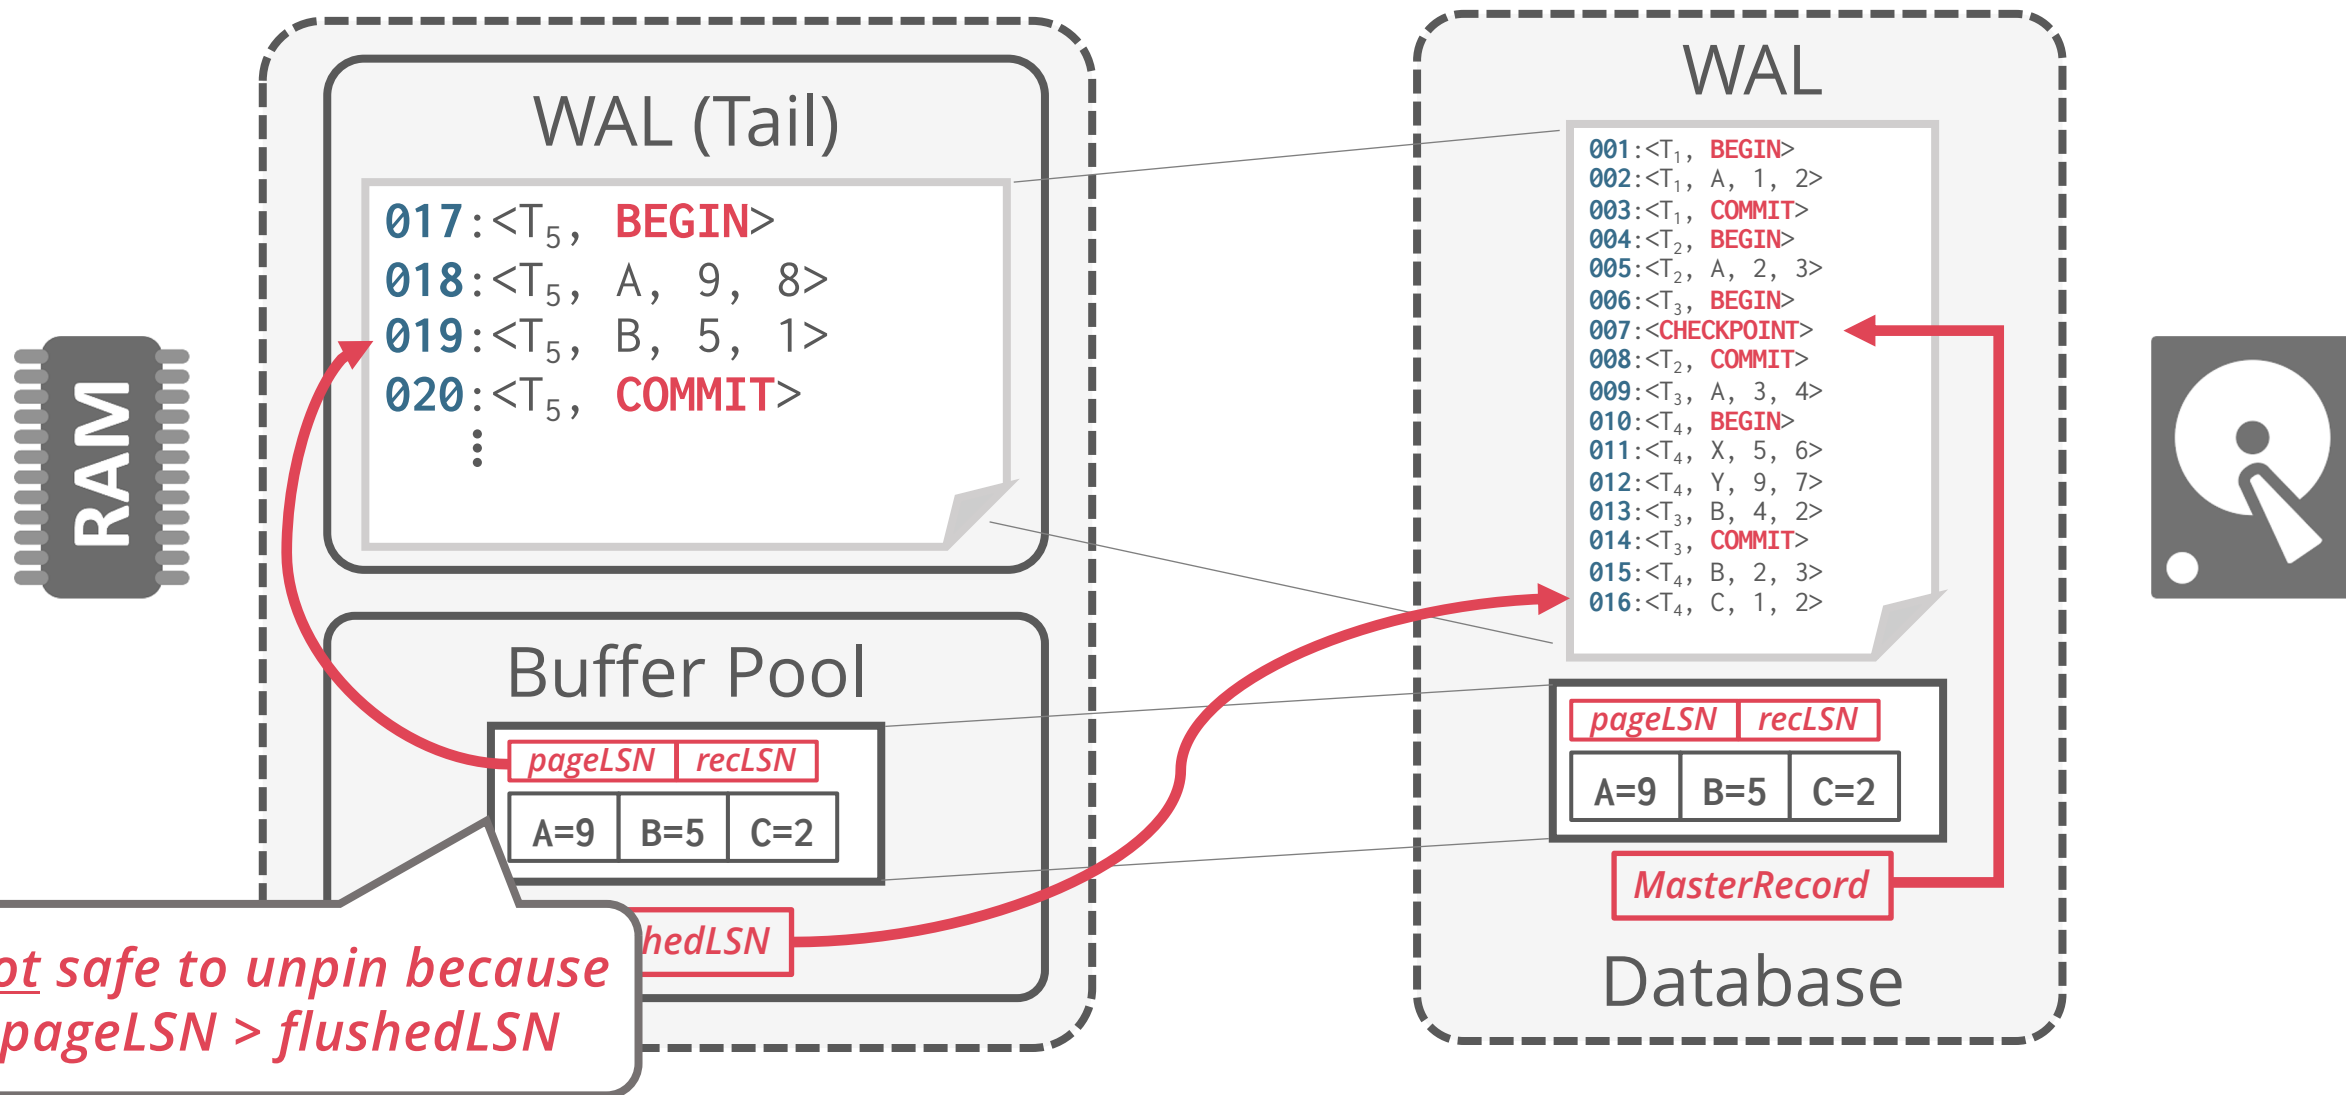
# WRITING LOG RECORDS



# WRITING LOG RECORDS



# WRITING LOG RECORDS



# NORMAL EXECUTION

Each txn invokes a sequence of reads and writes,  
followed by commit or abort

Assumptions in this lecture:

- All log records fit within a single page

- Disk writes are atomic

- Single-versioned tuples with Strict 2PL

- STEAL** + **NO-FORCE** buffer management with WAL



# TRANSACTION COMMIT

Write **COMMIT** record to log

All log records up to txn's **COMMIT** record are flushed to disk

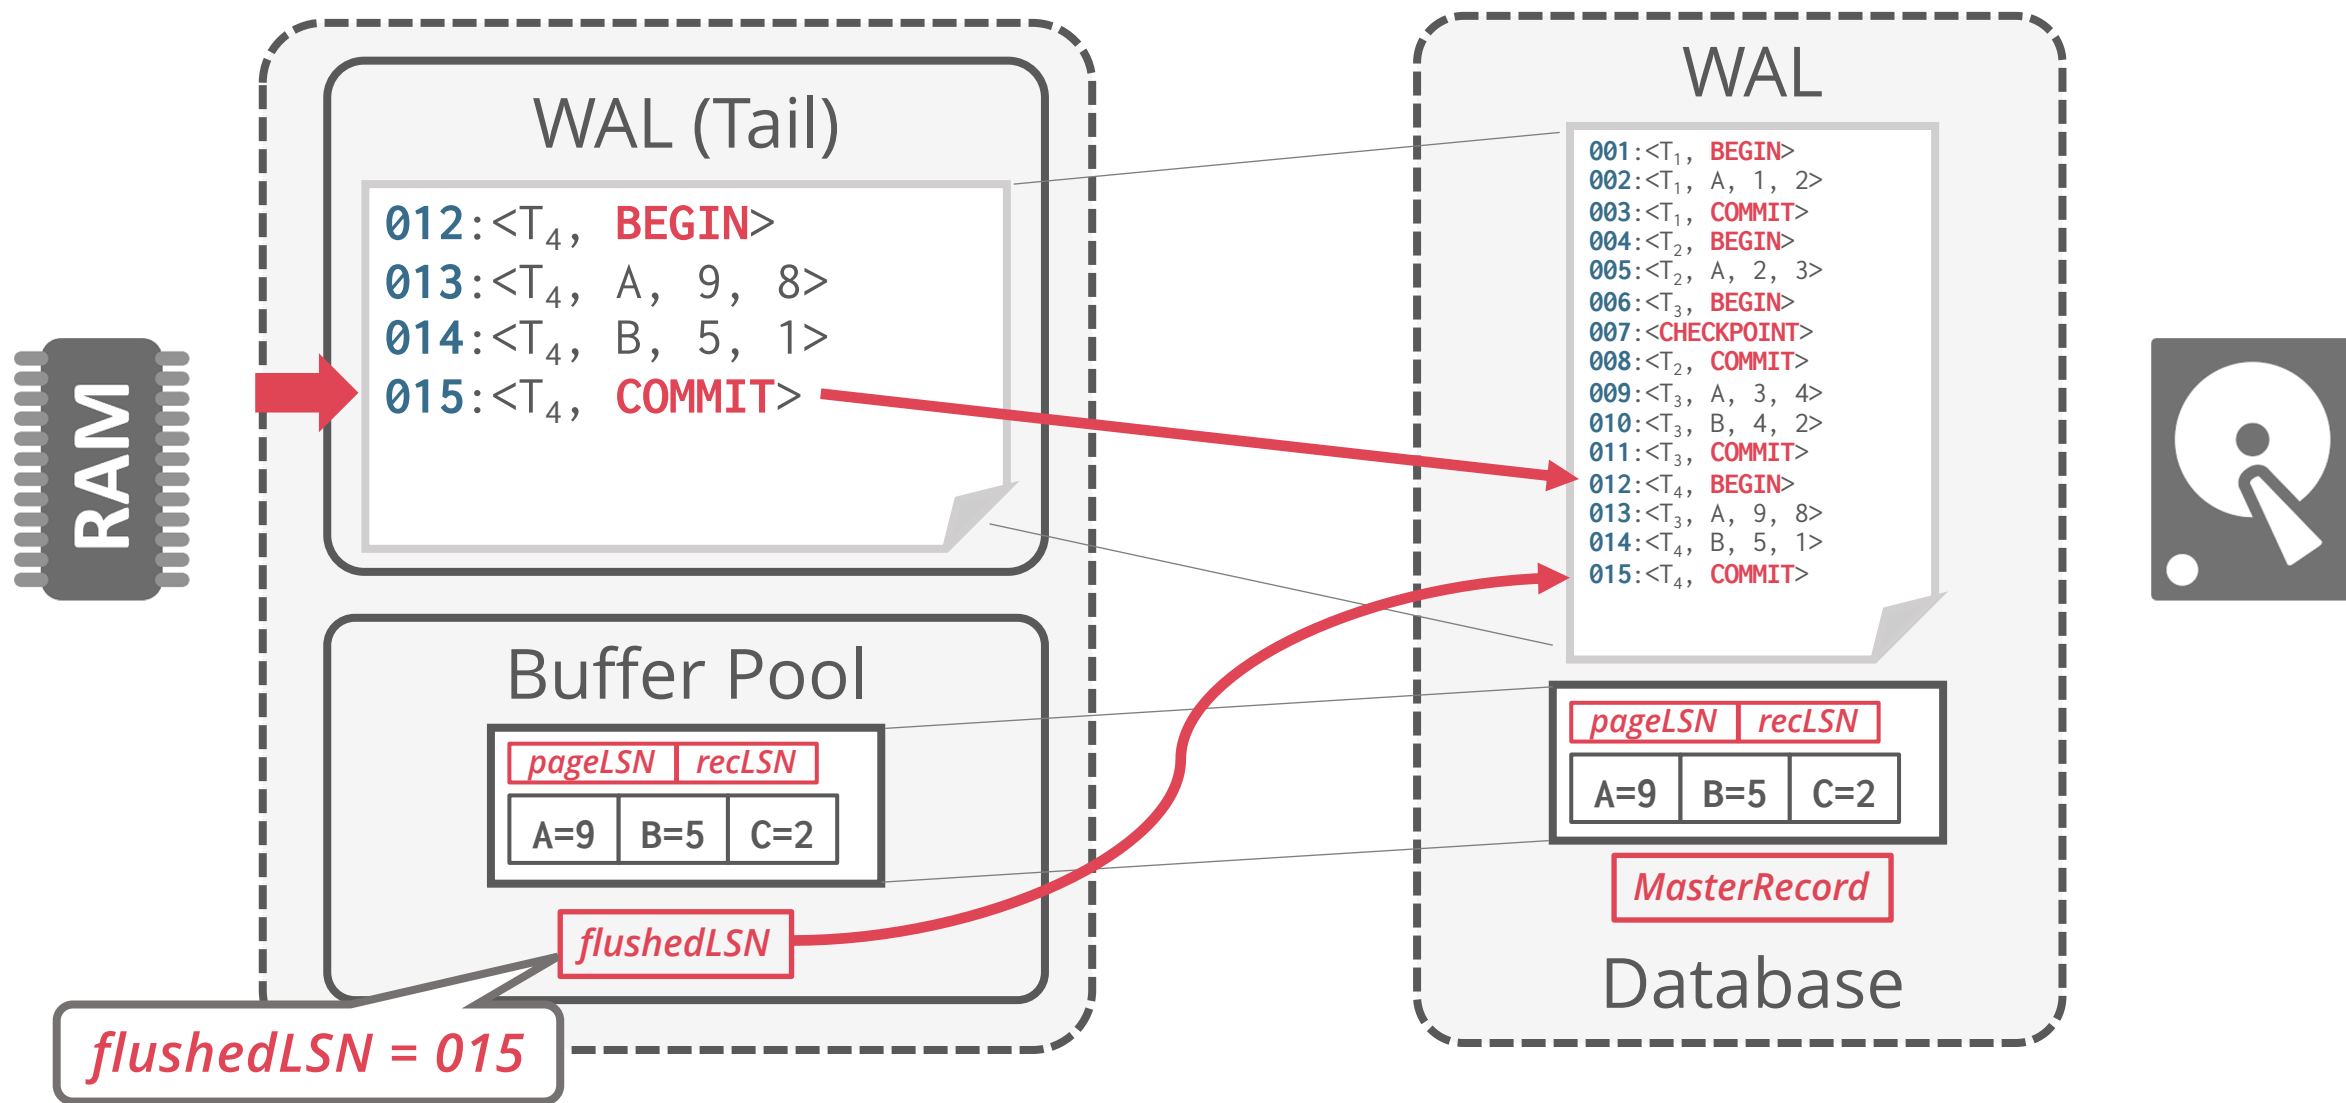
Note that log flushes are sequential, synchronous writes to disk

Many log records per log page

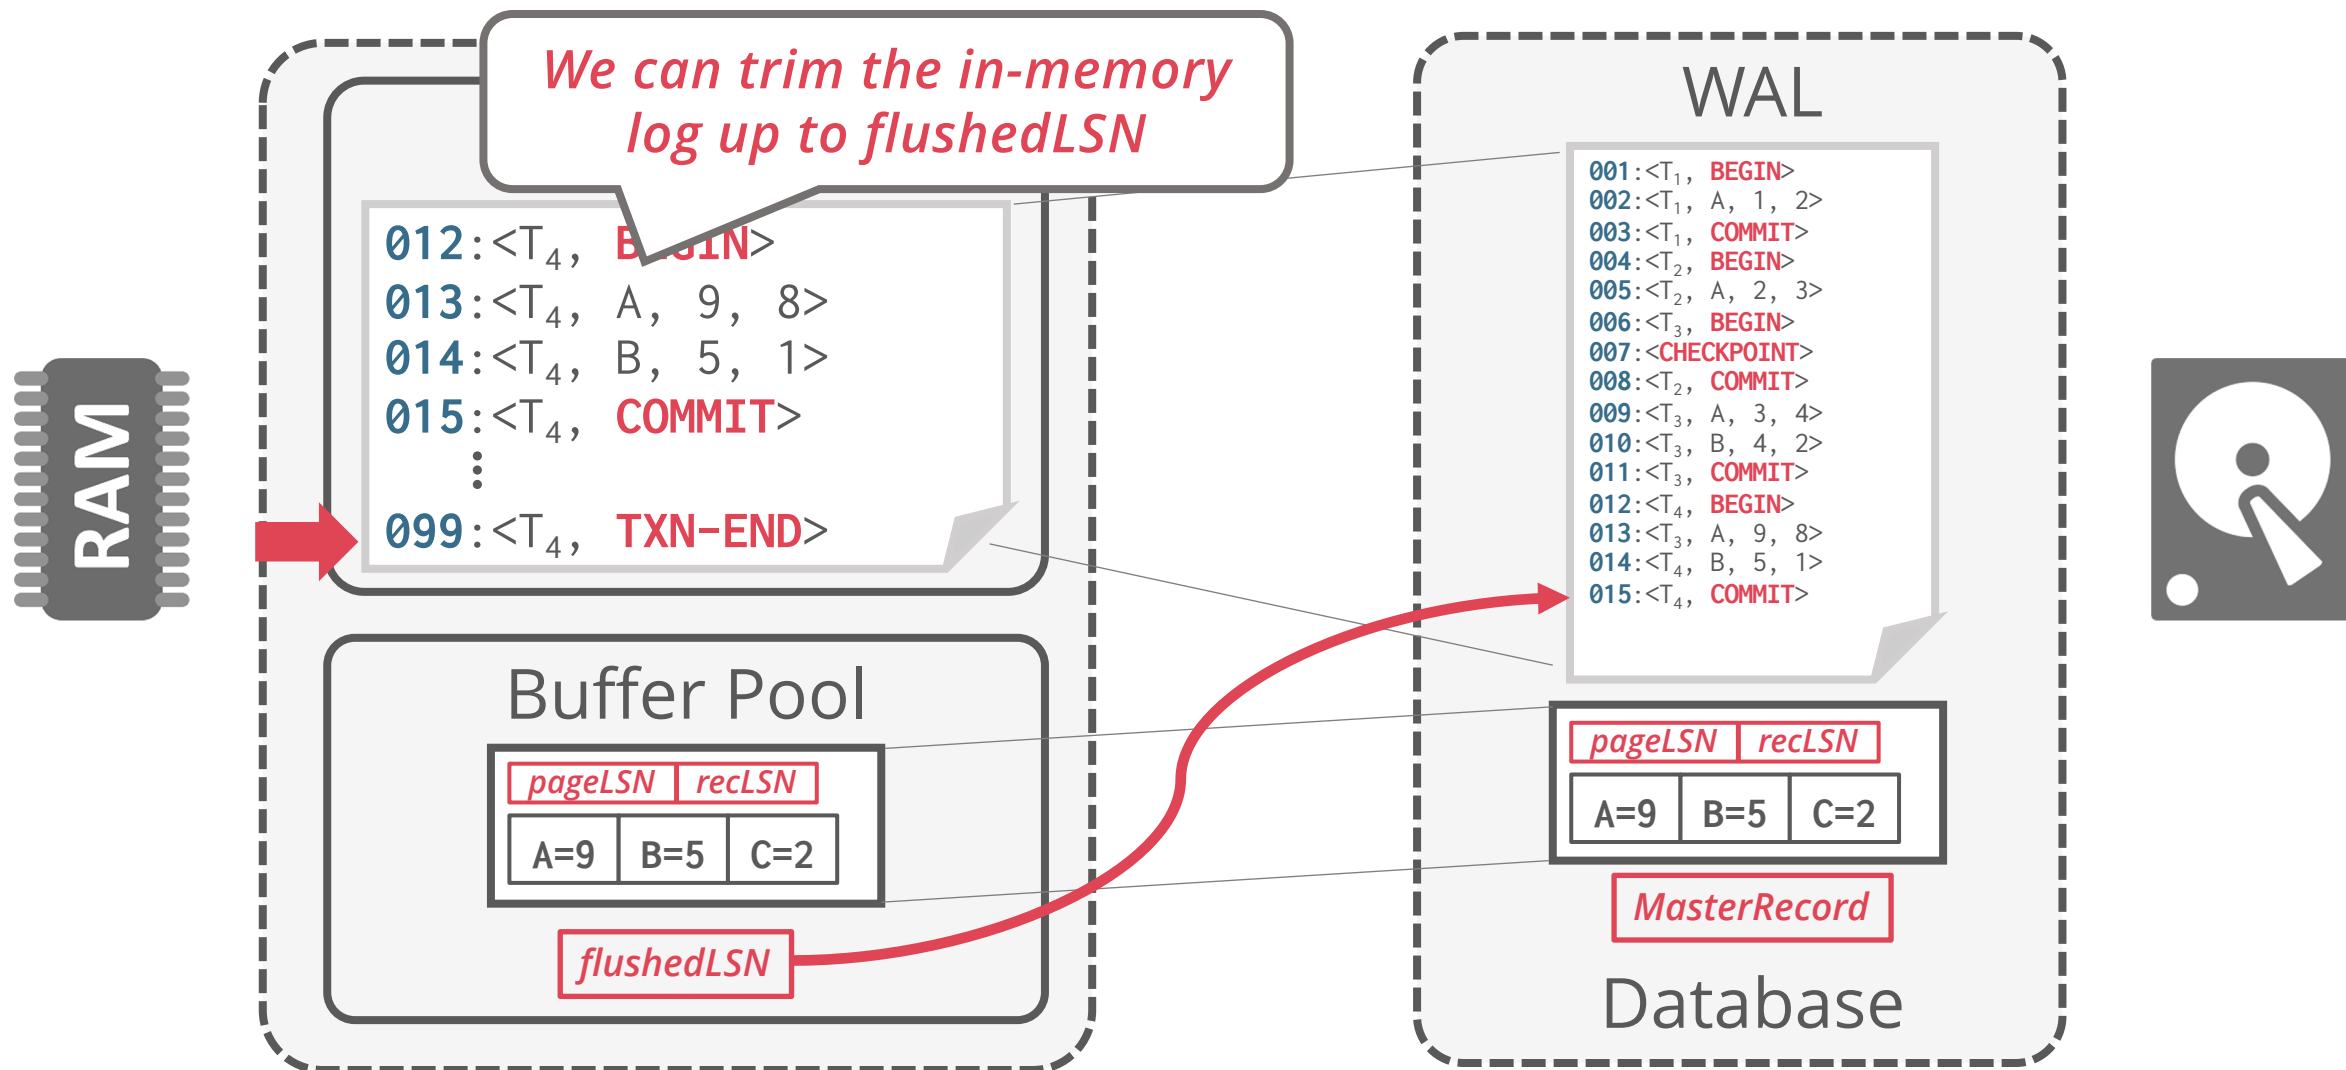
When the commit succeeds, write a special **TXN-END** record to log

This does not need to be flushed immediately

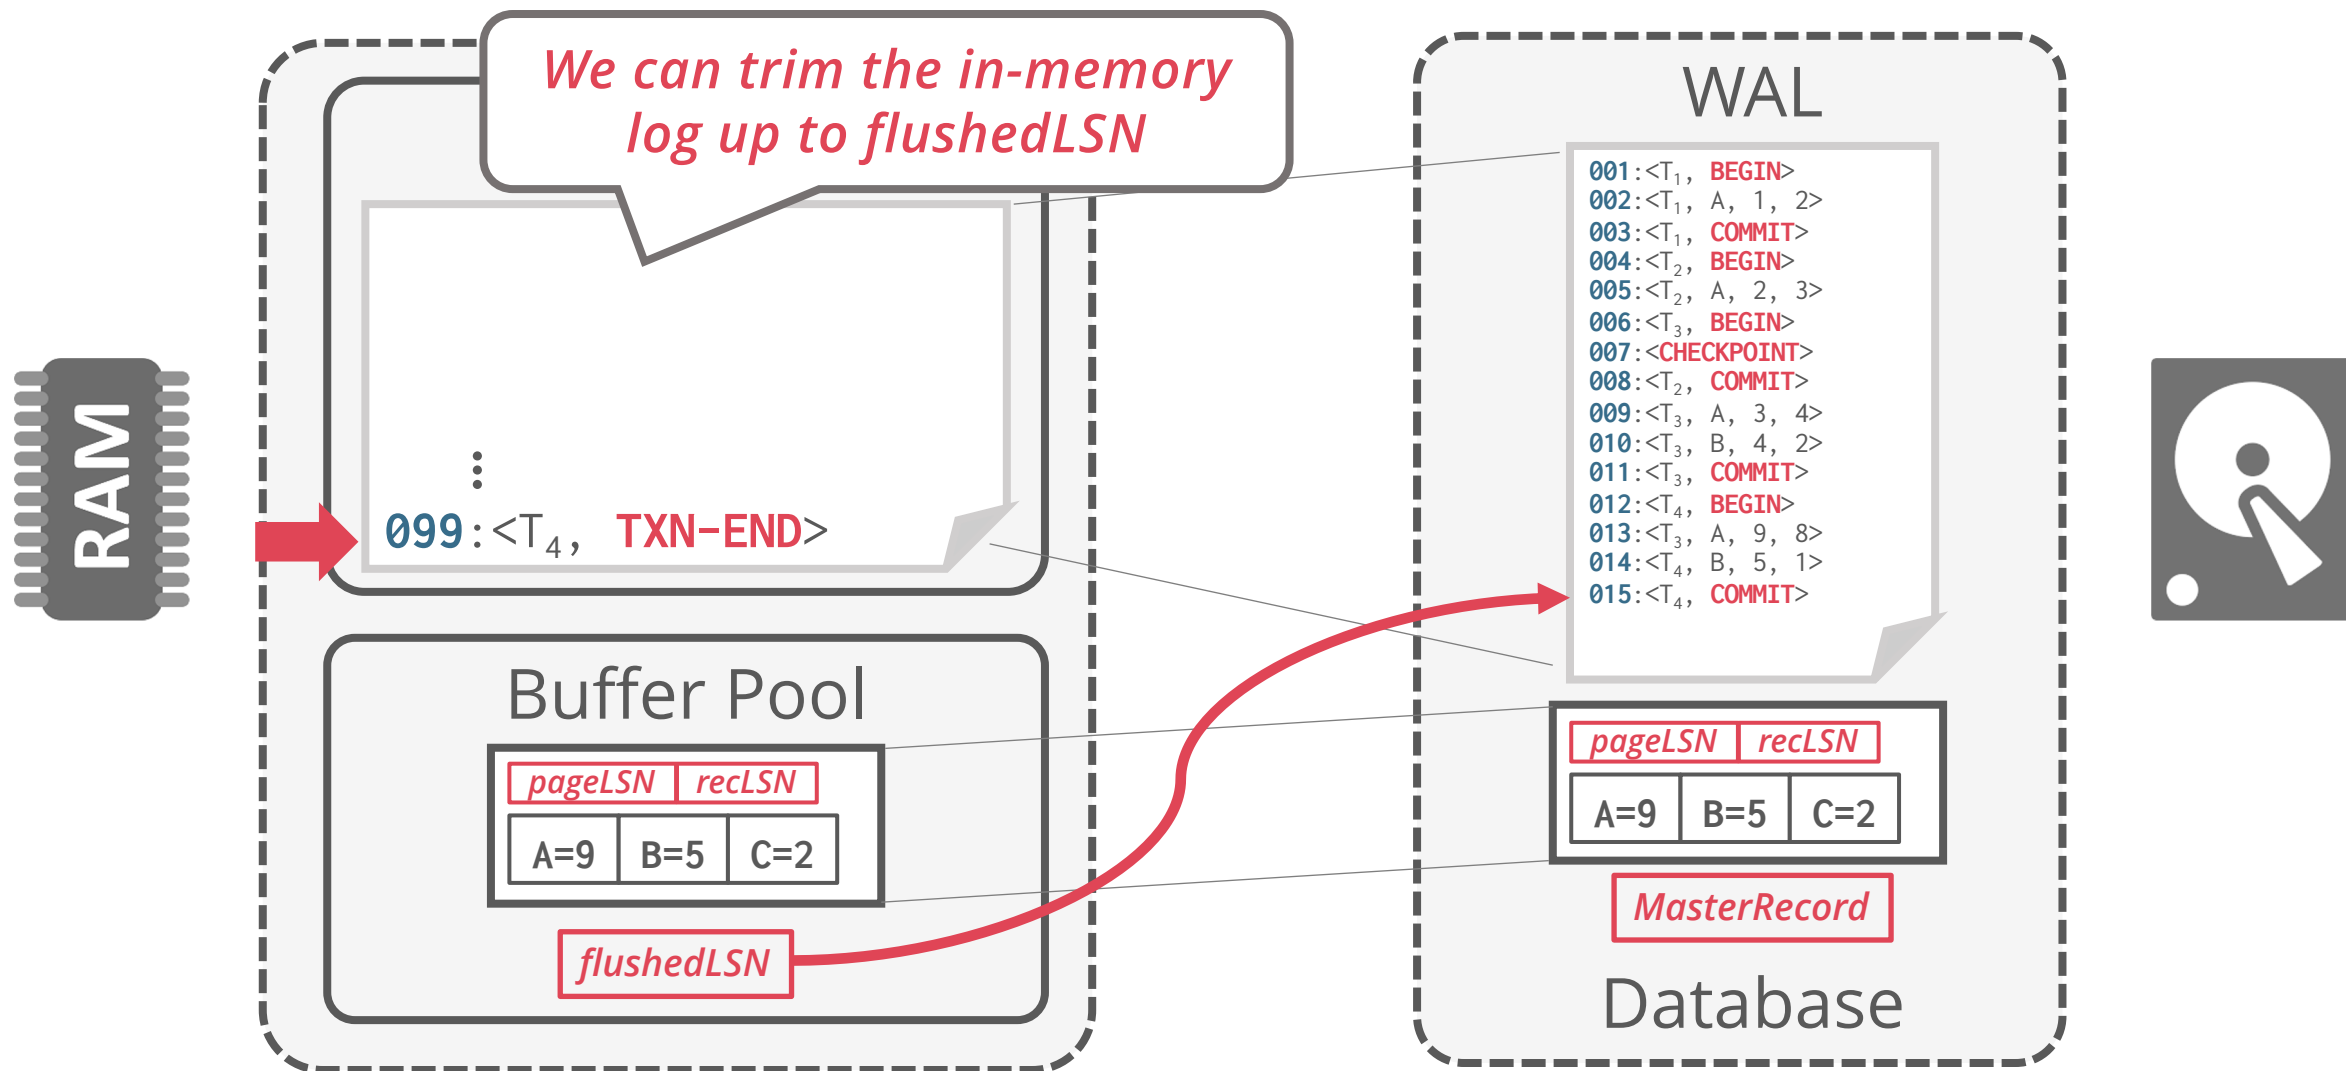
# TRANSACTION COMMIT



# TRANSACTION COMMIT



# TRANSACTION COMMIT



# TRANSACTION ABORT

We need to add another field to our log records

**prevLSN**: The previous *LSN* for the txn

Maintains a linked list for each txn that makes it easy to walk through its records

We want to “play back” the log in reverse order, **UNDO**ing updates

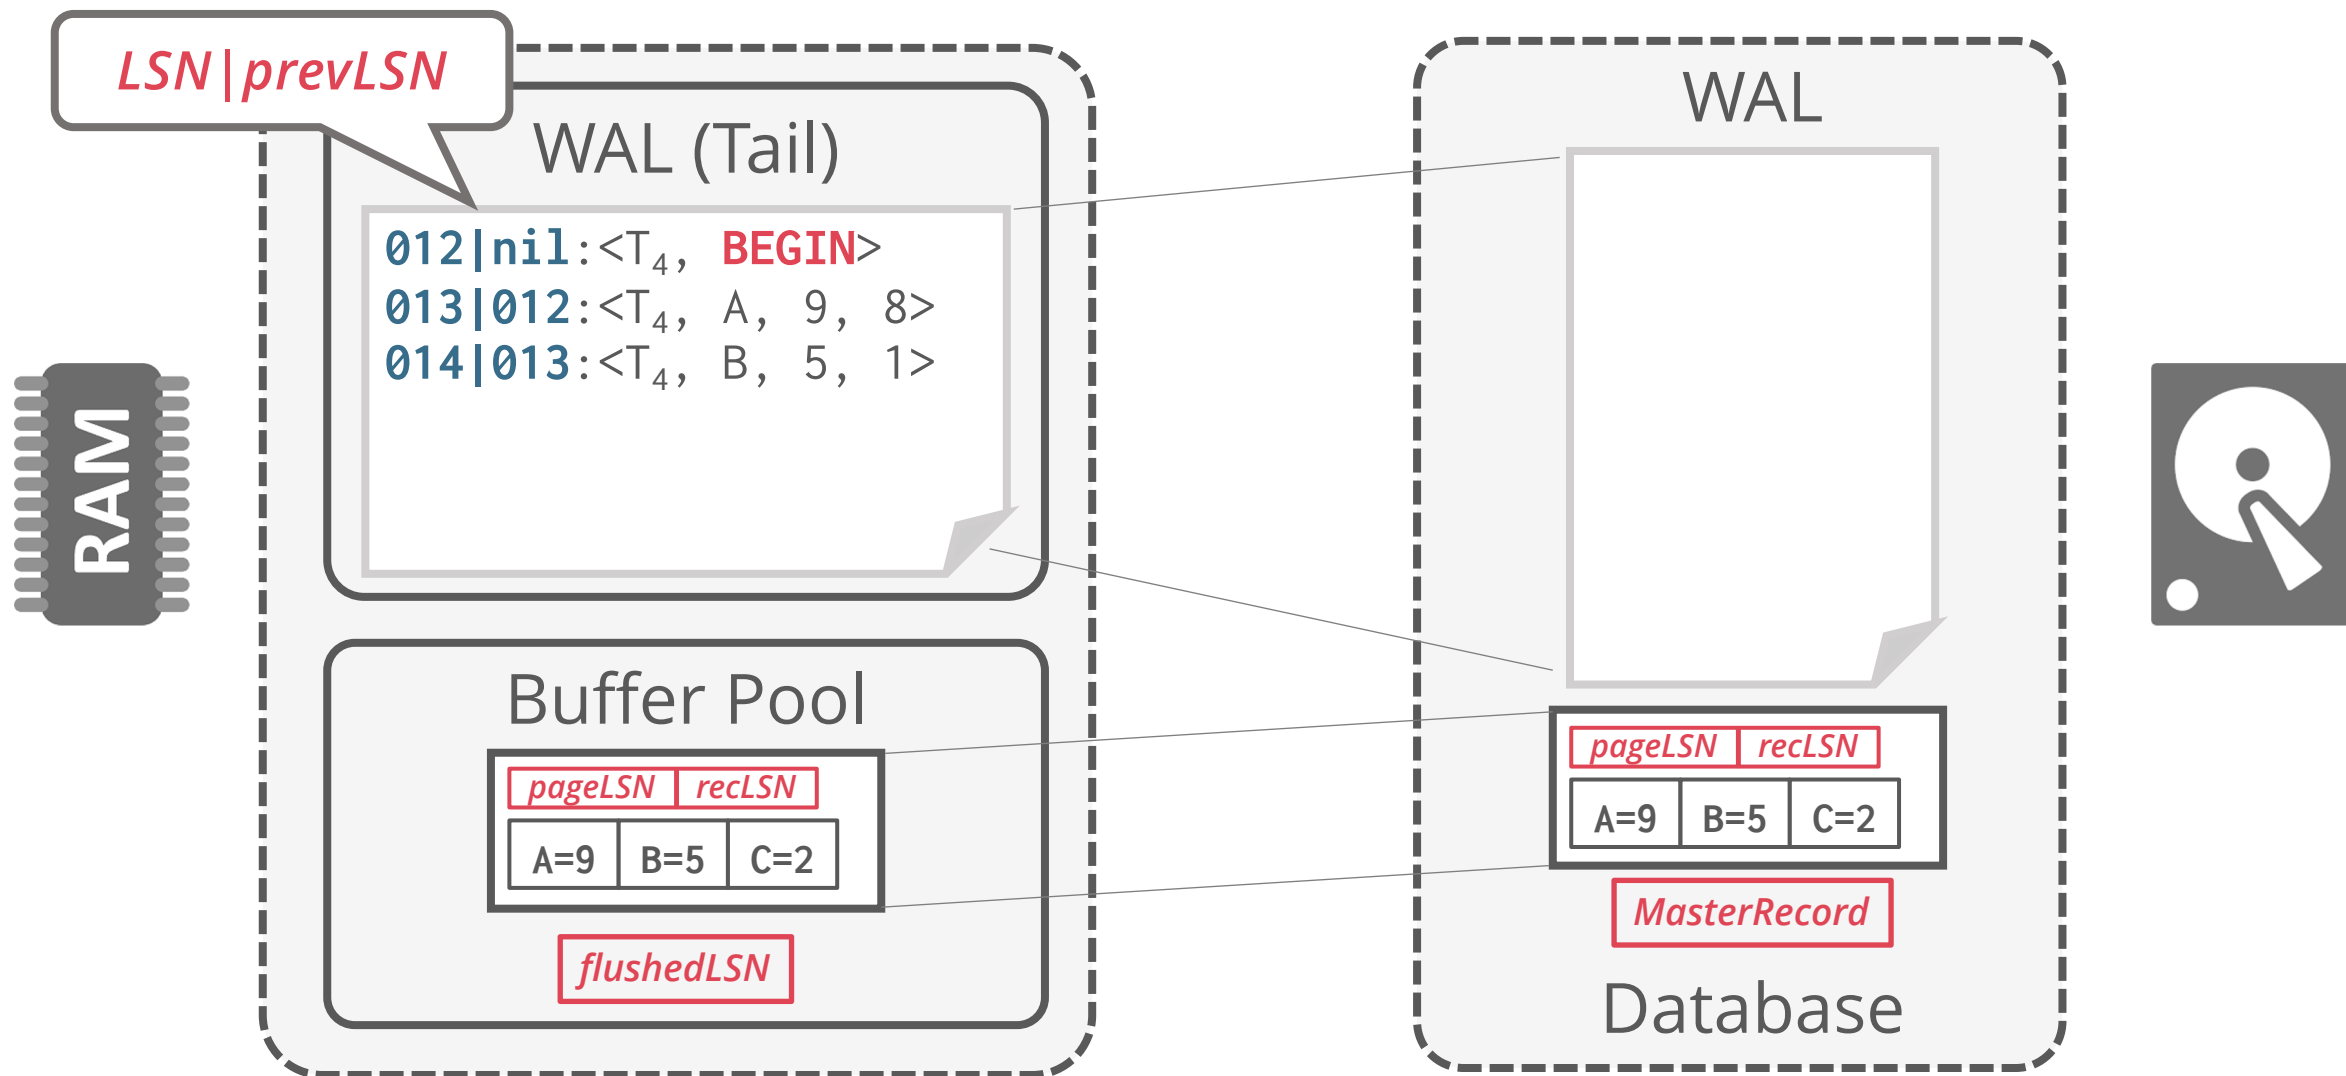
Get **lastLSN** of txn stored in the active txn table (more details later)

Can follow chain of log records backward via the **prevLSN** field

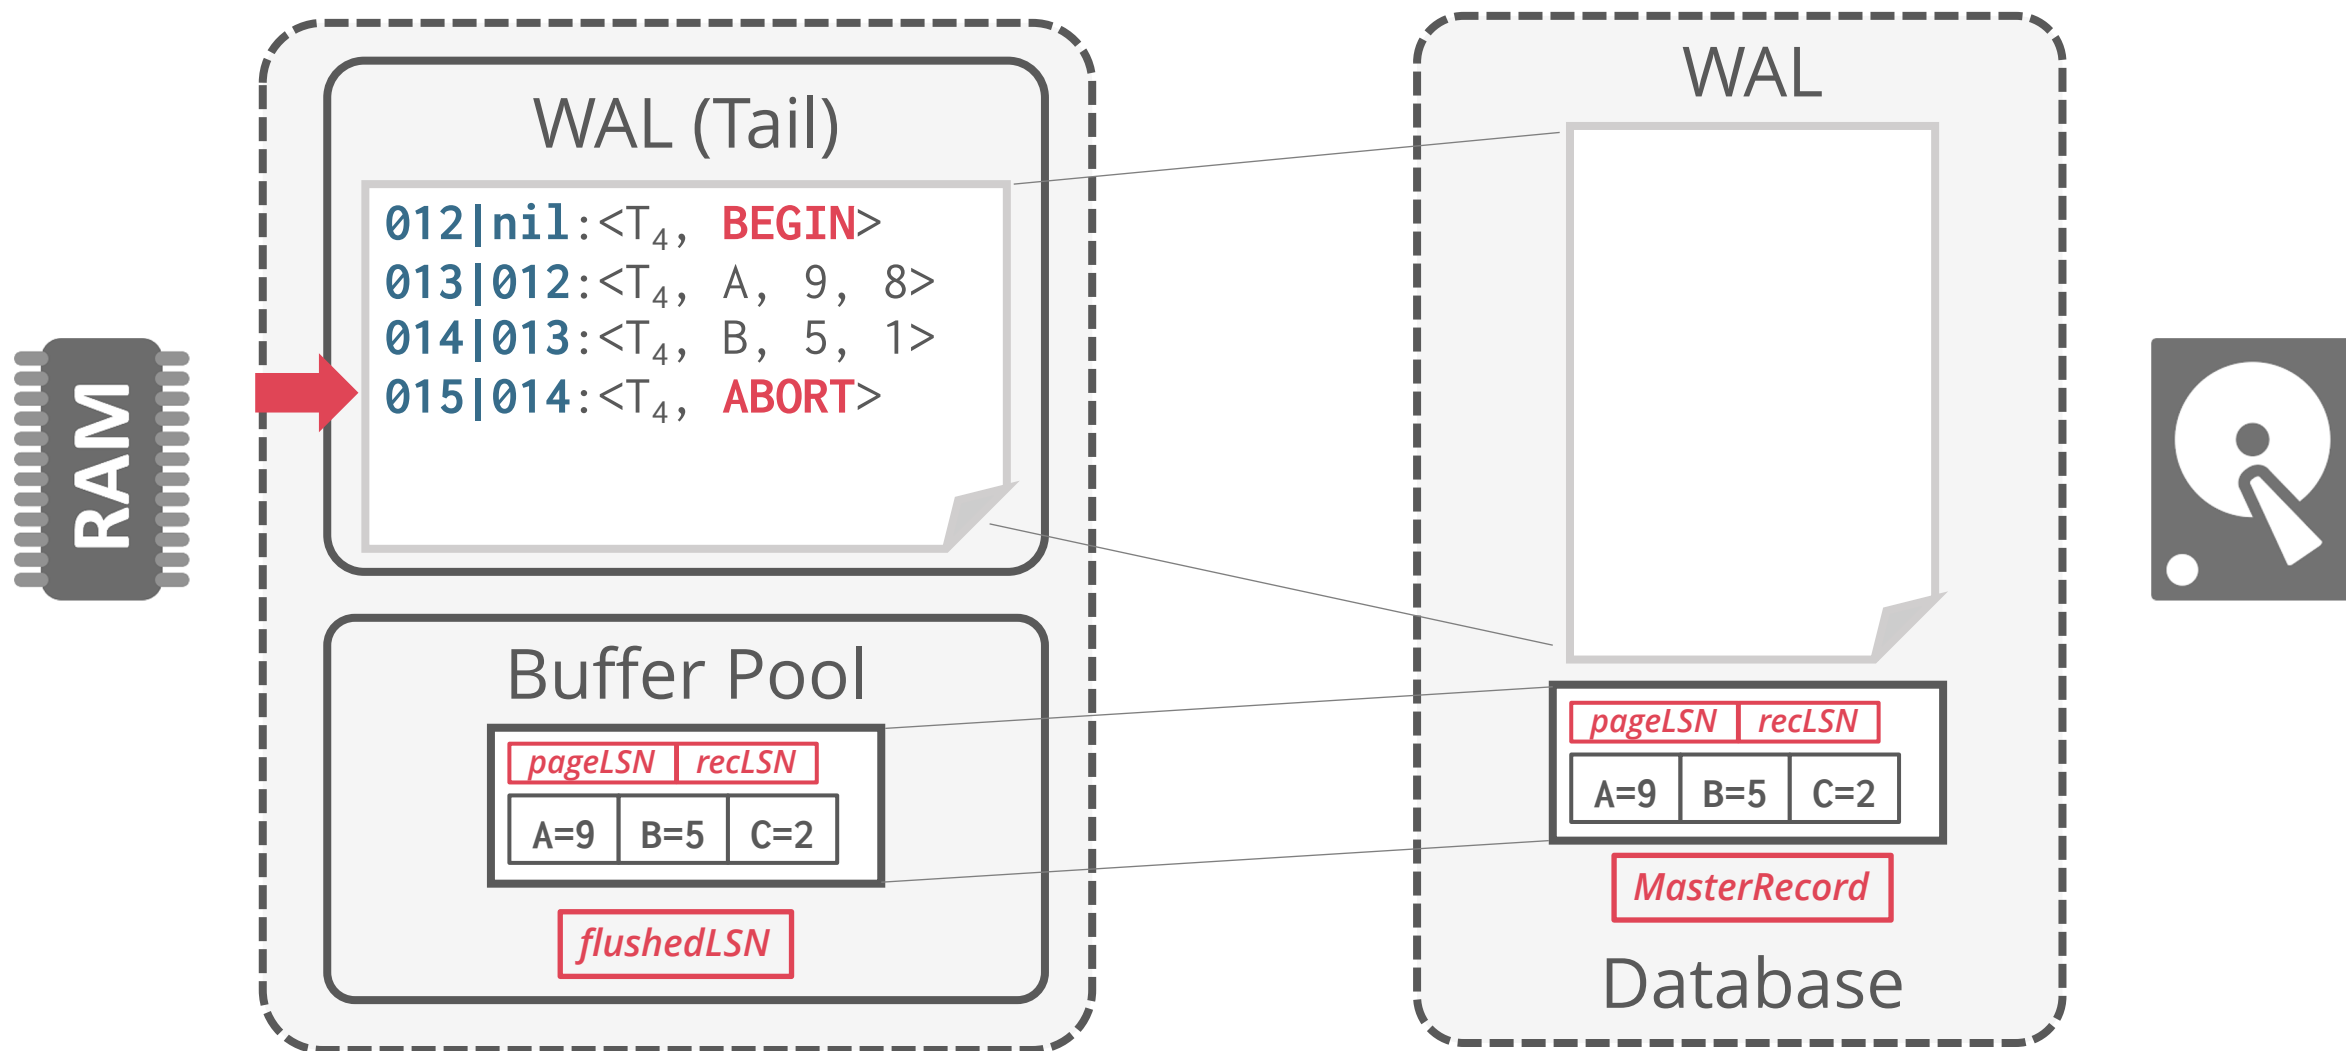
Before starting UNDO, write an **ABORT** log record

*For recovering from crash during UNDO!*

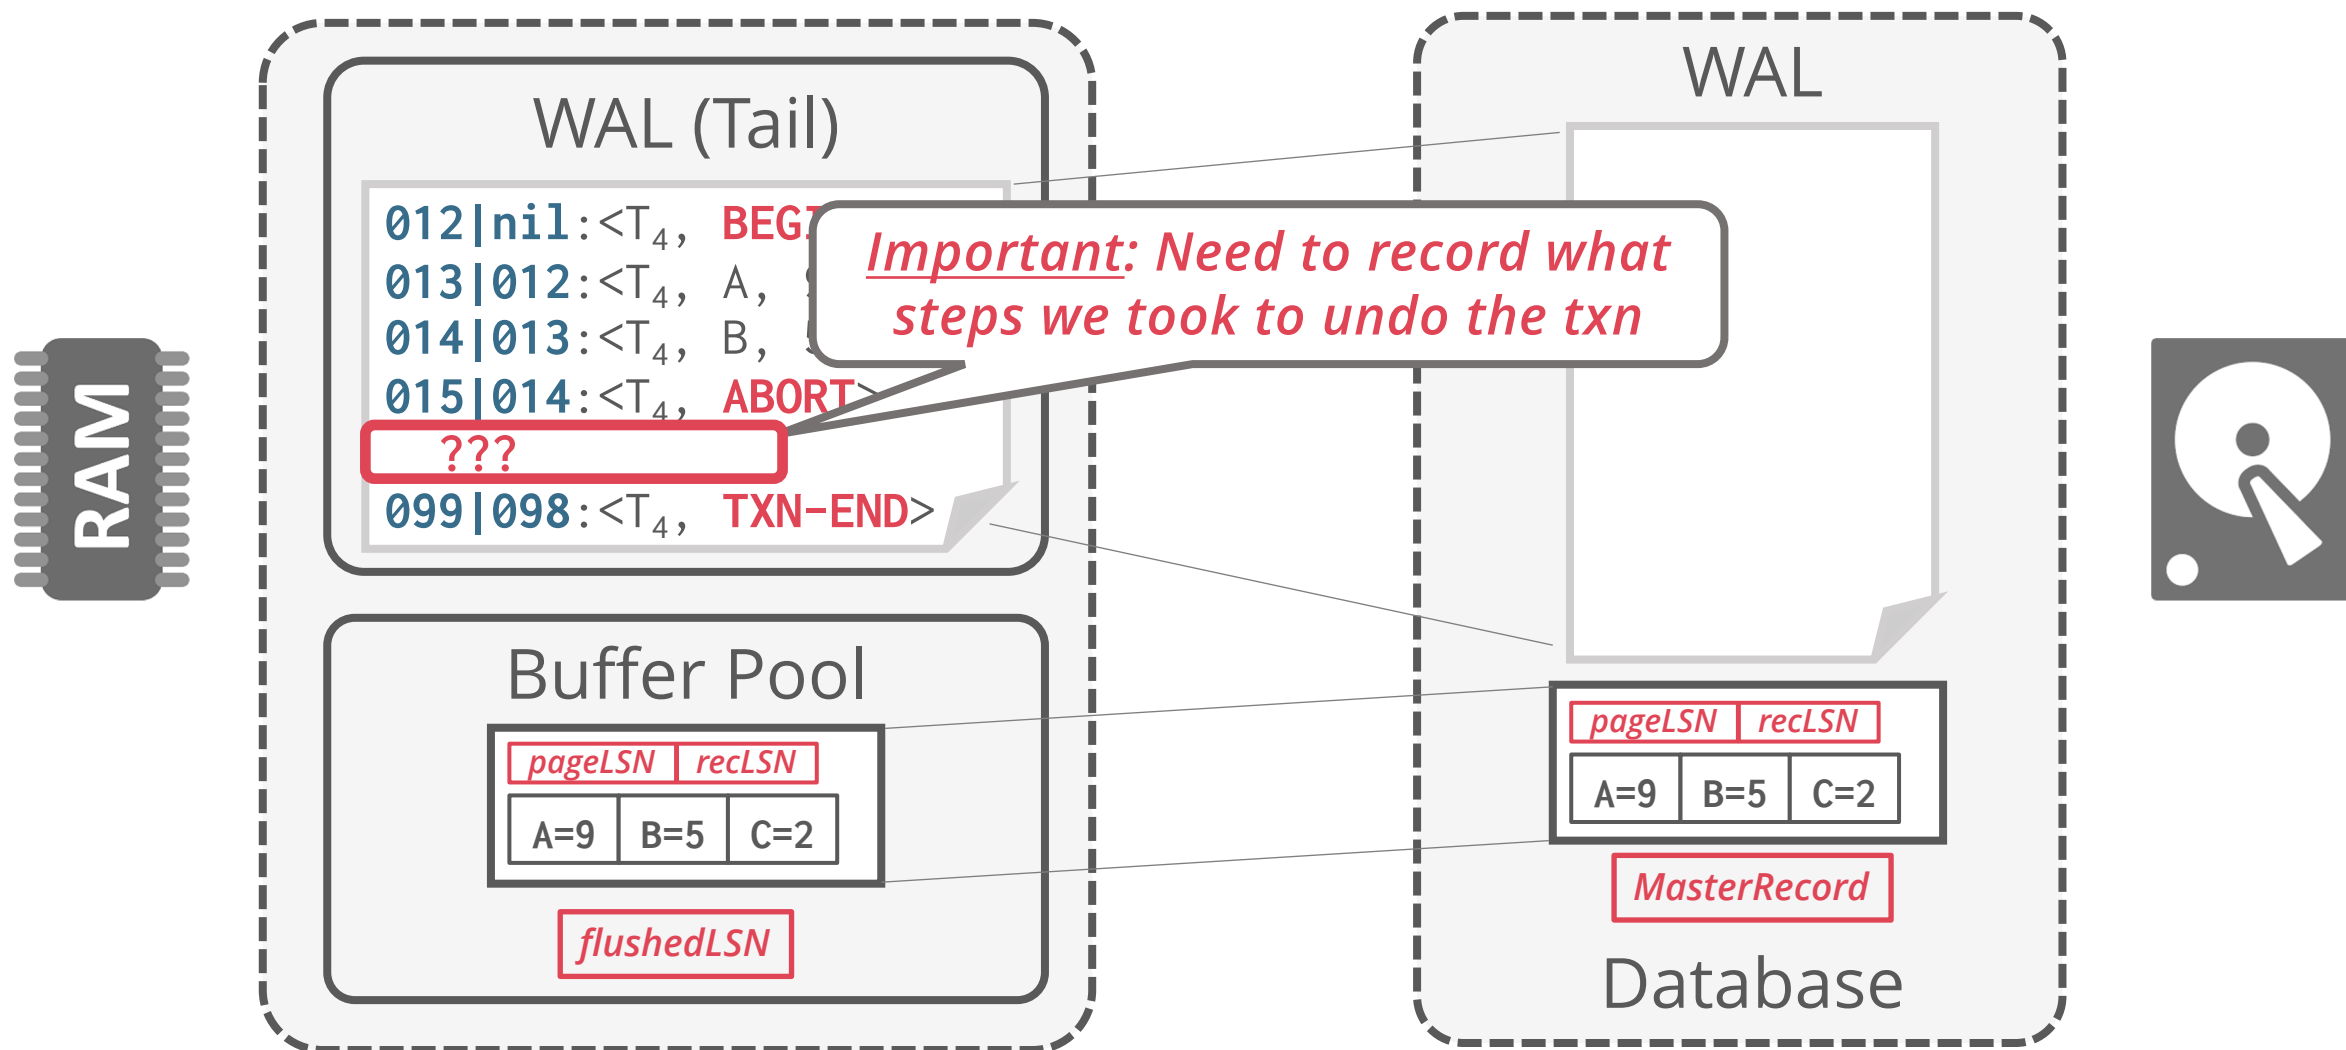
# TRANSACTION ABORT



# TRANSACTION ABORT



# TRANSACTION ABORT






# COMPENSATION LOG RECORDS

A CLR describes the actions taken to undo the actions of a previous update record

It has all the fields of an update log record plus the **undoNext** pointer (the next-to-be-undone LSN)

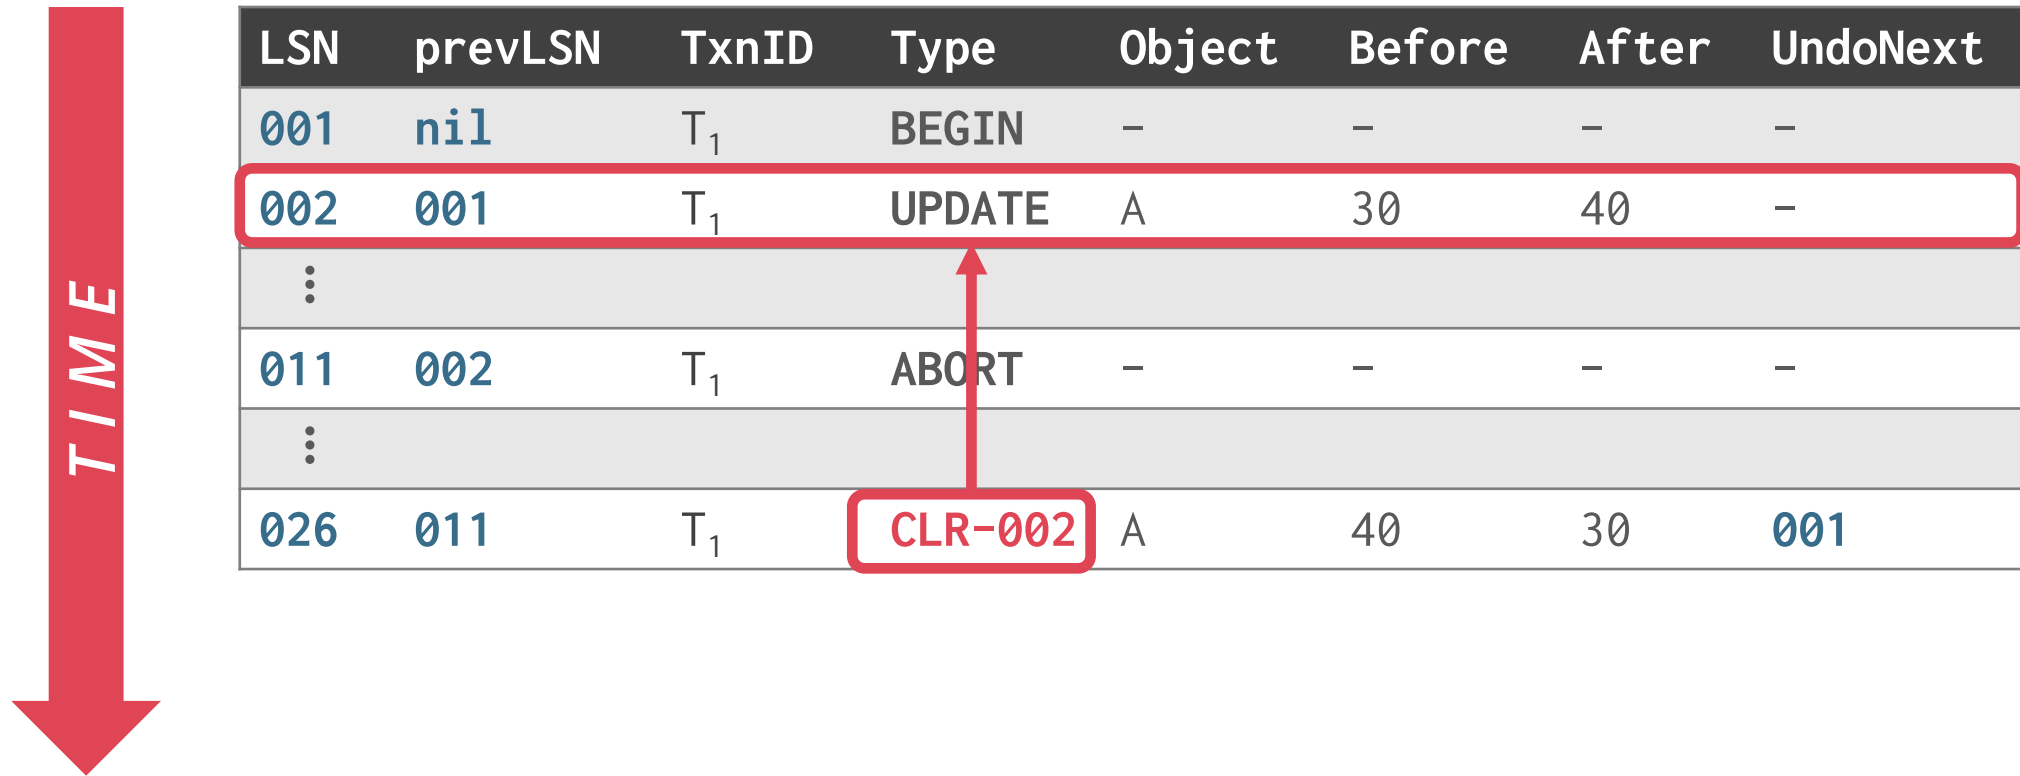
*CLRs* are added to log like any other record

# TRANSACTION ABORT – CLR EXAMPLE



LSN	prevLSN	TxnID	Type	Object	Before	After	UndoNext
001	nil	T <sub>1</sub>	BEGIN	–	–	–	–
002	001	T <sub>1</sub>	UPDATE	A	30	40	–
⋮							
011	002	T <sub>1</sub>	ABORT	–	–	–	–

# TRANSACTION ABORT – CLR EXAMPLE




A diagram illustrating a transaction log (LSN, prevLSN, TxnID, Type, Object, Before, After, UndoNext) during a transaction abort. A large red arrow on the left points downwards, labeled "TIME". The log entries are as follows:

LSN	prevLSN	TxnID	Type	Object	Before	After	UndoNext
001	nil	T <sub>1</sub>	BEGIN	–	–	–	–
002	001	T <sub>1</sub>	UPDATE	A	30	40	–
⋮							
011	002	T <sub>1</sub>	ABORT	–	–	–	–
⋮							
026	011	T <sub>1</sub>	CLR-002	A	40	30	001

The entry at LSN 002 (UPDATE) is highlighted with a red box. A red arrow points from the entry at LSN 026 (CLR-002) to the entry at LSN 002 (UPDATE), indicating the undo operation.

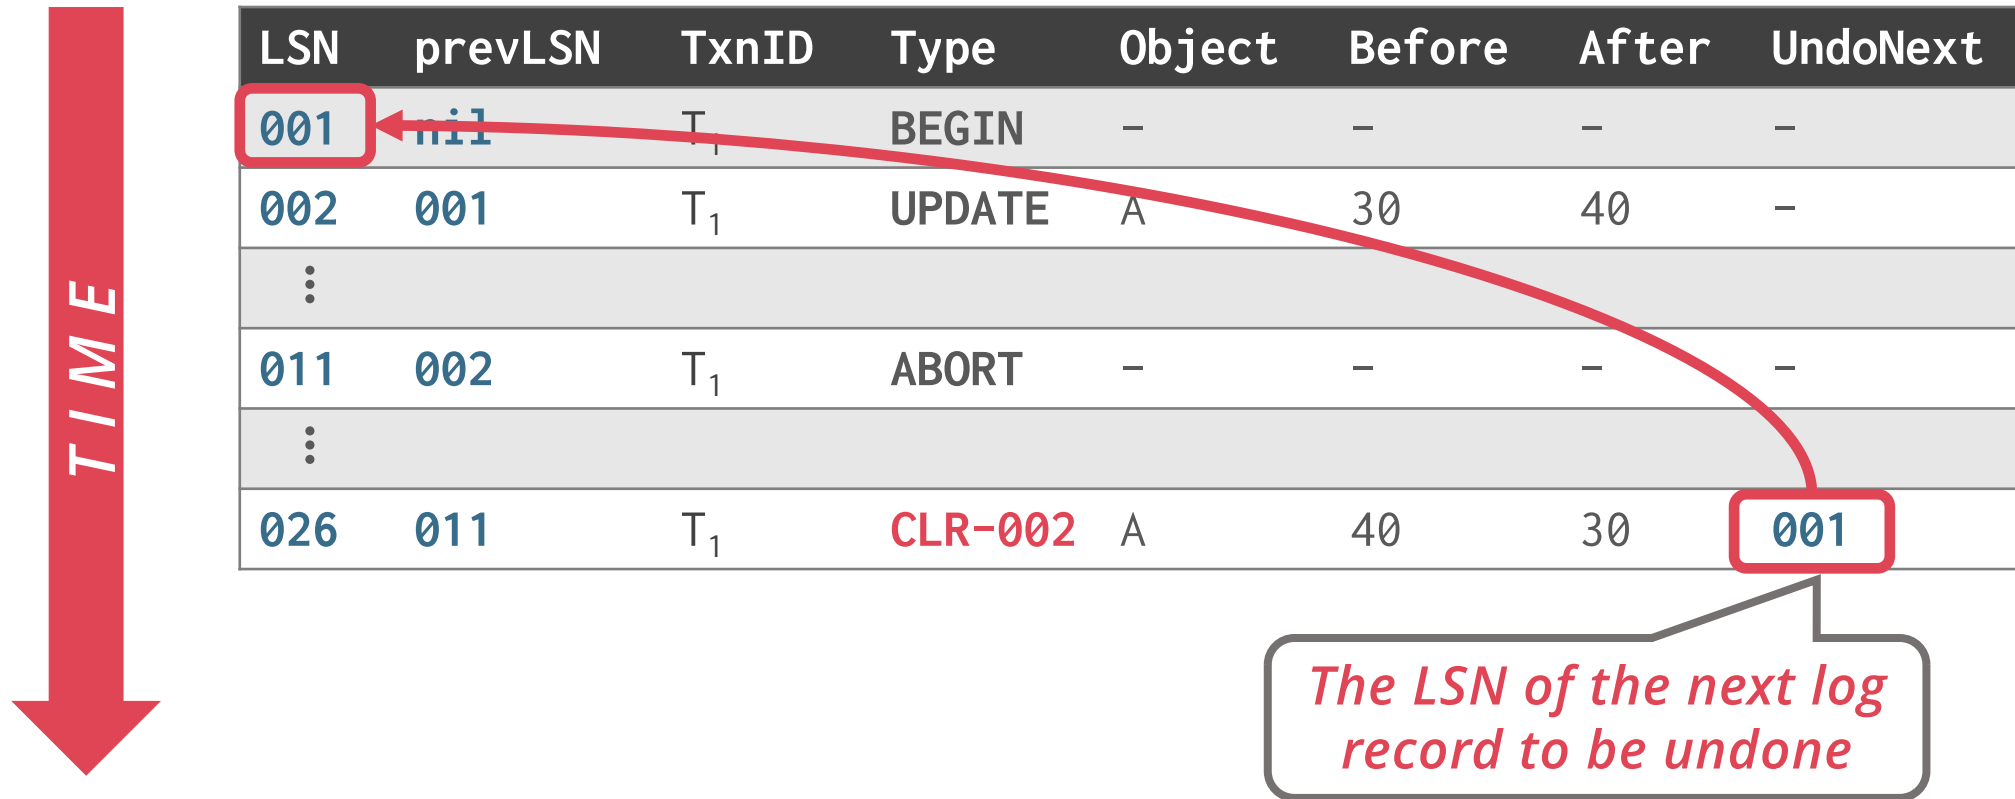
# TRANSACTION ABORT – CLR EXAMPLE




LSN	prevLSN	TxnID	Type	Object	Before	After	UndoNext
001	nil	T <sub>1</sub>	BEGIN	–	–	–	–
002	001	T <sub>1</sub>	UPDATE	A	30	40	–
⋮							
011	002	T <sub>1</sub>	ABORT	–	–	–	–
⋮							
026	011	T <sub>1</sub>	CLR-002	A	40	30	001

Diagram illustrating a transaction abort and the corresponding CLR (Current Log Record) entry. The log table shows the sequence of operations for transaction T<sub>1</sub>. The CLR entry (LSN 026) is generated after the abort, with its 'Before' value (40) and 'After' value (30) swapped relative to the original update operation (LSN 002). Red arrows and boxes highlight this swap, indicating the undo operation required to restore the original state.

# TRANSACTION ABORT – CLR EXAMPLE



# TRANSACTION ABORT – CLR EXAMPLE



LSN	prevLSN	TxnID	Type	Object	Before	After	UndoNext
001	nil	T <sub>1</sub>	BEGIN	–	–	–	–
002	001	T <sub>1</sub>	UPDATE	A	30	40	–
⋮							
011	002	T <sub>1</sub>	ABORT	–	–	–	–
⋮							
026	011	T <sub>1</sub>	CLR-002	A	40	30	001
027	026	T <sub>1</sub>	TXN-END	–	–	–	nil

# ABORT ALGORITHM

First write an **ABORT** record to log for the txn

Then play back updates in reverse order. For each update record:

- Write a **CLR** entry

- Restore old value

At end, write a **TXN-END** log record

Notice: CLRs **never** need to be undone

- but they might be redone when repeating history: guarantees atomicity!

# OTHER LOG-RELATED STATE

## Active Transaction Table (ATT)

One entry per currently active txn. Entry removed when txn commits or aborts

**txnId**: unique txn identifier

**status**: the current “mode” of the txn (running / committed / undo)

**lastLSN**: most recent LSN created by txn

## Dirty Page Table (DPT)

Record which pages in the buffer pool contain changes from uncommitted txns

One entry per dirty page in the buffer pool

**recLSN**: the LSN of the log record that first caused the page to be dirty



# CHECKPOINTING

Periodically, the DBMS creates a **checkpoint** to minimize the time taken to recover in the event of a system crash

“Fuzzy checkpointing” writes to log:

**CHECKPOINT-BEGIN** record: Indicates when checkpoint began

**CHECKPOINT-END** record: Contains current txn table and dirty page table

The recorded state is accurate as of the time of the **CHECKPOINT-BEGIN** record

No attempt to force dirty pages to disk! (too expensive)

Store LSN of checkpoint record in a safe place (**MasterRecord**)

# ARIES – RECOVERY PHASES

## Phase #1 – Analysis

Read WAL from last checkpoint to identify dirty pages in the buffer pool and active txns at the time of the crash

## Phase #2 – Redo

Repeat all actions starting from an appropriate point in the log (even txns that will abort)

## Phase #3 – Undo

Reverse the actions of txns that did not commit before the crash

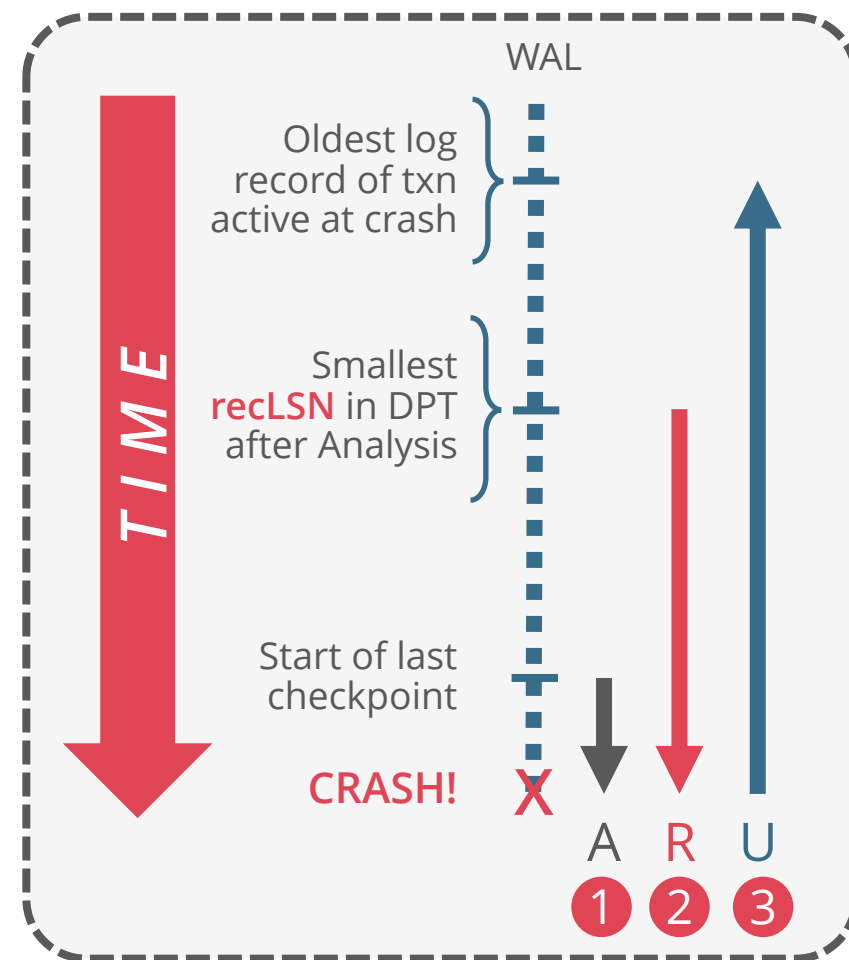
# ARIES – OVERVIEW

Start from last **BEGIN-CHECKPOINT** found via **MasterRecord**

**Analysis:** Figure out which txns committed or failed since checkpoint

**Redo:** Repeat all actions

**Undo:** Reverse effects of failed txns



# ANALYSIS PHASE

Scan log forward from last successful checkpoint

If you find a **TXN-END** record, remove its corresponding txn from **ATT**

All other records:

- Add txn to **ATT** with status **UNDO** if not already there

- Set **lastLSN** to the **LSN** of the current log record

- On commit, change status to **COMMIT**

For **UPDATE** records:

- If page **X** not in **DPT**, add **X** to **DPT** and set its **recLSN=LSN**

# ANALYSIS PHASE

At the end of the analysis phase:

**ATT** tells the DBMS which txns were active at time of crash

**DPT** tells the DBMS which dirty pages might not have made it to disk

# REDO PHASE

**Goal:** repeat history to reconstruct state at the moment of the crash

Reapply all updates (even aborted txns!) and redo **CLRs**

Scan forward from the log record containing smallest **recLSN** in **DPT**

For each update log record or **CLR** with a **LSN**, redo the action unless:

Affected page is not in the **DPT**, or

Affected page is in **DPT** but its **recLSN** is greater than the record's **LSN**, or

Affected **pageLSN** (on disk)  $\geq$  the record's **LSN**

# REDO PHASE

To redo an action

- Reapply logged action

- Set **pageLSN** to log record's *LSN*

- No additional logging, no forced flushes!

At the end of the Redo Phase, write **TXN-END** log records for all txns with status **COMMIT** and remove them from the **ATT**

# UNDO PHASE

Undo all txns that were active at the time of crash and therefore will never commit

These are all txns with **UNDO** status in the **ATT** after the Analysis Phase

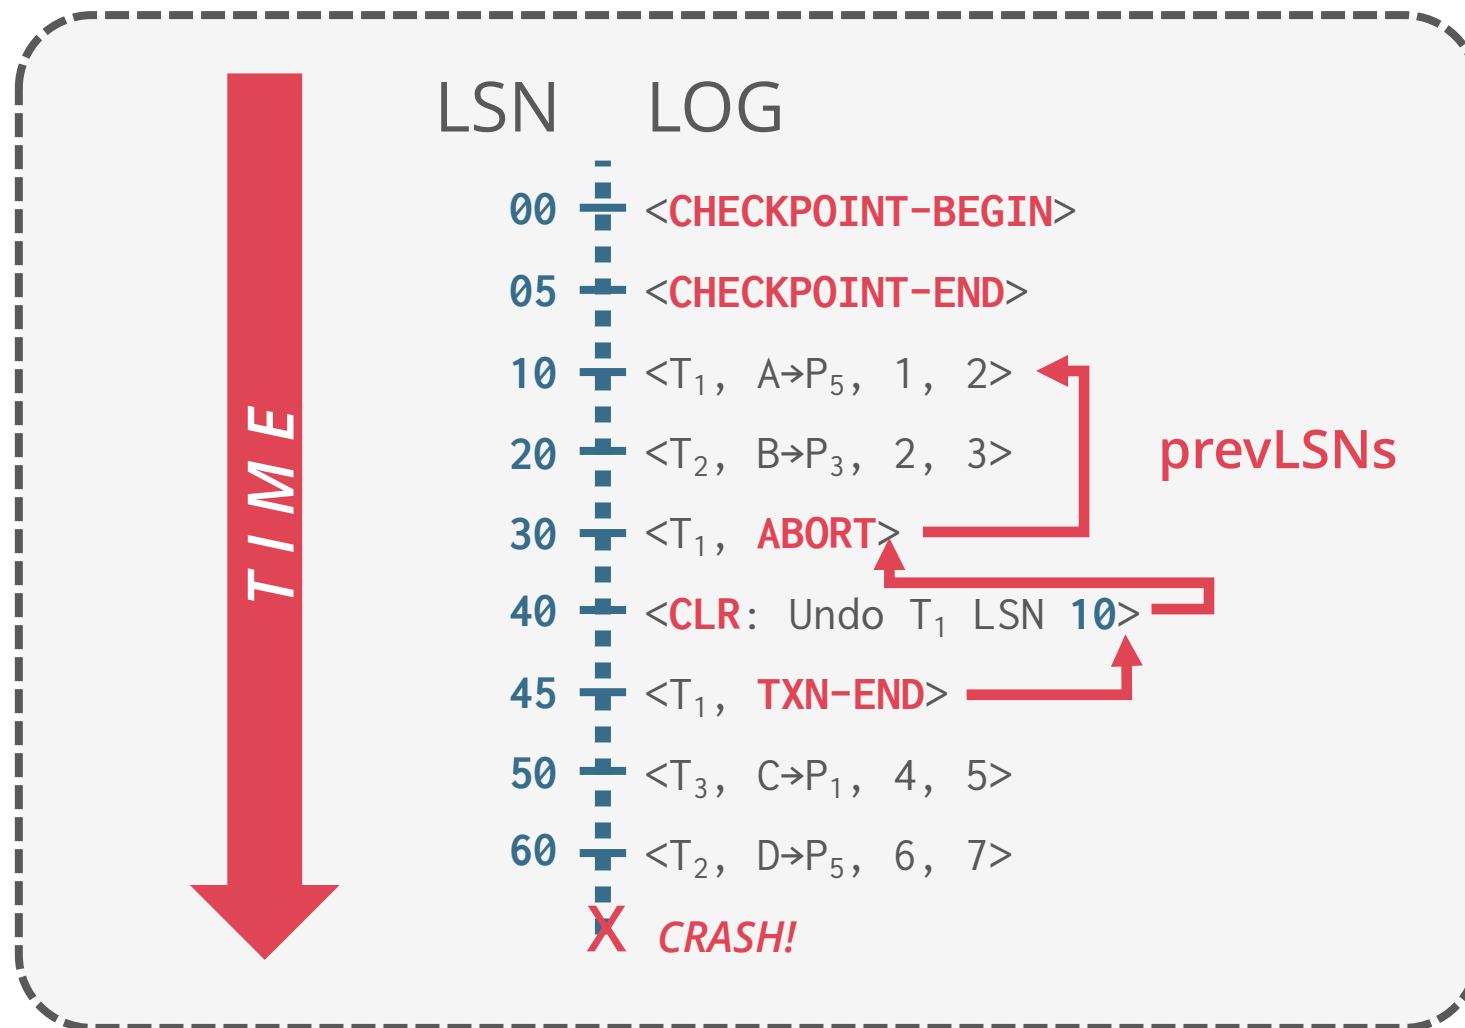
Process them in reverse **LSN** order using **lastLSN** to speed up traversal

Write a **CLR** for every modification

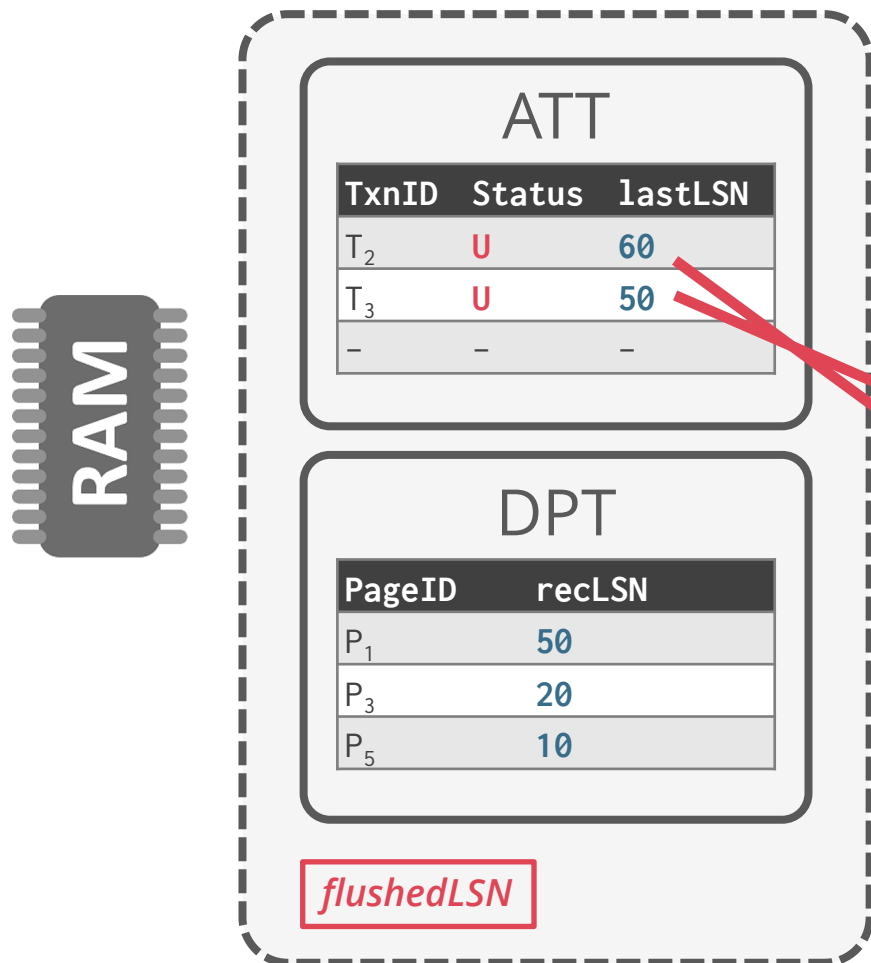
Aborting a txn is a special case of Undo Phase applied to only one txn



# FULL EXAMPLE




# FULL EXAMPLE




LSN	LOG
00,05	<CHECKPOINT-BEGIN>, <CHECKPOINT-END>
10	<T <sub>1</sub> , A→P <sub>5</sub> , 1, 2>
20	<T <sub>2</sub> , B→P <sub>3</sub> , 2, 3>
30	<T <sub>1</sub> , ABORT>
40,45	<CLR: Undo T <sub>1</sub> LSN 10>, <T <sub>1</sub> , TXN-END>
50	<T <sub>3</sub> , C→P <sub>1</sub> , 4, 5>
60	<T <sub>2</sub> , D→P <sub>5</sub> , 6, 7>
	<b>X CRASH! RESTART!</b>

# FULL EXAMPLE



LSN	LOG
00,05	<CHECKPOINT-BEGIN>, <CHECKPOINT-END>
10	<T <sub>1</sub> , A→P <sub>5</sub> , 1, 2>
20	<T <sub>2</sub> , B→P <sub>3</sub> , 2, 3>
30	<T <sub>1</sub> , ABORT>
40,45	<CLR: Undo T <sub>1</sub> LSN 10>, <T <sub>1</sub> , TXN-END>
50	<T <sub>3</sub> , C→P <sub>1</sub> , 4, 5>
60	<T <sub>2</sub> , D→P <sub>5</sub> , 6, 7>
	<b>X CRASH! RESTART!</b>
70	<CLR: Undo T <sub>2</sub> LSN 60, UndoNext 20>

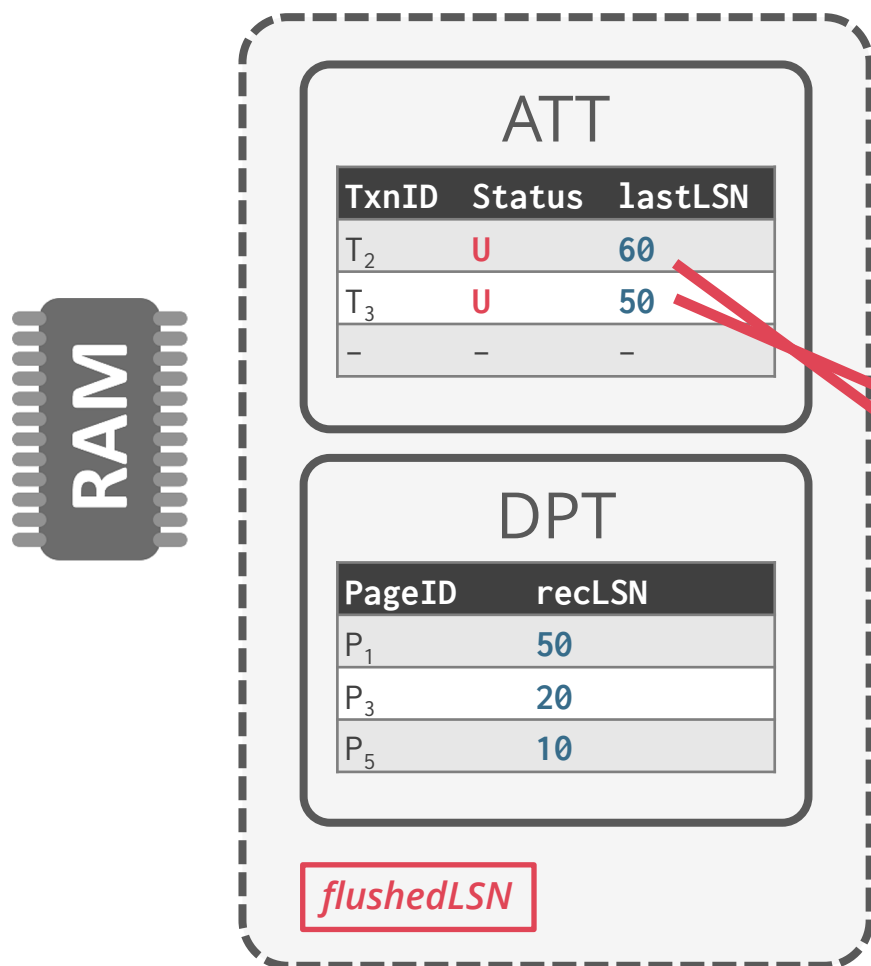
# FULL EXAMPLE



LSN	LOG
00,05	<CHECKPOINT-BEGIN>, <CHECKPOINT-END>
10	<T <sub>1</sub> , A→P <sub>5</sub> , 1, 2>
20	<T <sub>2</sub> , B→P <sub>3</sub> , 2, 3>
30	<T <sub>1</sub> , ABORT>
40,45	<CLR: Undo T <sub>1</sub> LSN 10>, <T <sub>1</sub> , TXN-END>
50	<T <sub>3</sub> , C→P <sub>1</sub> , 4, 5>
60	<T <sub>2</sub> , D→P <sub>5</sub> , 6, 7>
	<b>X CRASH! RESTART!</b>
70	<CLR: Undo T <sub>2</sub> LSN 60, UndoNext 20>
80,85	<CLR: Undo T <sub>3</sub> LSN 50> <T <sub>3</sub> , TXN-END>

*Flush dirty pages + WAL to disk!*


# FULL EXAMPLE



LSN	LOG
00,05	<CHECKPOINT-BEGIN>, <CHECKPOINT-END>
10	<T <sub>1</sub> , A→P <sub>5</sub> , 1, 2>
20	<T <sub>2</sub> , B→P <sub>3</sub> , 2, 3>
30	<T <sub>1</sub> , ABORT>
40,45	<CLR: Undo T <sub>1</sub> LSN 10>, <T <sub>1</sub> , TXN-END>
50	<T <sub>3</sub> , C→P <sub>1</sub> , 4, 5>
60	<T <sub>2</sub> , D→P <sub>5</sub> , 6, 7>
	<b>X CRASH! RESTART!</b>
70	<CLR: Undo T <sub>2</sub> LSN 60, UndoNext 20>
80,85	<CLR: Undo T <sub>3</sub> LSN 50> <T <sub>3</sub> , TXN-END>
	<b>X CRASH! RESTART!</b>

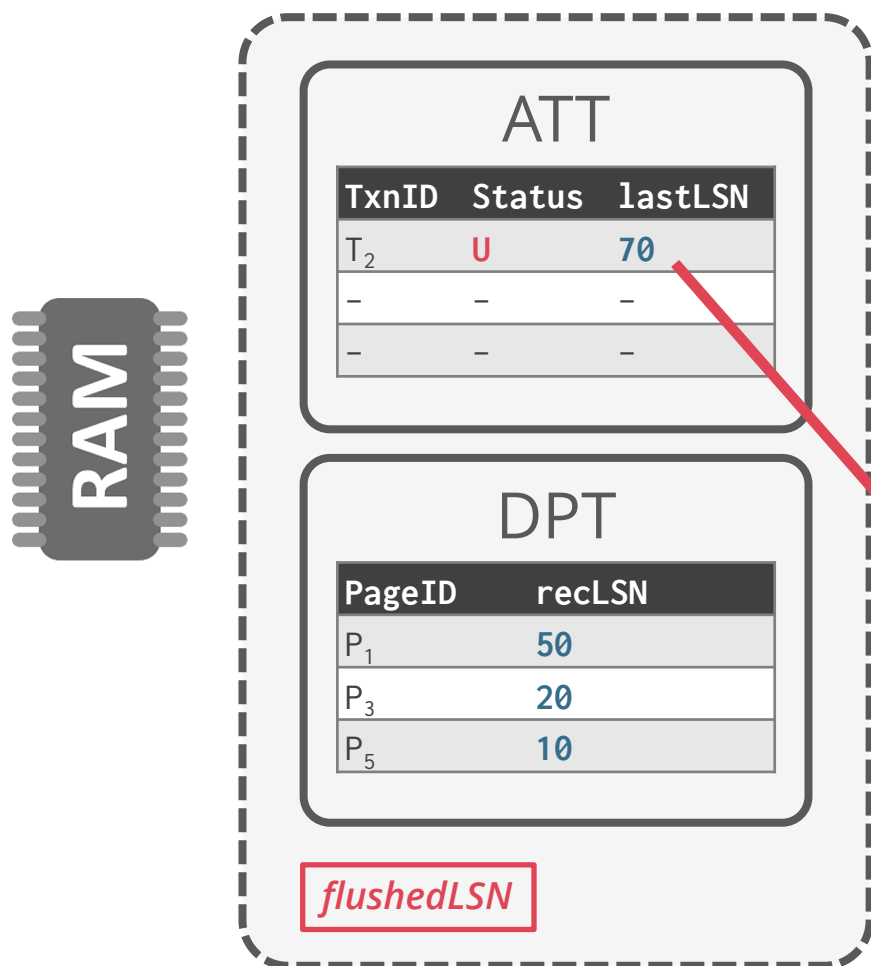
*Flush dirty pages + WAL to disk!*

# FULL EXAMPLE




LSN	LOG
00,05	<CHECKPOINT-BEGIN>, <CHECKPOINT-END>
10	<T <sub>1</sub> , A→P <sub>5</sub> , 1, 2>
20	<T <sub>2</sub> , B→P <sub>3</sub> , 2, 3>
30	<T <sub>1</sub> , ABORT>
40,45	<CLR: Undo T <sub>1</sub> LSN 10>, <T <sub>1</sub> , TXN-END>
50	<T <sub>3</sub> , C→P <sub>1</sub> , 4, 5>
60	<T <sub>2</sub> , D→P <sub>5</sub> , 6, 7>
	<b>X CRASH! RESTART!</b>
70	<CLR: Undo T <sub>2</sub> LSN 60, UndoNext 20>
80,85	<CLR: Undo T <sub>3</sub> LSN 50> <T <sub>3</sub> , TXN-END>
	<b>X CRASH! RESTART!</b>

# FULL EXAMPLE



LSN	LOG
00,05	<CHECKPOINT-BEGIN>, <CHECKPOINT-END>
10	<T <sub>1</sub> , A→P <sub>5</sub> , 1, 2>
20	<T <sub>2</sub> , B→P <sub>3</sub> , 2, 3>
30	<T <sub>1</sub> , ABORT>
40,45	<CLR: Undo T <sub>1</sub> LSN 10>, <T <sub>1</sub> , TXN-END>
50	<T <sub>3</sub> , C→P <sub>1</sub> , 4, 5>
60	<T <sub>2</sub> , D→P <sub>5</sub> , 6, 7>
	<del>X CRASH! RESTART!</del>
70	<CLR: Undo T <sub>2</sub> LSN 60, UndoNext 20>
80,85	<CLR: Undo T <sub>3</sub> LSN 50> <T <sub>3</sub> , TXN-END>
	<del>X CRASH! RESTART!</del>

# FULL EXAMPLE



LSN	LOG
00,05	<CHECKPOINT-BEGIN>, <CHECKPOINT-END>
10	<T <sub>1</sub> , A→P <sub>5</sub> , 1, 2>
20	<T <sub>2</sub> , B→P <sub>3</sub> , 2, 3>
30	<T <sub>1</sub> , ABORT>
40,45	<CLR: Undo T <sub>1</sub> LSN 10>, <T <sub>1</sub> , TXN-END>
50	<T <sub>3</sub> , C→P <sub>1</sub> , 4, 5>
60	<T <sub>2</sub> , D→P <sub>5</sub> , 6, 7>
	<del>X CRASH! RESTART!</del>
70	<CLR: Undo T <sub>2</sub> LSN 60, UndoNext 20>
80,85	<CLR: Undo T <sub>3</sub> LSN 50> <T <sub>3</sub> , TXN-END>
	<del>X CRASH! RESTART!</del>
90,95	<CLR: Undo T <sub>2</sub> LSN 20> <T <sub>2</sub> , TXN-END>



# ADDITIONAL CRASH ISSUES

What does the DBMS do if it crashes during recovery in Analysis Phase?

Nothing. Just run recovery again

What does the DBMS do if it crashes during recovery in Redo Phase?

Again nothing. Redo everything again

# CONCLUSION: ARIES

## Write-Ahead Logging with **STEAL** + **NO-FORCE**

Any change is recorded in log on stable storage before the database change is written to disk

Fuzzy checkpointing (snapshot of dirty pageIDs and currently active txns)

## Repeating History During Redo

On restart, redo everything since the earliest dirty page

## Undo txns that never commit

Write CLRs when undoing to survive failures during restarts

# ARIES

Recovery algorithm developed at IBM Research in early 1990s

## Write-Ahead Logging

Any change is recorded in log on stable storage before the change is written to disk

Must use **STEAL** + **NO-FORCE** buffer pool policies

## Repeating History During Redo

On restart, retrace actions and restore database to exact state before crash

## Logging Changes During Undo

Log also undo actions to ensure action is not repeated in the event of repeated failures